

Transactors

Simpler Concurrency

Transactors

Simpler Concurrency

VIKTOR KLANG
R&D DIRECTOR



The Short Story

- Developer since 1998
- Started on C
- Been on the JVM since 2001
- Background in business software
- Pragmatist
- OSS Enthusiast
- Concurrency Maniac

About Akka

Akka is the platform for the next generation event-driven, scalable and fault-tolerant architectures on the Java Virtual Machine

Before we start

- This will be **JVM**-focused
- There is **no** general optimal solution
- Problems, tools and the problem with tools
- I expect a **continuous flow** of questions!

This presentation will hurt

But it **will** be worth it

What is Concurrency?

What about Parallelism?

Computer programs
are lifeless collections
of instructions

Processes

A Process is

- An **instance** of a computer program
- Has its own **timeline**
- Call **Stack**
- Address space (**Heap**)
- Security attributes and more...

Process gotchas

- OS Processes **can be expensive to create** (win32)
- Inter-Process **Communication has overhead**
- Context switching **costs**
- **Hangs onto resources** when idling
- **Error management and recovery is hard**

Threads

A Thread is

- An standalone execution path within a process
- Has its own **timeline**
- Call **Stack**
- Shares Address space (**Heap**) with other Threads within the same process

Thread gotchas

- Expensive to create
- Context switching costs
- Suffers from diminishing returns
- Hangs onto resources when idling
- Lifecycle is hard to control from the outside
- Error handling and recovery is hard

Others

- **Actors**
- Fibers
- Executor services
- Dataflow concurrency
- Parallel collections
- Agents
- and more...

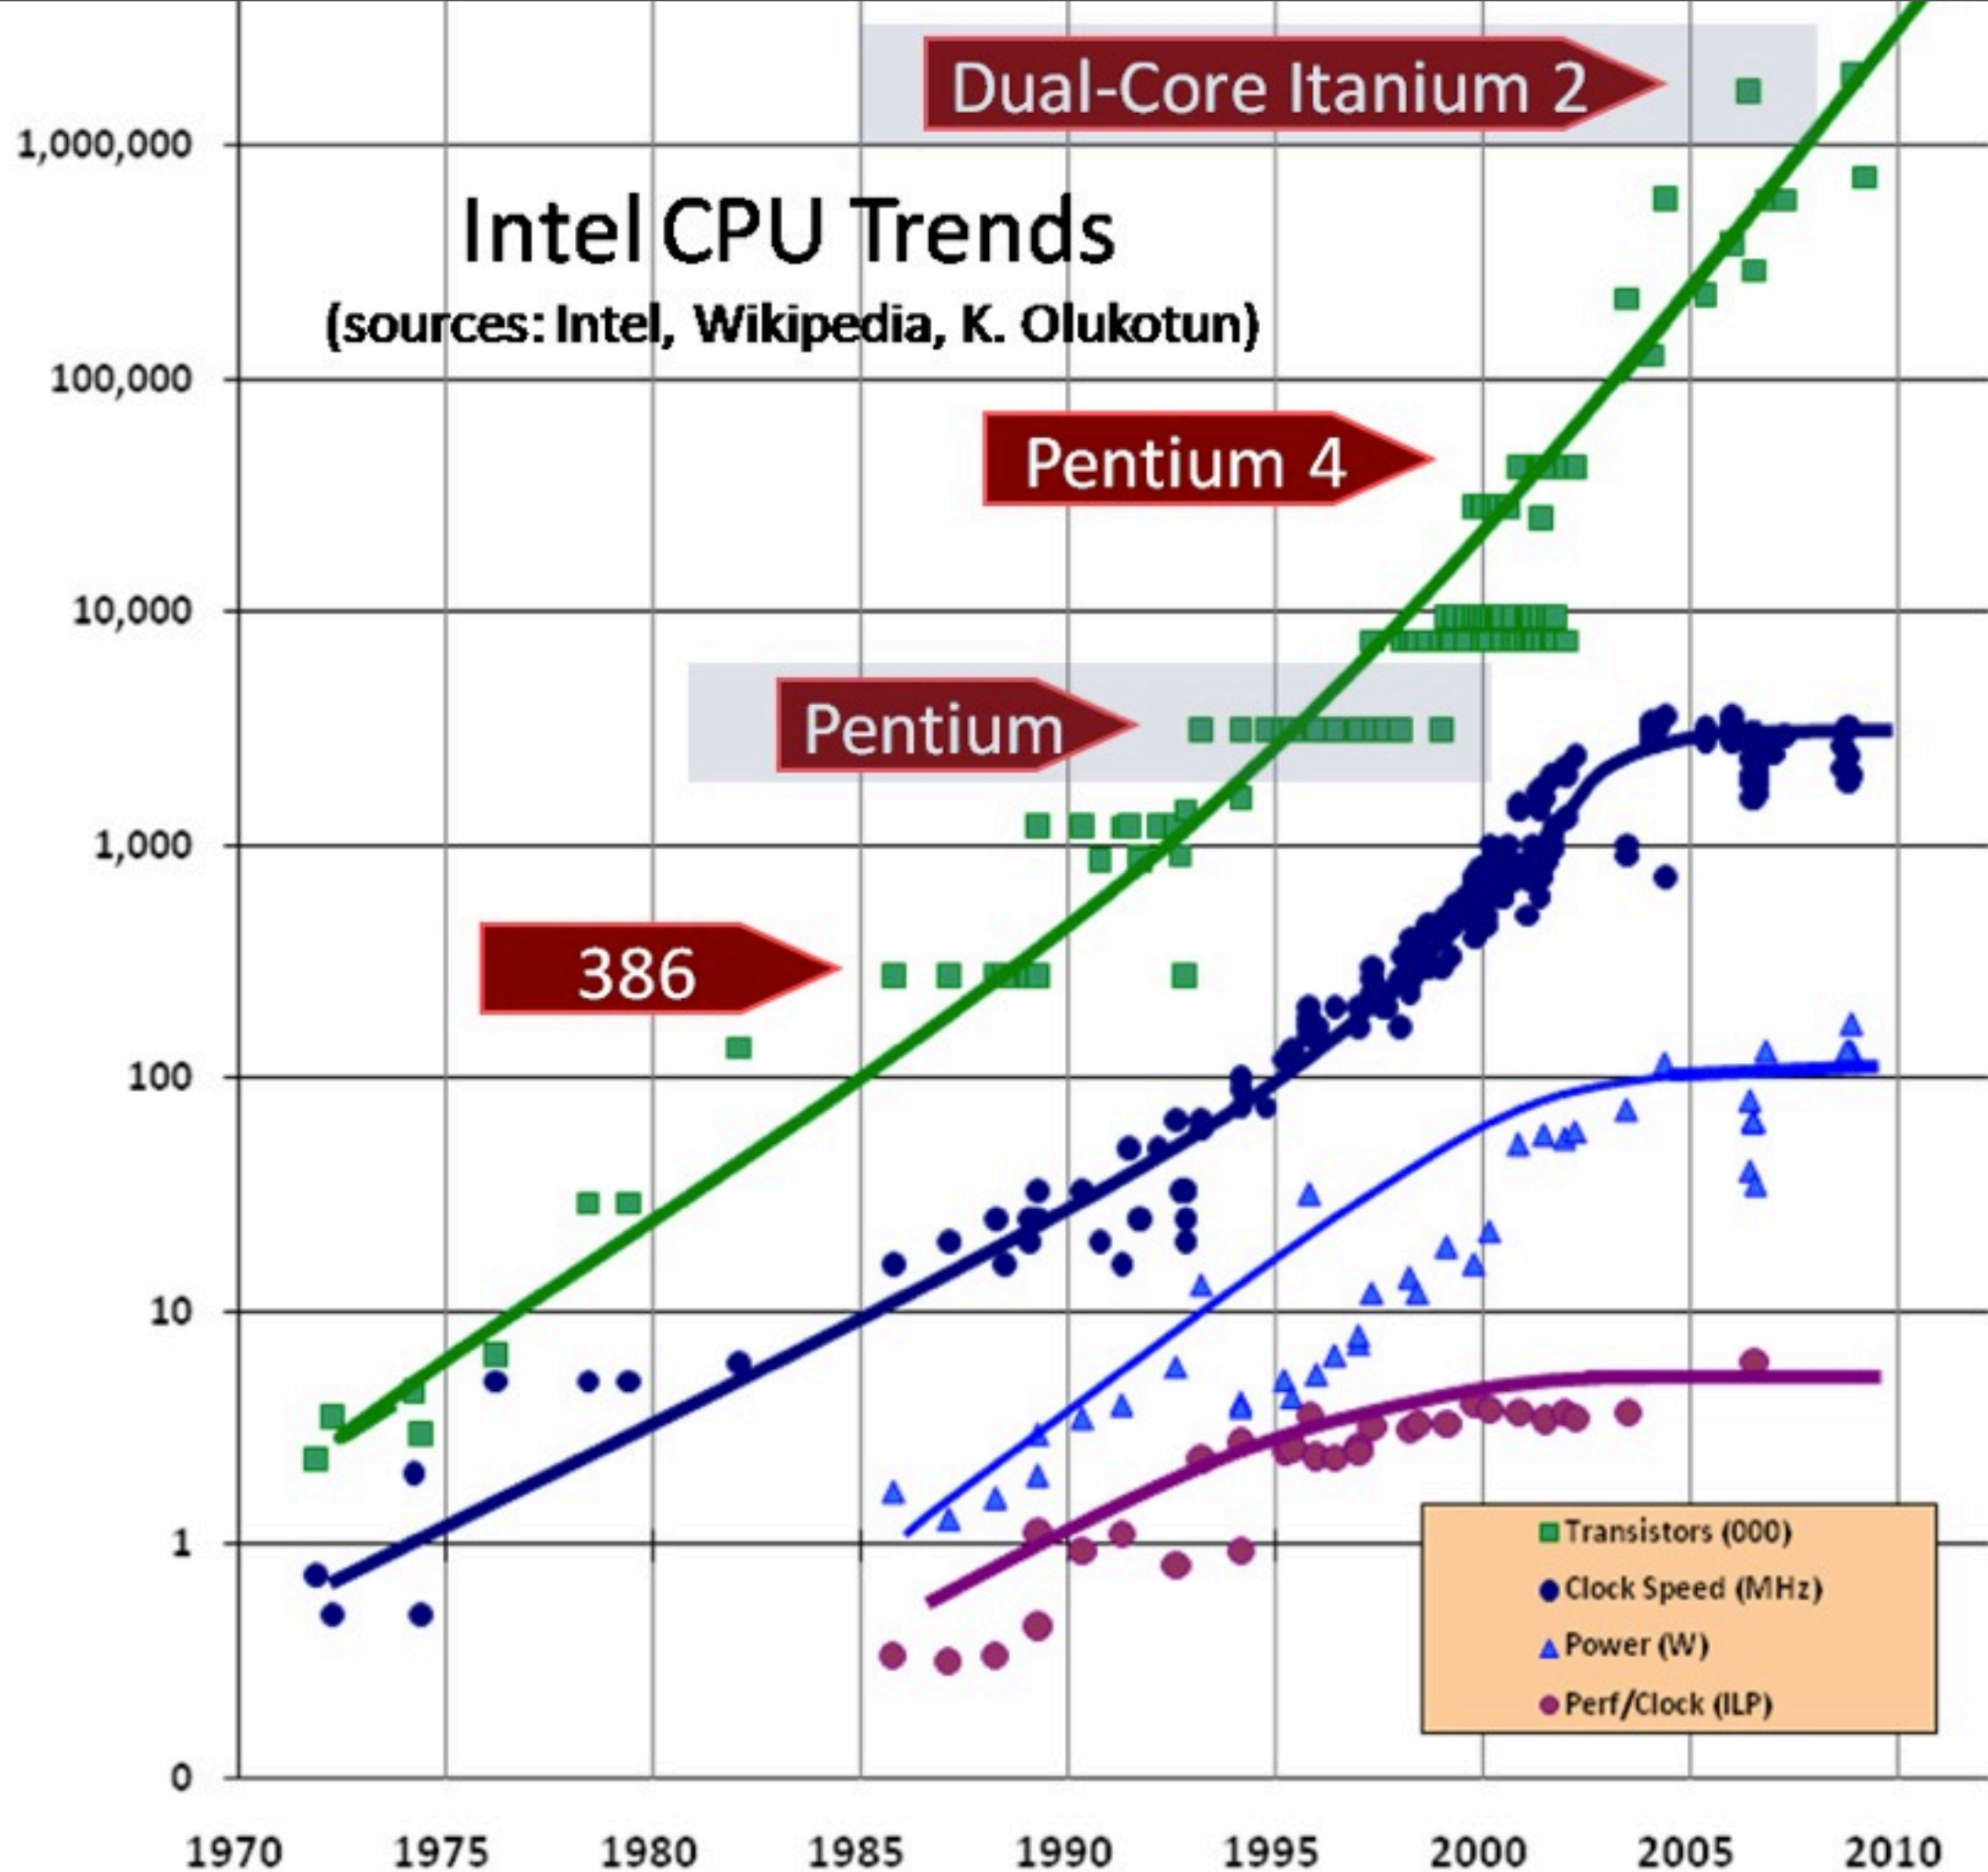
Schedulers

- The job of the **OS** scheduler is to allocate available CPU time **between processes**
- The **JVM** has its own scheduler to distribute its share of CPU time between threads

Some
bad news guys...

The free lunch is over

Herb Sutter 2005 <http://www.gotw.ca/publications/concurrency-ddj.htm>



Chip manufacturers will
focus on multicore technology
as the way to
improve performance

As a result -

Software developers will be
forced to write massively
multithreaded software

Some

good news...

Multicore

gives us

parallelism

Sync vs Async

The SLA

- 4 Services
 1. 3 seconds
 2. 4 seconds
 3. 1 second
 4. 6 seconds
- The services are unrelated
- Results need to be joined and returned
- Your budget is 7 seconds

The SLA: Sync

```
def callServices = {  
  val first    = callService1 //3 seconds  
  val second  = callService2 //4 seconds  
  val third    = callService3 //1 seconds  
  val fourth   = callService4 //6 seconds  
  return first + second + third + fourth  
}
```

The SLA: Async

```
def callServices = {  
    val first    = callService1 //3 seconds  
    val second  = callService2 //4 seconds  
    val third    = callService3 //1 seconds  
    val fourth   = callService4 //6 seconds  
    return first.await + second.await +  
           third.await + fourth.await  
}
```

Result: The SLA

- Synchronous calls

14 seconds

- Asynchronous calls

6 seconds

The Store



Rules!

- Current stock is non-negative
- Must not allow overselling
- Must not miscommunicate to the customer
- More than 1 customer needs to be able to simultaneously use the shop

It's a minefield out
there!

Example pitfalls

- Race-conditions
- Deadlocks
- Lock Convoys
- Starvation
 - Livelocks
- Data corruption
- Visibility and ordering issues

Race-conditions

REPL time!

Deadlocks

ThreadA:

lockA.lock

lockB.lock

ThreadB:

lockB.lock

lockA.lock

REPL time!

Lock Convoys

Starvation & Livelocks

Data corruption

REPL time!

Visibility and ordering

REPL time!

```
class SomeClass {  
    var first = 0  
    var second = 0  
  
    def foo() {  
        first = 1  
        second = 1  
    }  
  
    def bar() {  
        if (second == 1 && first == 0)  
            println("can't happen")  
    }  
}
```

The usual suspects

- Mutable shared state
- Side effects
- Wrong use of concurrency control structures/techniques

```
class Counter {  
    private val count: Long = 1000000  
  
    def increment(): Long = count += 1  
  
    def decrement(): Long = count -= 1  
  
    def set(newCount: Long) {  
        count = newCount  
    }  
  
    def value(): Long = count  
}
```


Concurrency Control:

Encoding correct behavior

Concurrency primitives

Compare-And-Swap

- Atomic
- Conditional modification
- Low level construct
- **Very** useful

CTS example

```
class MyClass {  
  private var isActive = false  
  
  def start() {  
    if (!isActive) {  
      isActive = true  
      //Code to be run on start  
    }  
  }  
  
  def shutdown() {  
    if (isActive) {  
      isActive = false  
      //Code to be run on shutdown  
    }  
  }  
}
```

CAS example

```
class MyClass {  
  val isActive = new AtomicBoolean(false)  
  
  def start() {  
    if (isActive.compareAndSet(false,true)) {  
      //Code to be run on start  
    }  
  }  
  
  def shutdown() {  
    if (isActive.compareAndSet(true,false)) {  
      //Code to be run on shutdown  
    }  
  }  
}
```

Volatile variables

- Makes sure that reads and writes of the variable are against main memory
- Access synchronizes all cached copies of variables with main memory
- Reads and writes are atomic, but NOT lockable.

Monitor

- A thread safe? object
- **Serializes** execution of its methods
- Employs the use of **wait** and **notify**
- `java.lang.Object` is a Monitor but requires the use of the **synchronized** keyword to mark methods that needs mutual exclusion
- Access synchronizes **all cached copies of variables** with main memory

Locks

- Controls access to a common resource
- May support reentrancy
- May support readers/writer

The loo



Photograph © [Andrew Dunn](#), 1992.

Semaphore

- Controls access to a common resource
- Binary (MUTual EXclusion) or Counting

The Restaurant



Latches

- A condition **starting out as false**
- When set to true **remains true forever**
- Enables Threads to **wait for it** to become true

Barrier

- Any thread must **stop at the barrier**
- Cannot proceed until **enough threads** have arrived
- Can be viewed as **a threshold**

Threads... well...

The “**Threads & locks**”-combination is **still**
taught as **The Way to Do It**

The problem is

Knowing up front which locks need to be locked
Breaks encapsulation

All locks need to be taken **before** modification - in
the **correct order**

All locks **must be unlocked**

Thread are expensive and their lifecycle
management is **error prone**

“Threads are to Concurrency as
Witchcraft is to Physics”

“Hanging by a thread is the punishment
for Shared State Concurrency”

- Gilad Bracha

Java Language architect and maintainer and co-author of the Java Language Specification.

Software Transactional Memory (STM)

Transaction?

- A list of **operations**
- Atomic: **All or nothing**
- Usually
 - **A**tomtic
 - **C**onsistent
 - **I**solated
 - **D**urable
- Example: Database transaction

STM Overview

- View the memory (heap and stack) as a **transactional dataset**
- Similar to a database
 - begin
 - commit
 - abort/rollback
- Transactions are **retried automatically** upon collision
- **Rolls back** changes on abort

STM Overview

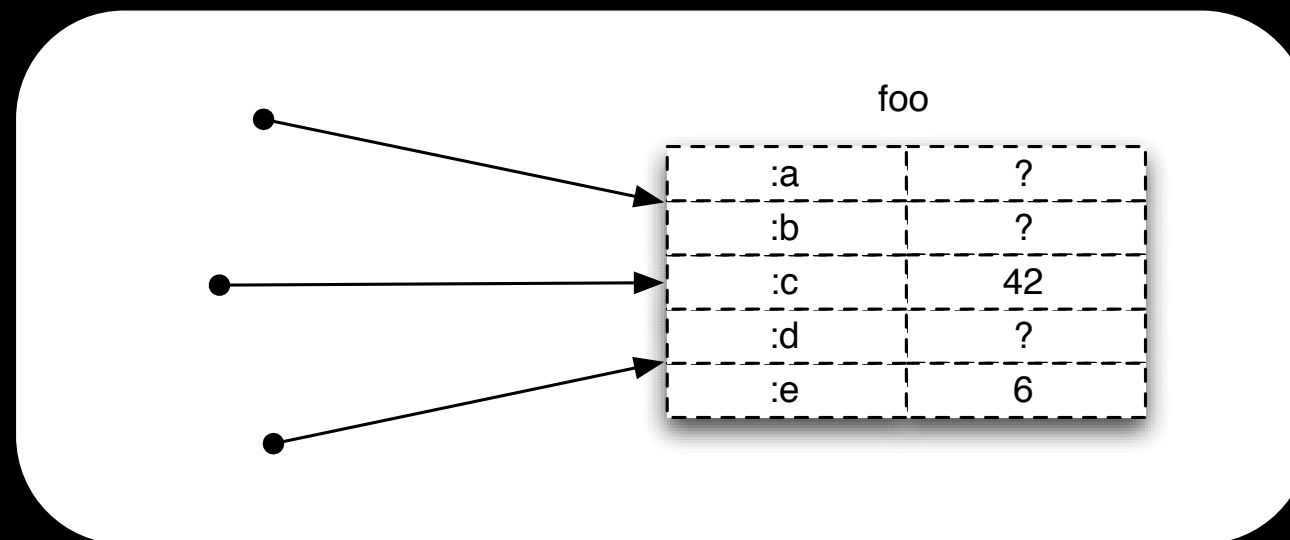
- Atomic
- Consistent
- Isolated
- Transactions can nest
- Transactions can compose!

STM restrictions

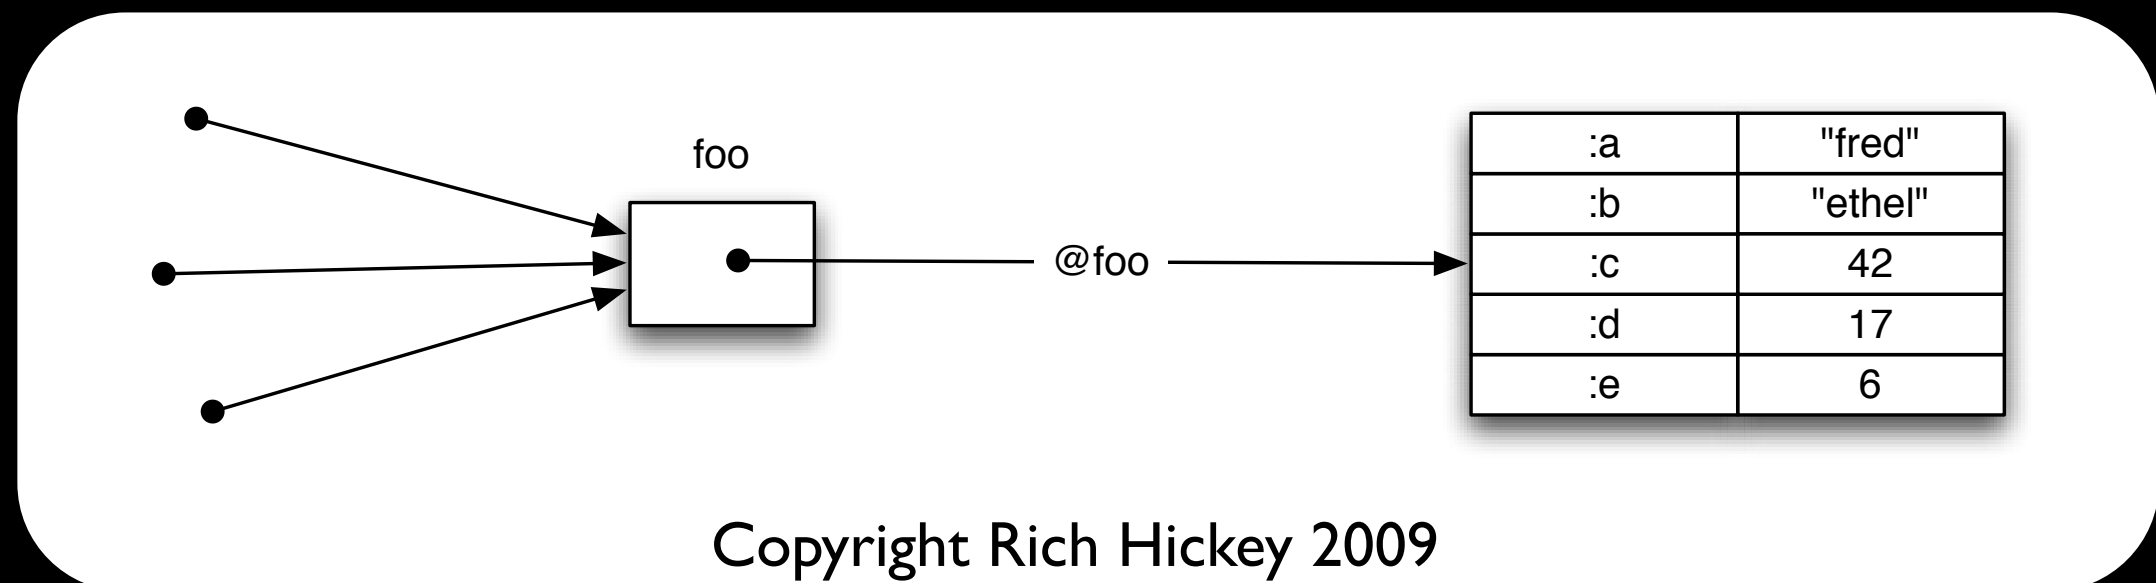
- Operations within the scope of a transaction
 - **need** to be idempotent
 - **can't** have side-effects

Managed References

Typical OO: direct access to mutable objects



Managed Reference: separates Identity & Value



Managed References

Separates **Identity** from **Value**

- **Values** are **immutable**
- **Identity** (Ref) holds **Values**

Change is a function

Compare-and-swap (**CAS**)

Abstraction of time

Must be used **within a transaction**

Transactional Reference

Creation:

```
val number = Ref(0)
```

Usage:

```
atomic {  
    number.alter( _ + 5 )  
    number.swap(10)  
    val result = number.get  
}
```


Transactional Map

Creation:

```
val beers = TransactionalMap[String,Float]( )
```

Usage:

```
atomic {  
    beers.put("Heineken",5.0f)  
    val alcoholPct = beers("Heineken")  
}
```

Transactional **Vector**

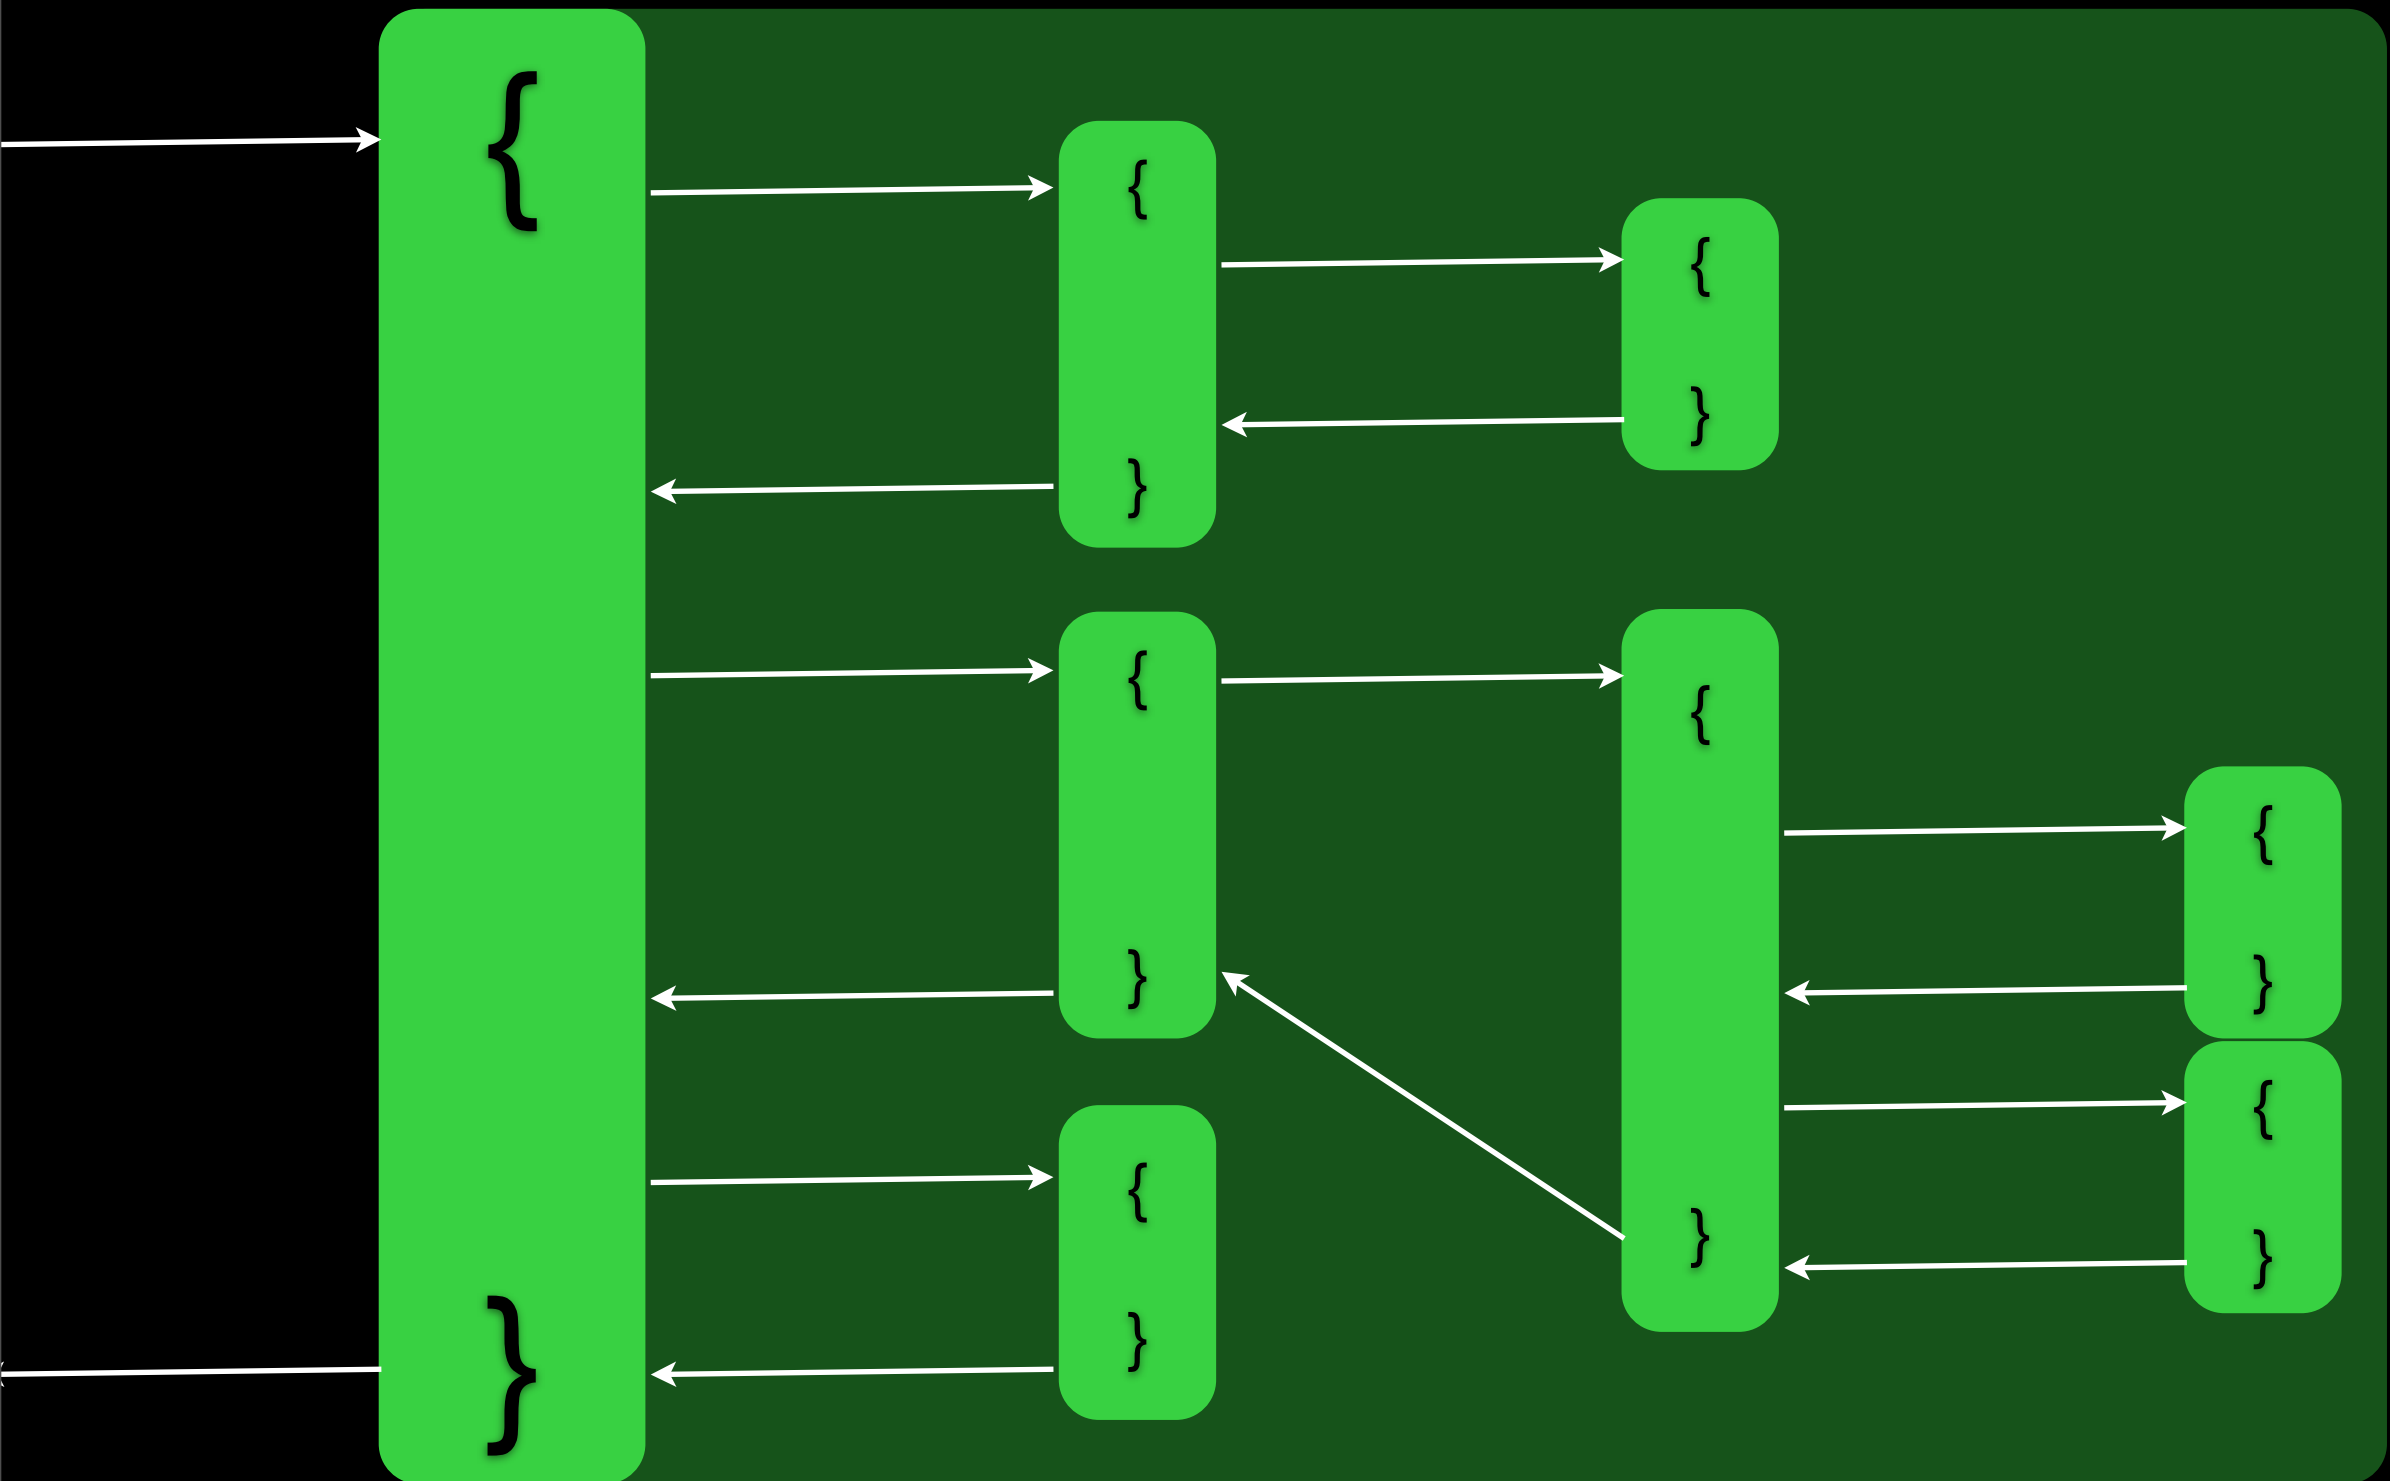
Creation:

```
val users = TransactionalVector(new User(""))
```

Usage:

```
atomic {  
    users.add(new User("Jonas"))  
    users.update(0, new User("Viktor"))  
    val first = users.get(0)  
}
```

Transaction composition

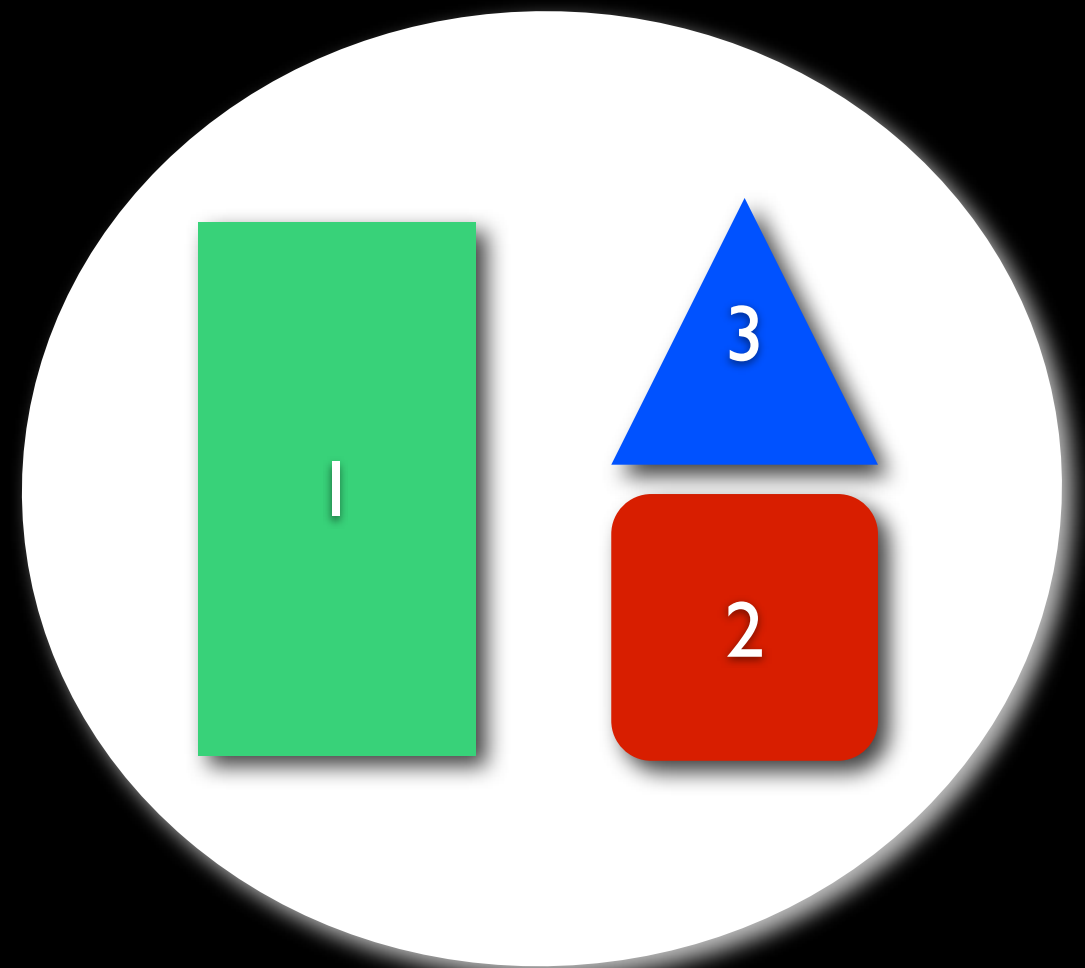


REPL time!

Actors

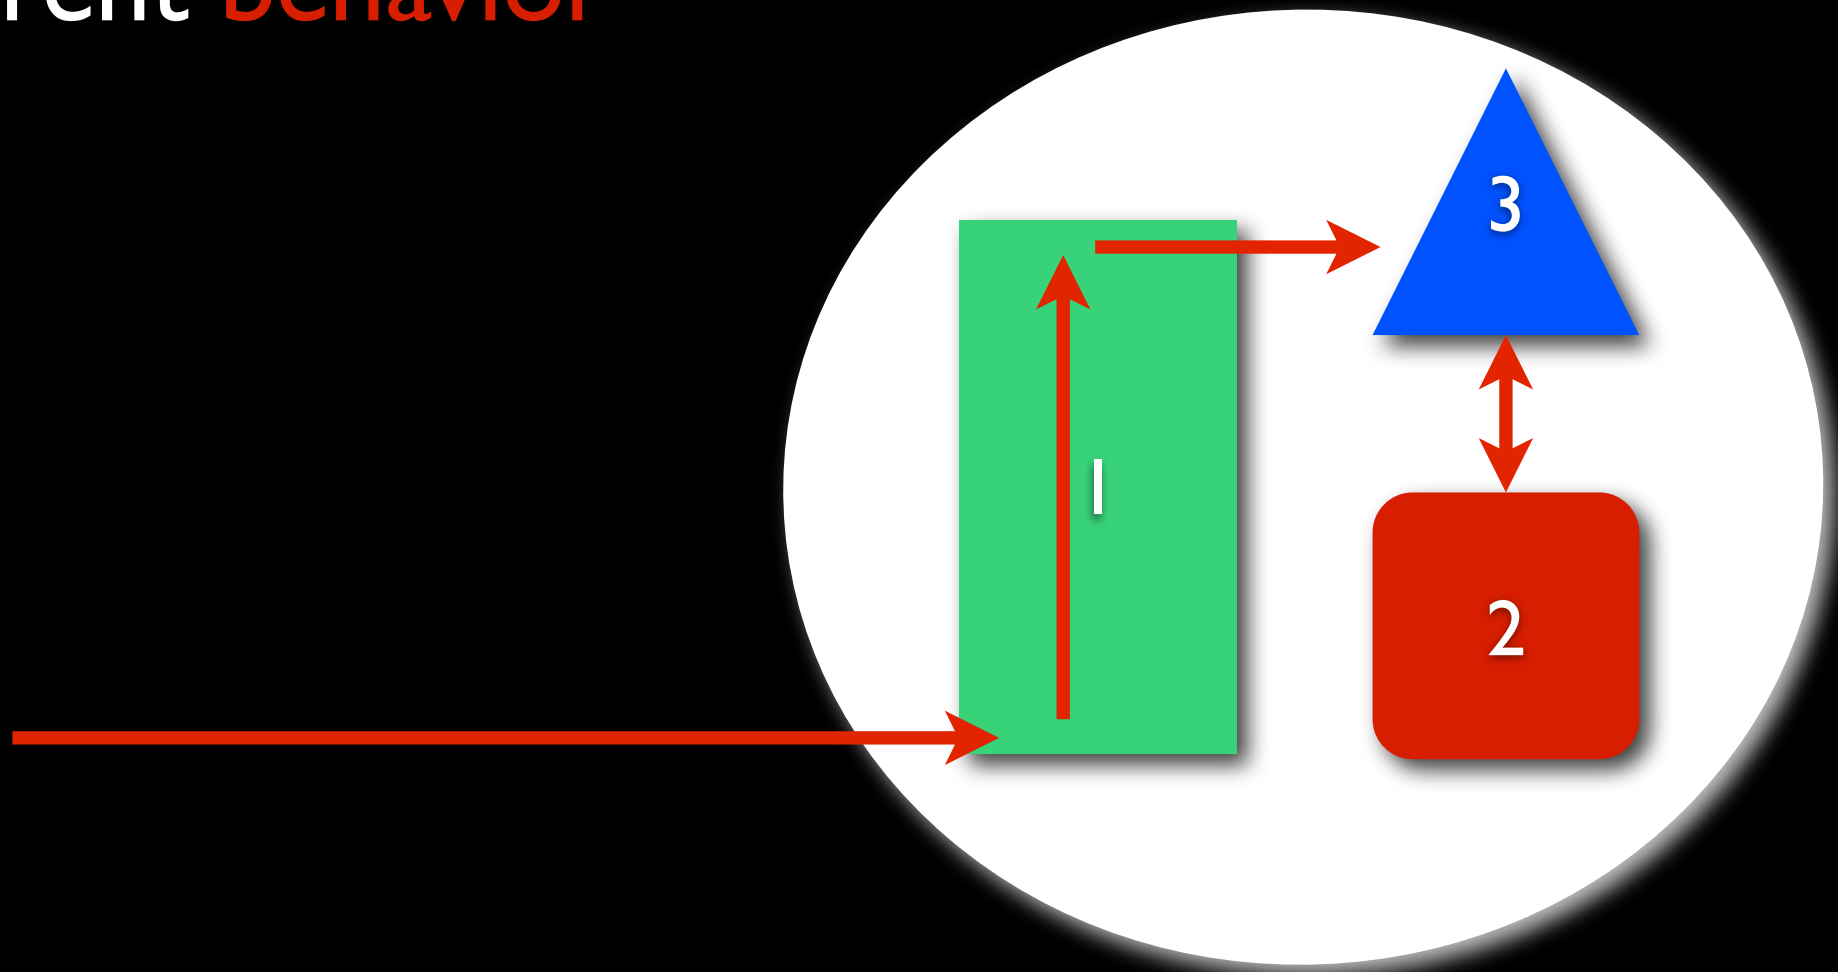
An Actor is

1. A mailbox
2. State (Optional)
3. Current behavior



Message flow

1. The mailbox
2. The State (Optional)
3. The Current behavior



Benefits of Actors

- One message at a time = **no locks**
- Inherently **asynchronous**
- **Cheap!**
- Literally **MILLIONS** at the same time

Actor Essentials

- An actor can send messages to other actors
- Can create new actors
- Can change behavior towards the next message
- Processes one message at a time

Lifecycle

- Unstarted
- Running
 - BeingRestarted
- Shutdown

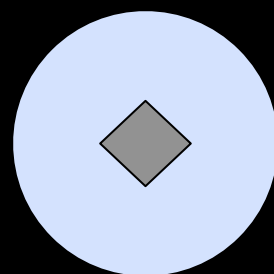
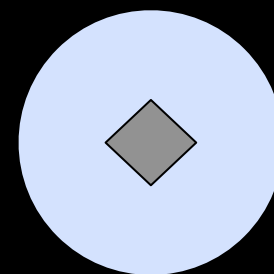
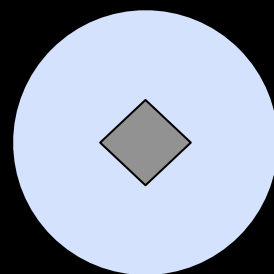
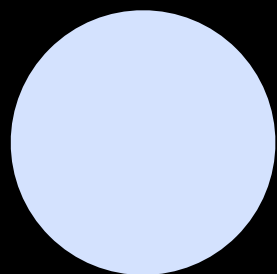
Defining an Actor

```
class MyActor extends Actor {  
  def receive = {  
    case "hello world" => reply("This is getting old man!")  
    case 5 => reply("I really don't like 5")  
    case i: Int => reply(i * i)  
    case _ => //I ignore other messages  
  }  
}
```

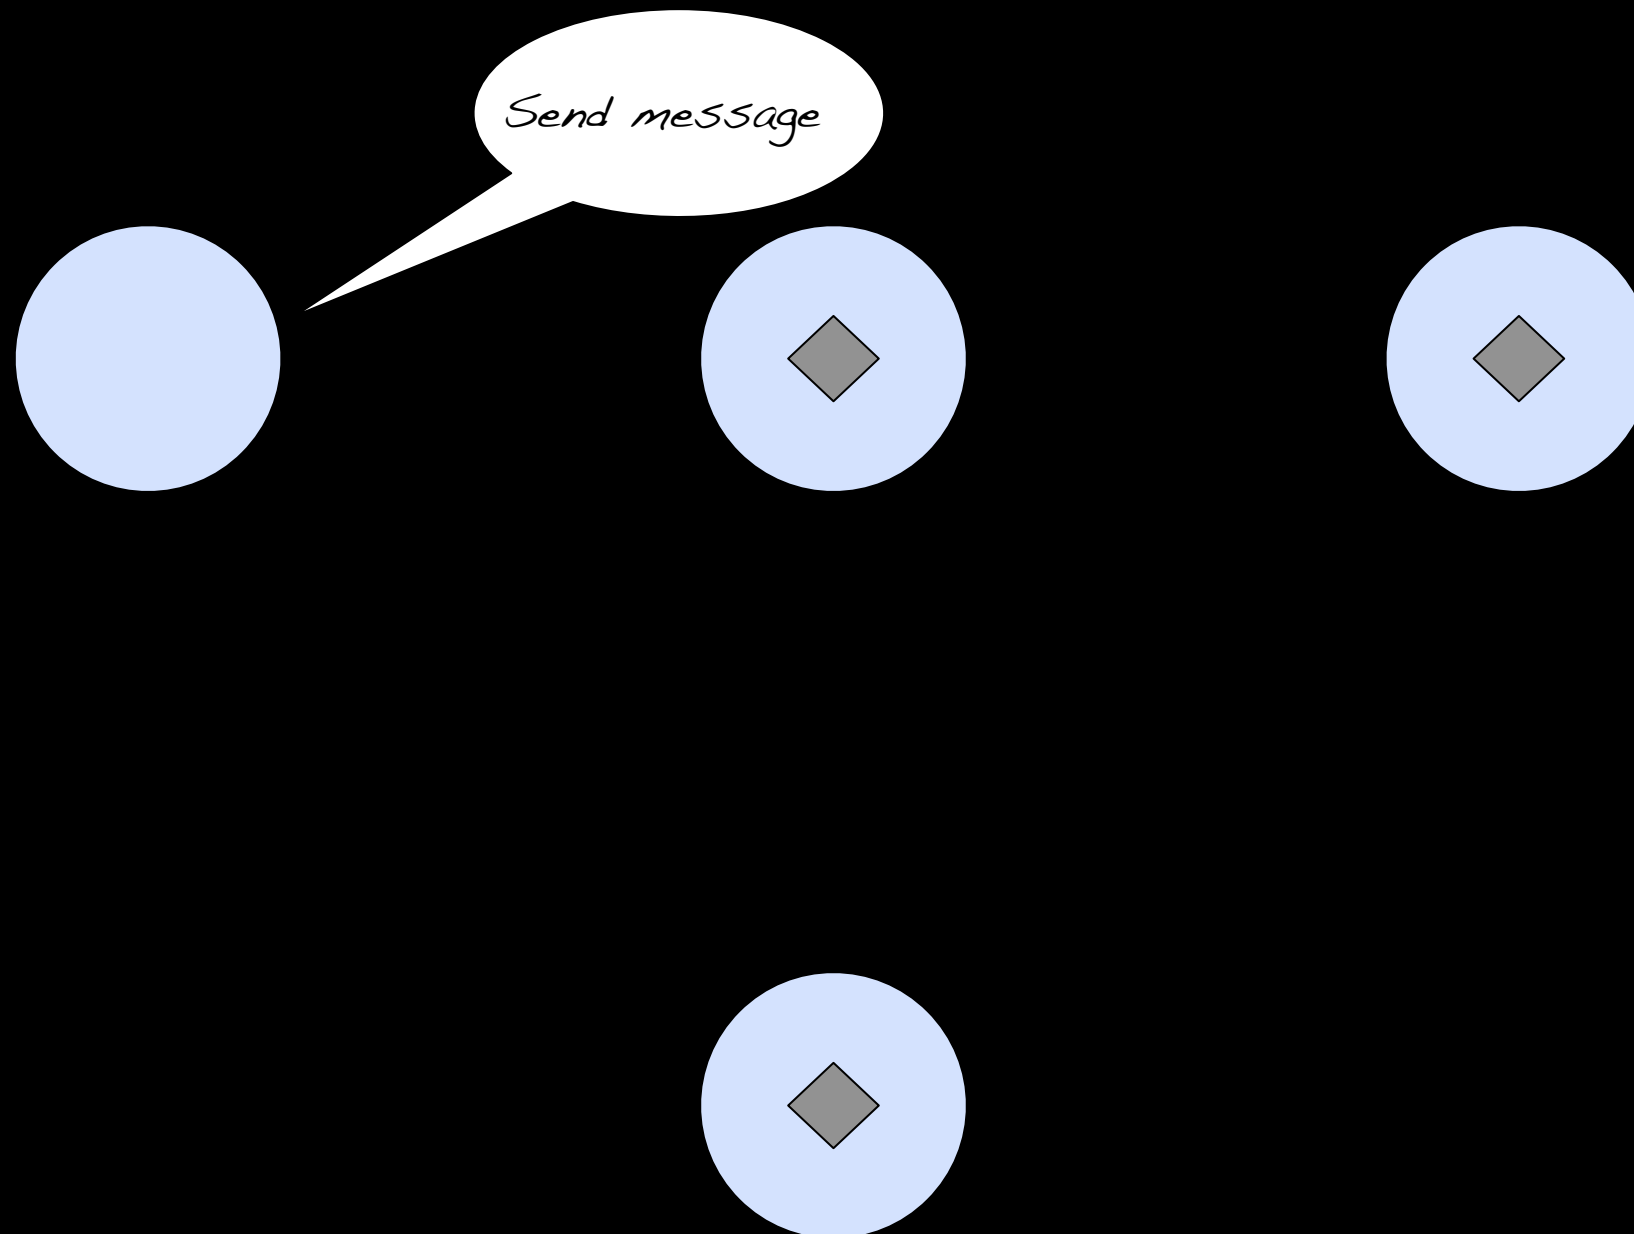
Sending messages

Method	Scala	Java
Send message	actor ! message	actor.sendOneWay(message)
Send reply within	actor !! message	actor.sendRequestReply(message)
Send reply eventually	actor !!! message	actor.sendRequestReplyFuture(message)

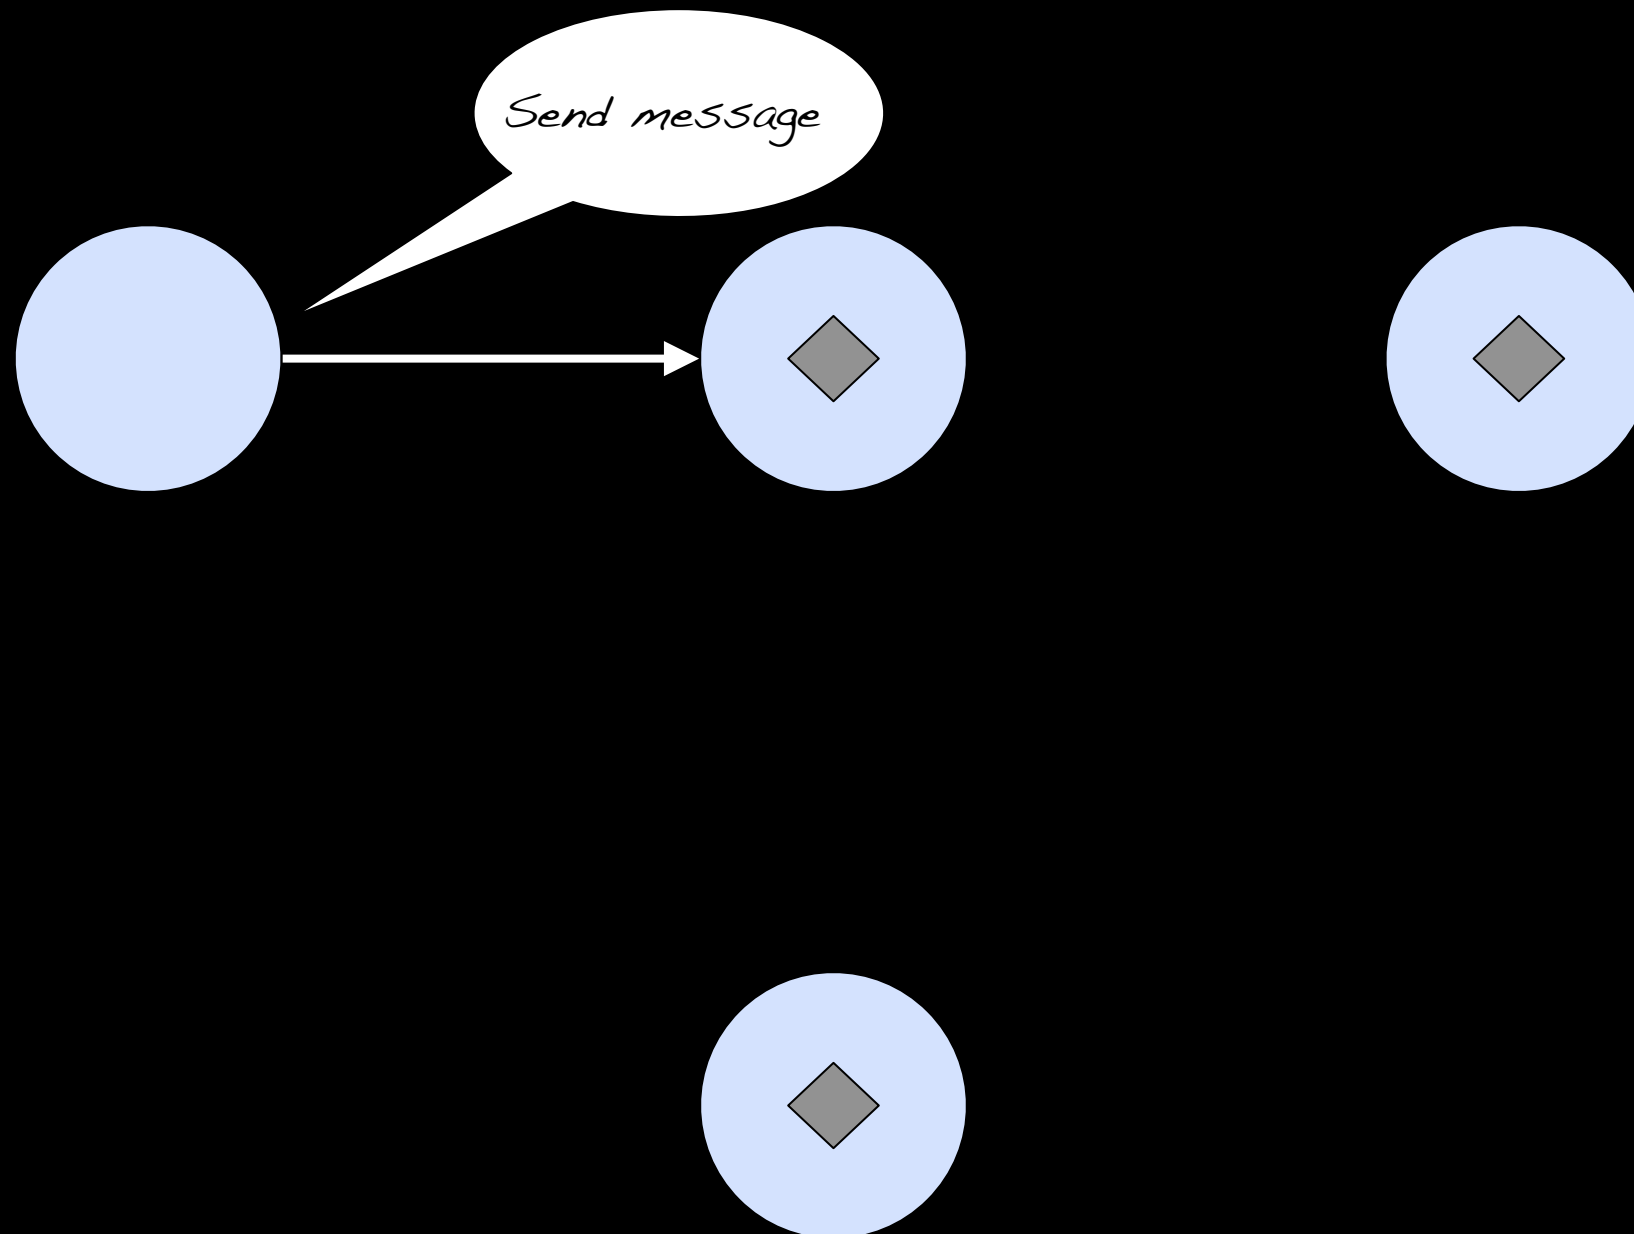
Actors



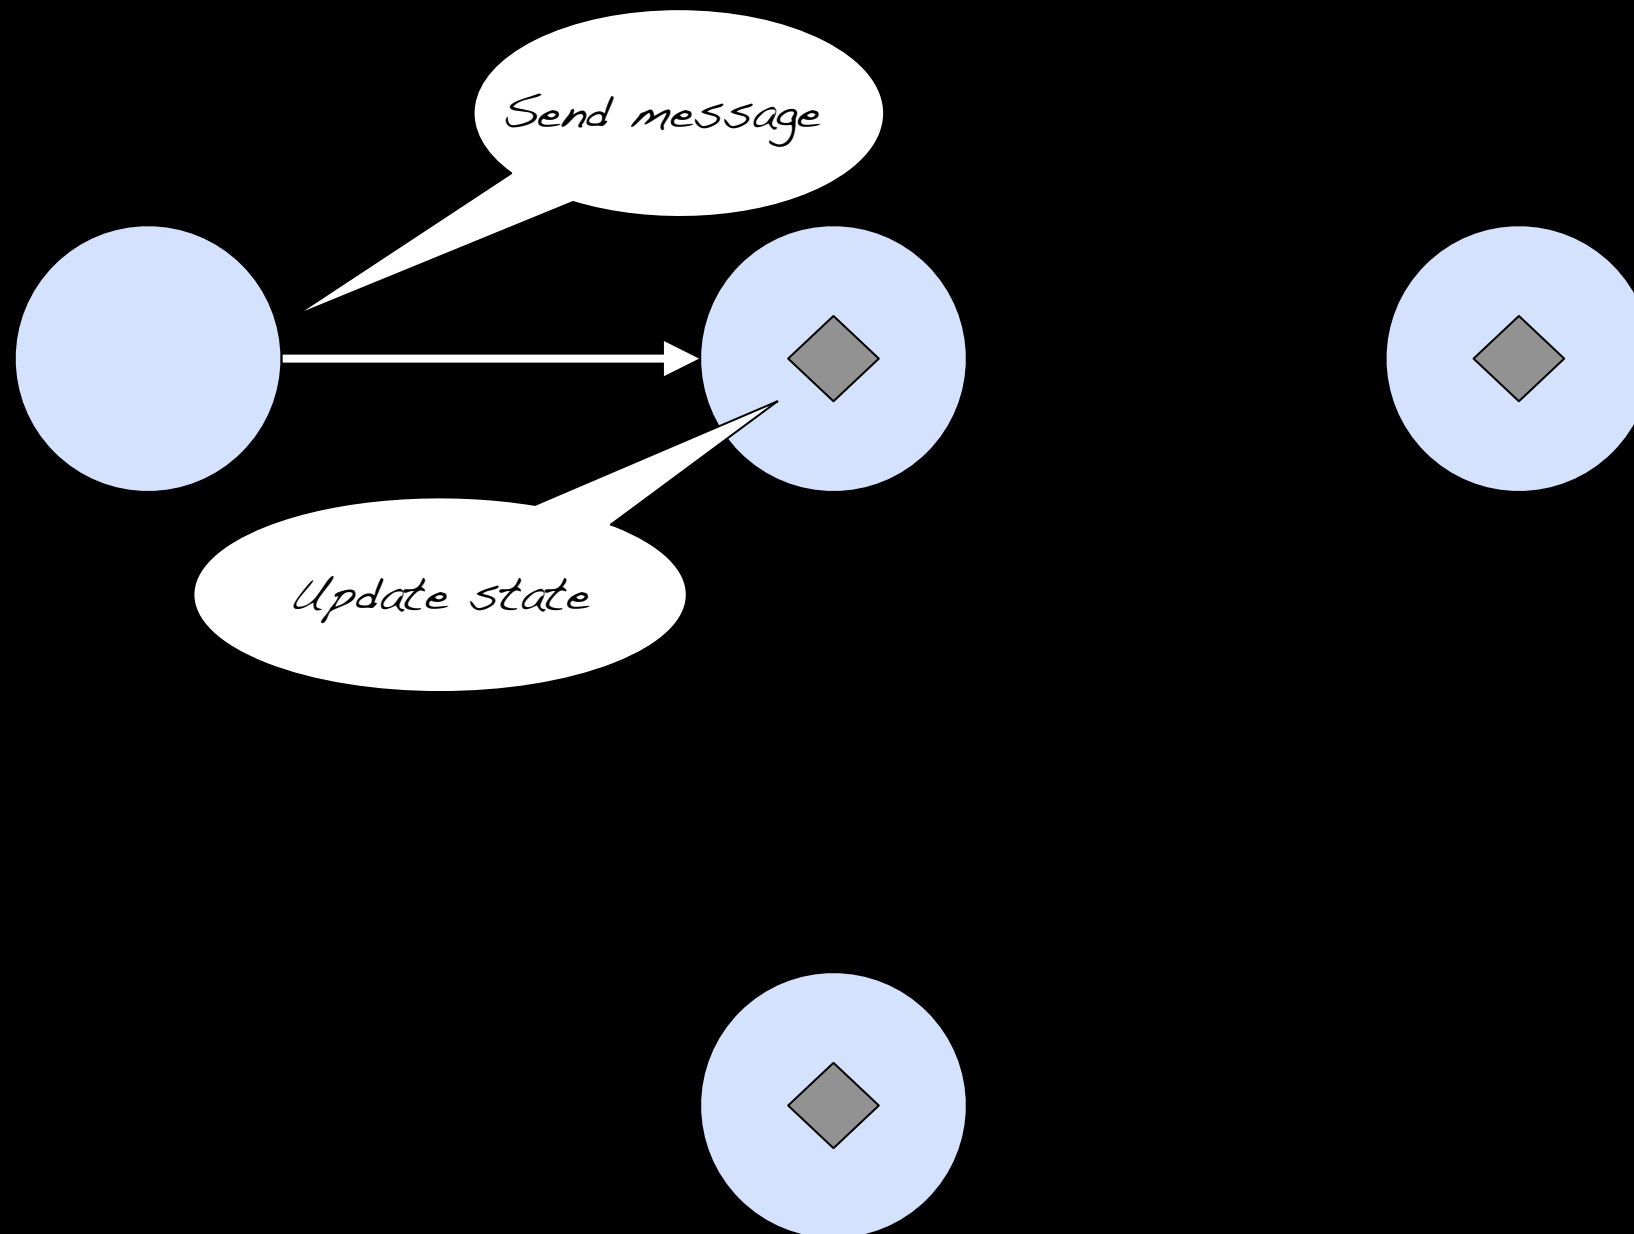
Actors



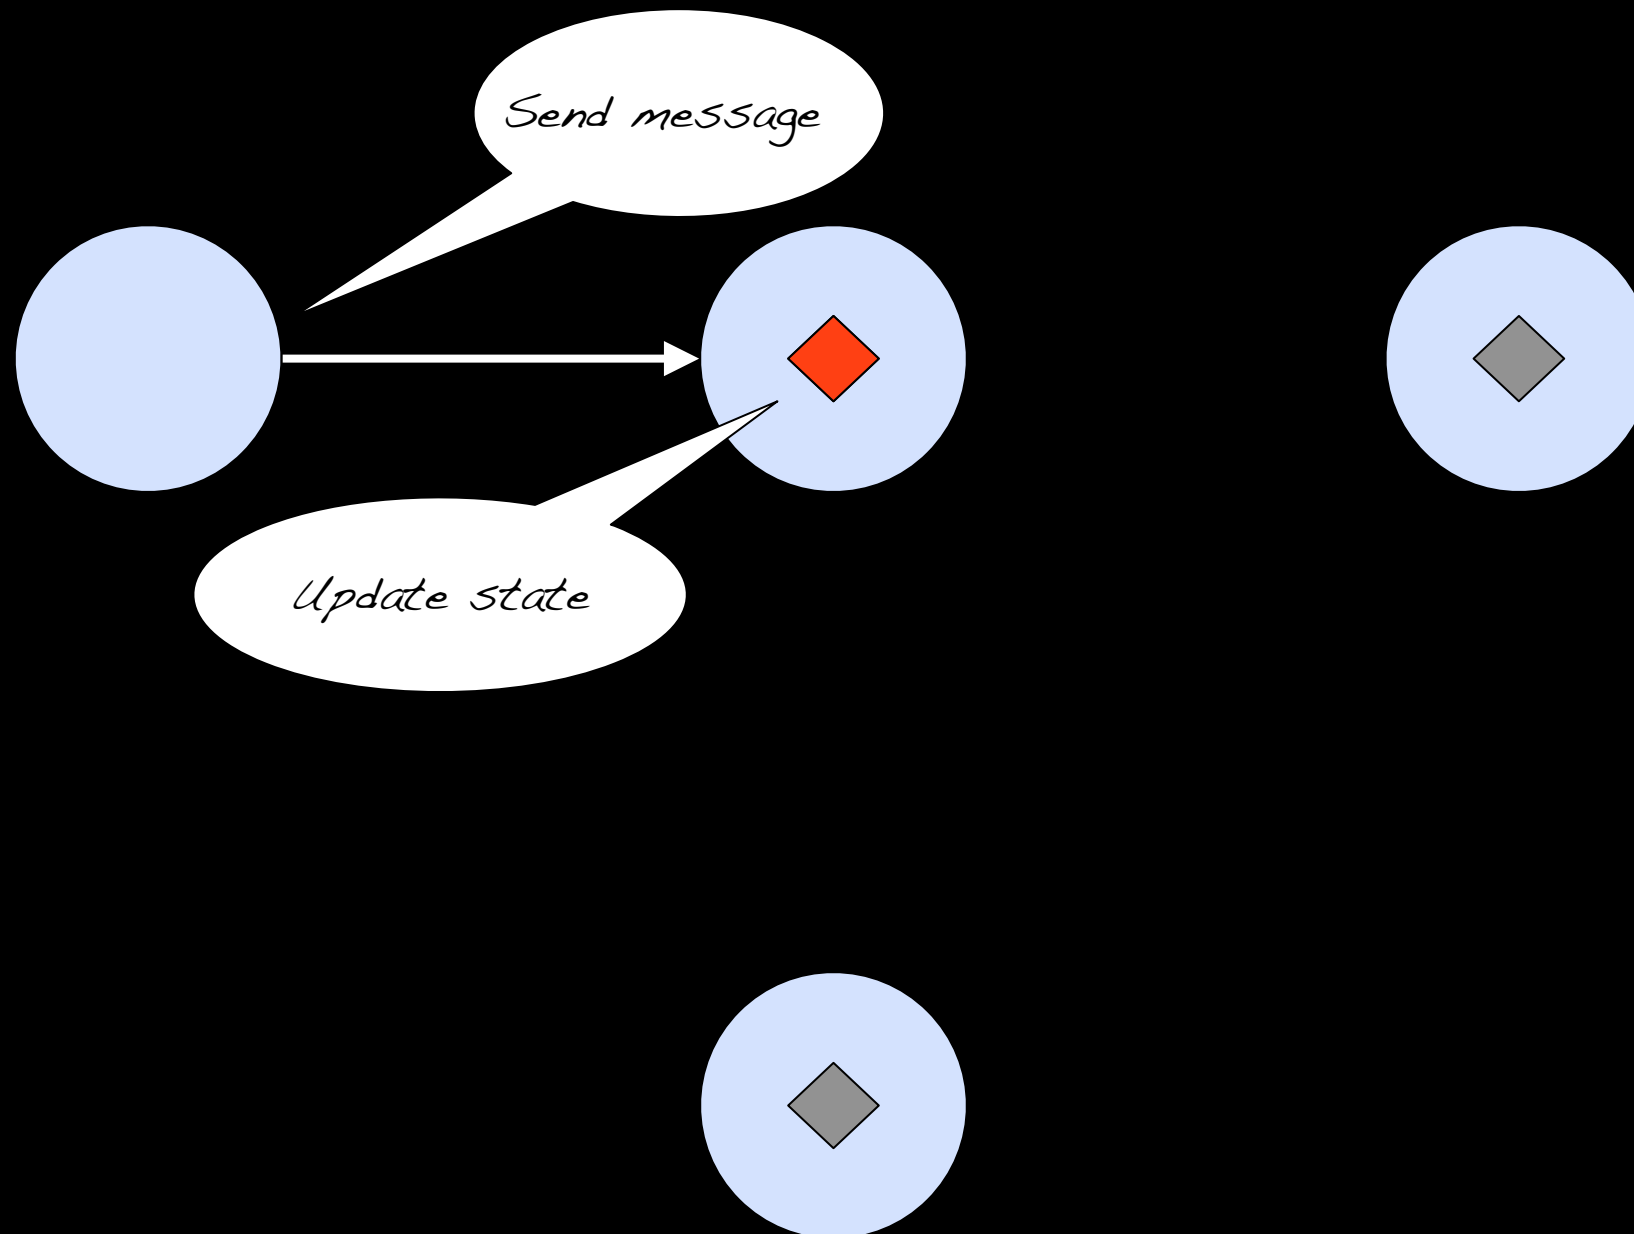
Actors



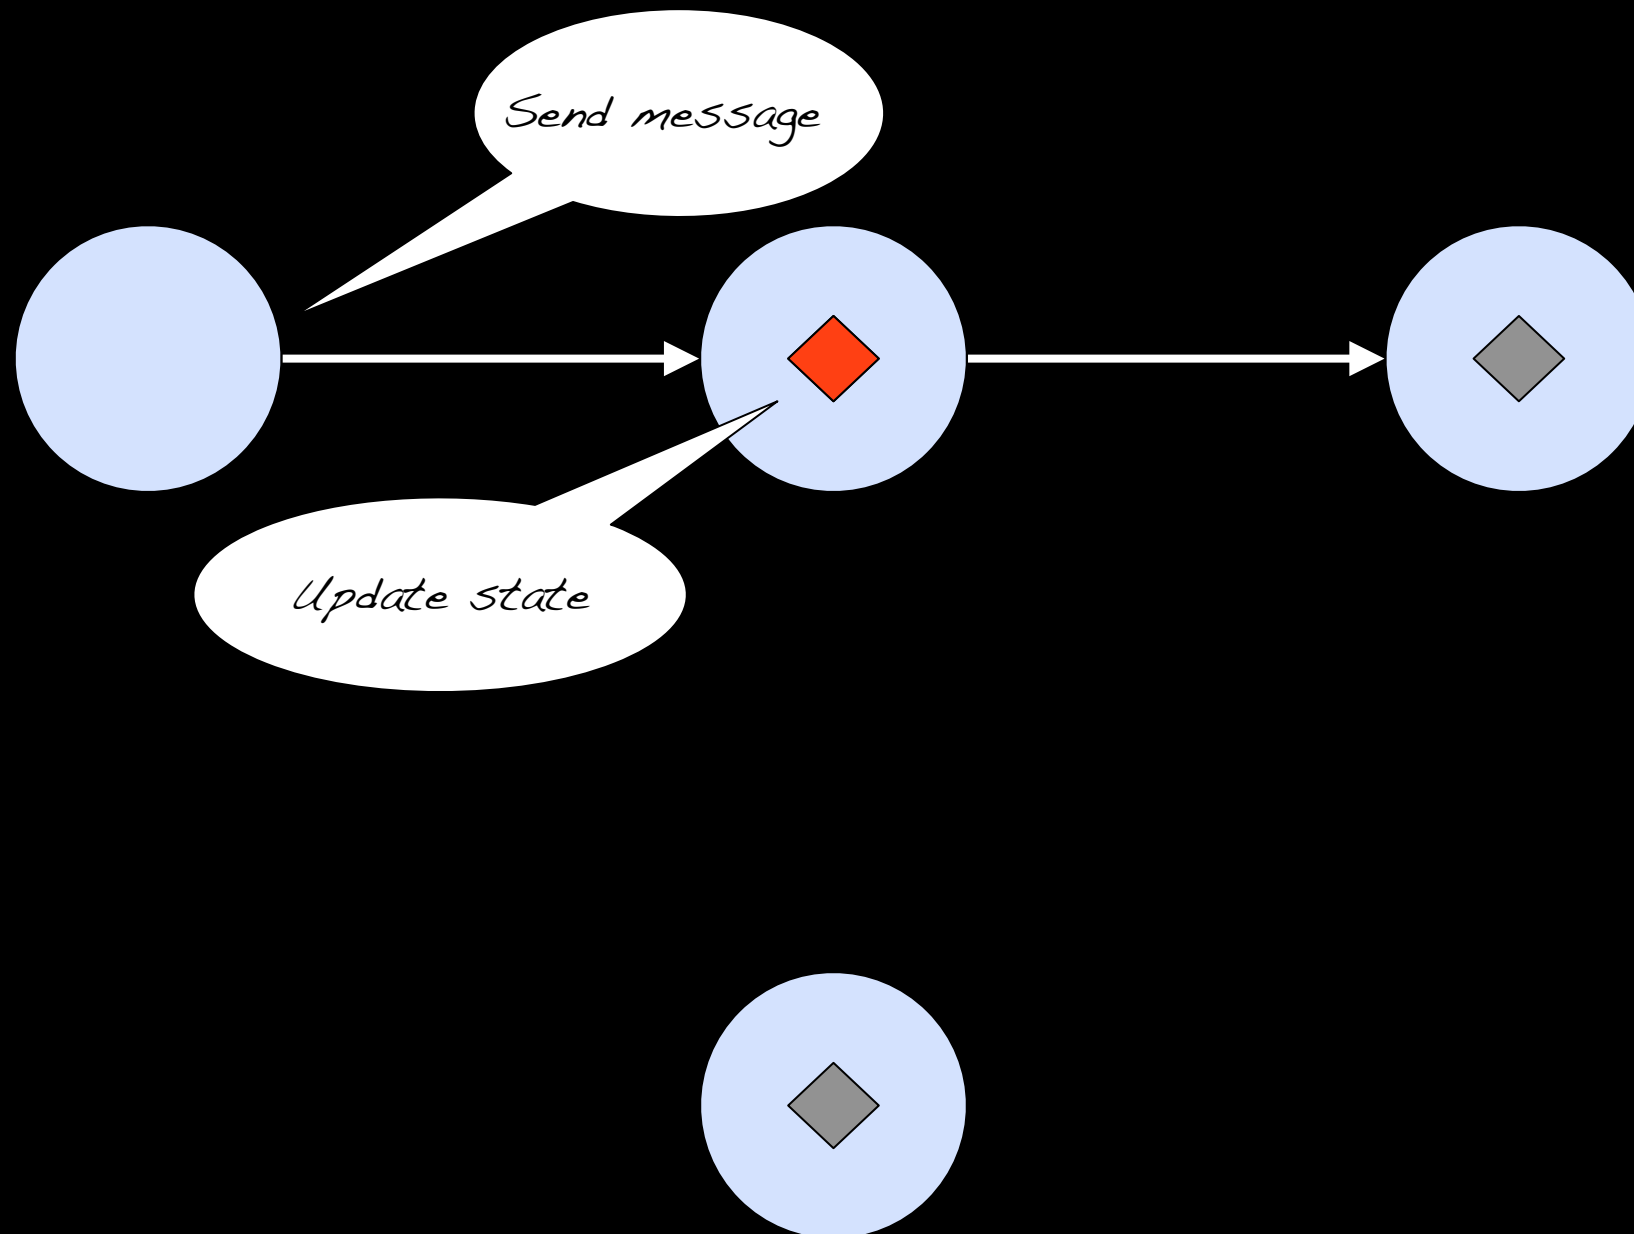
Actors



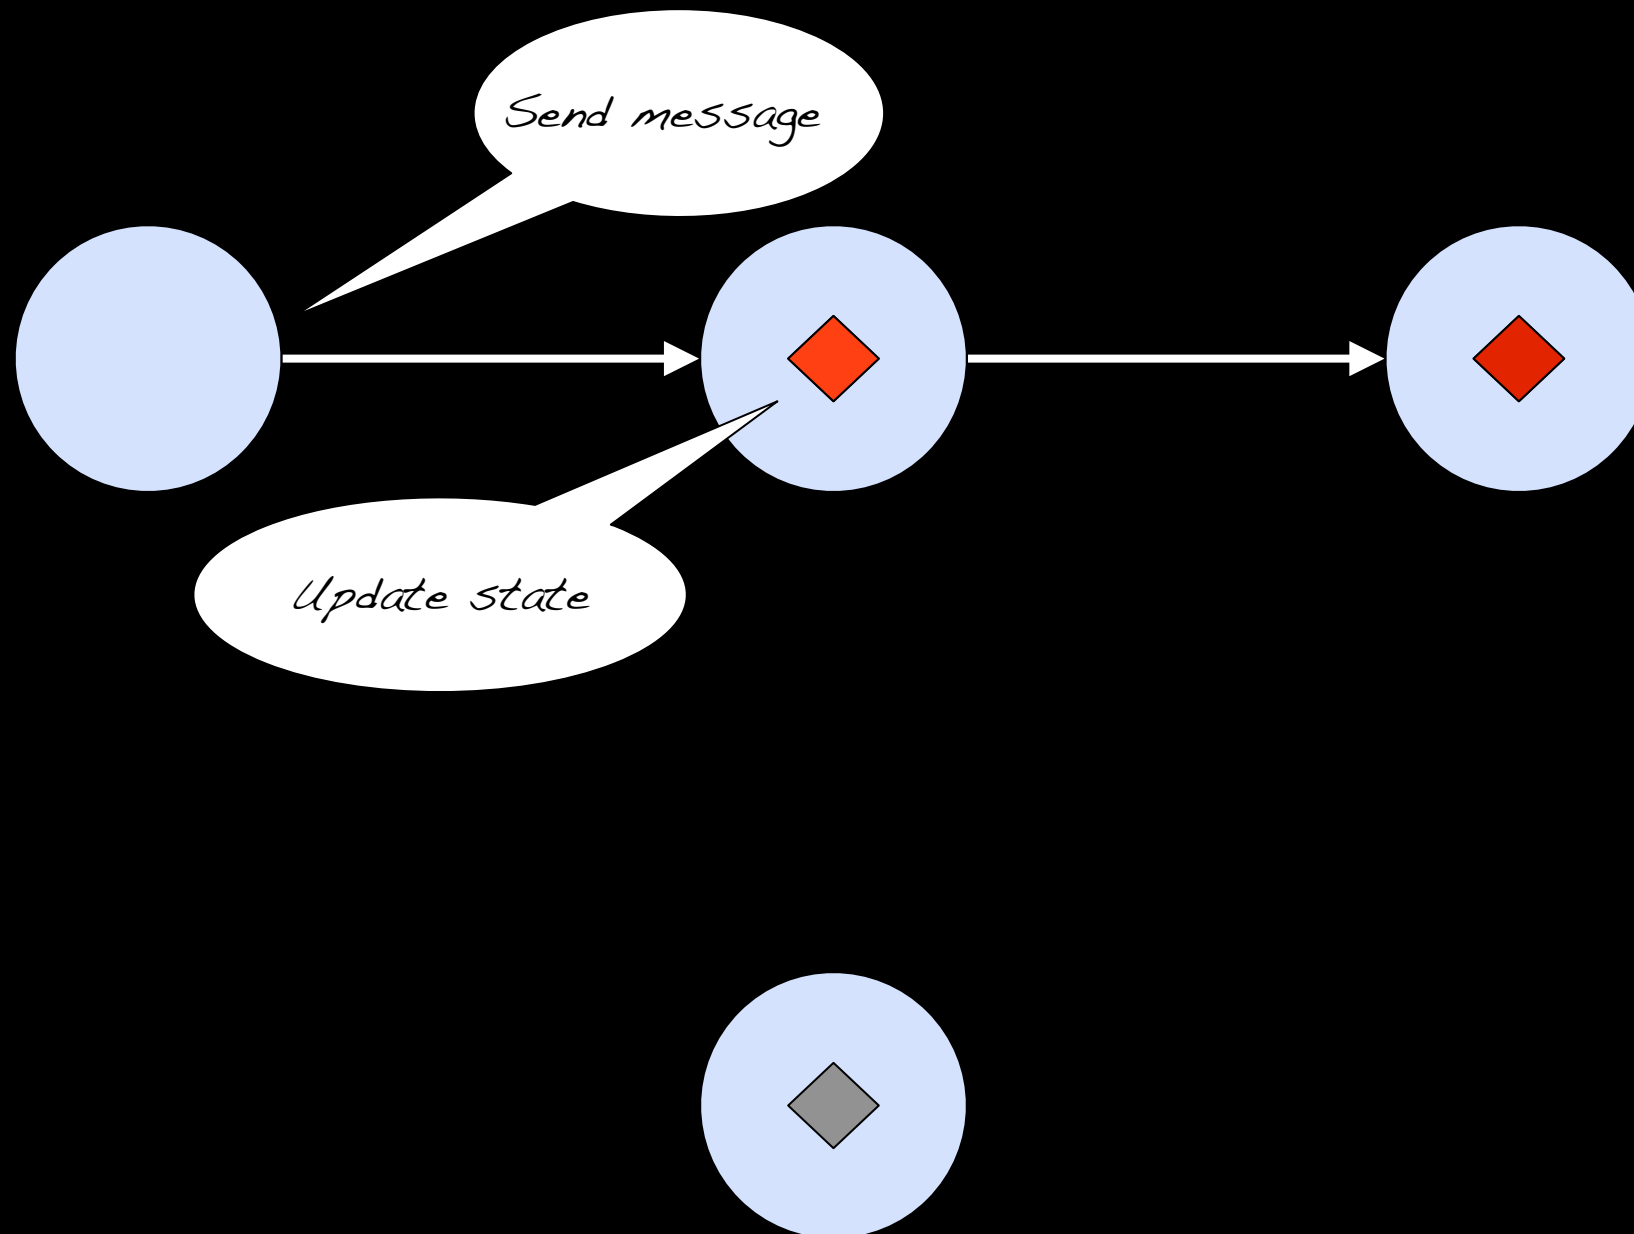
Actors



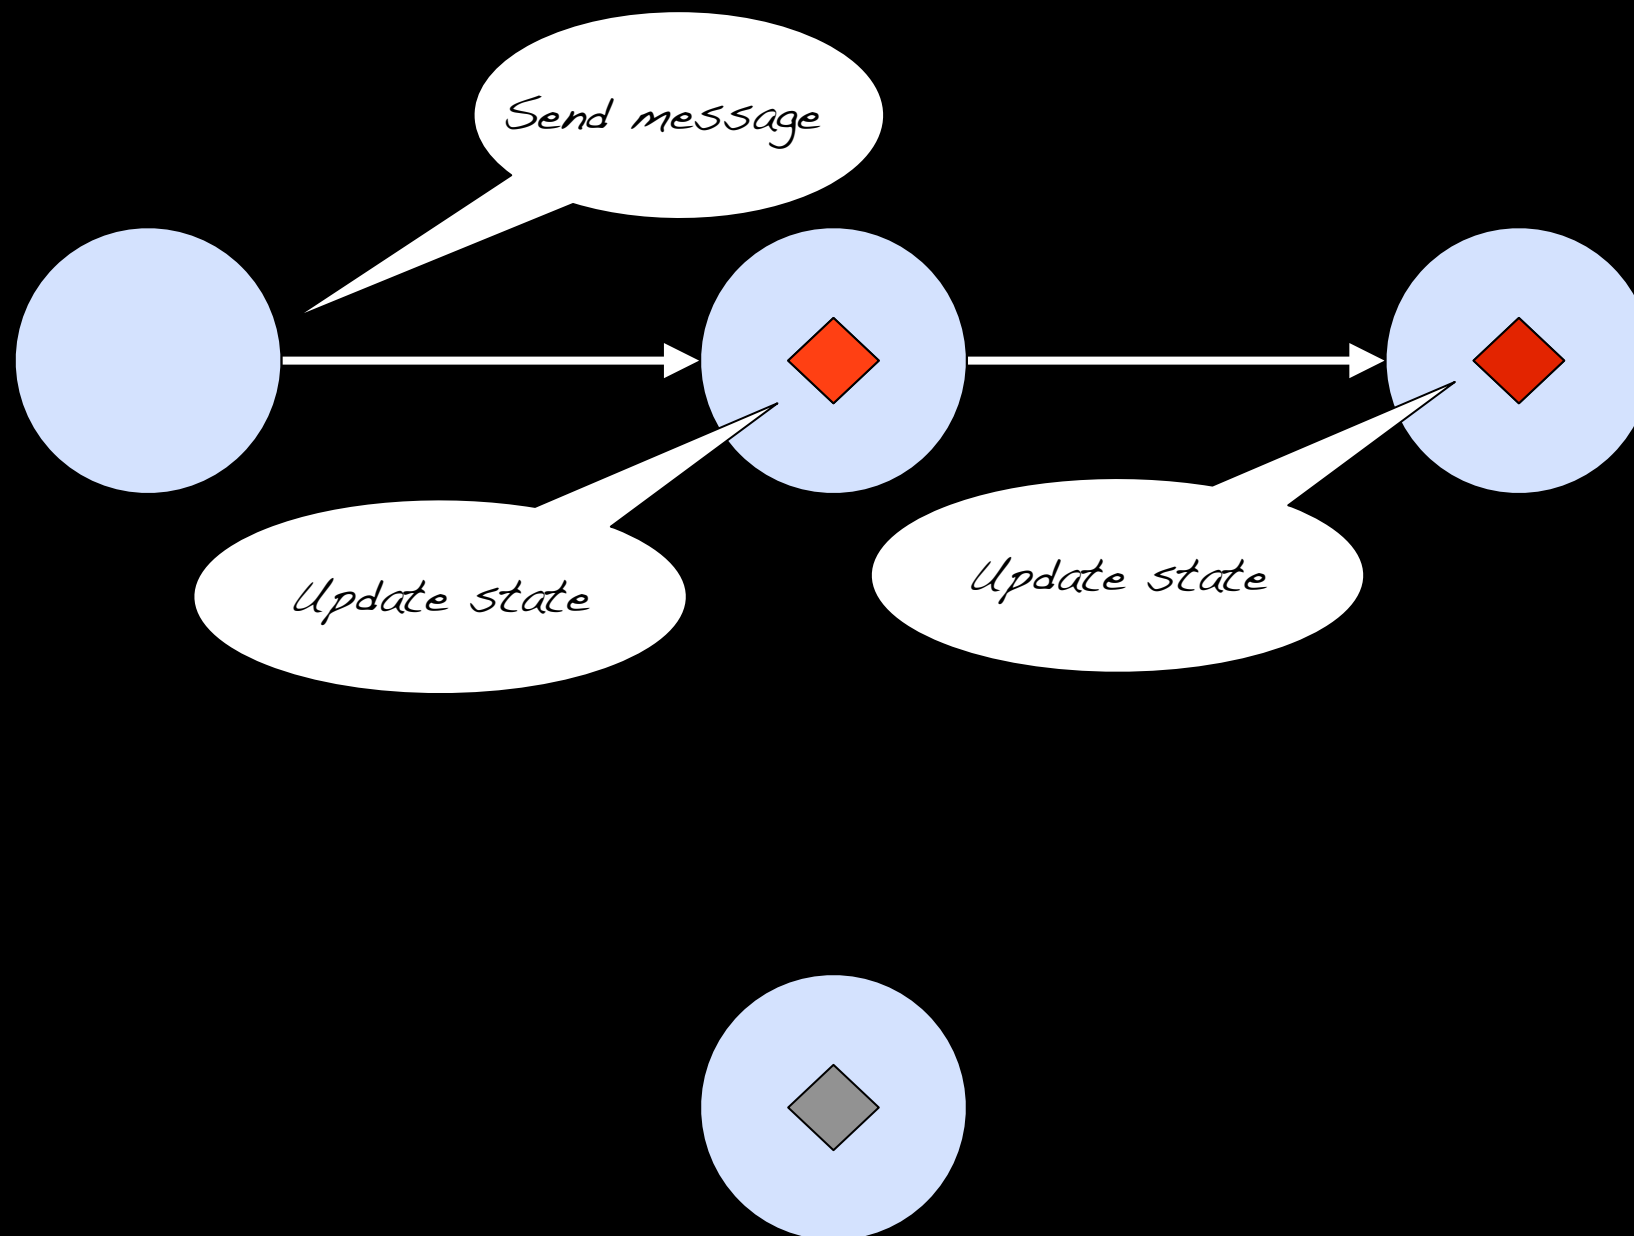
Actors



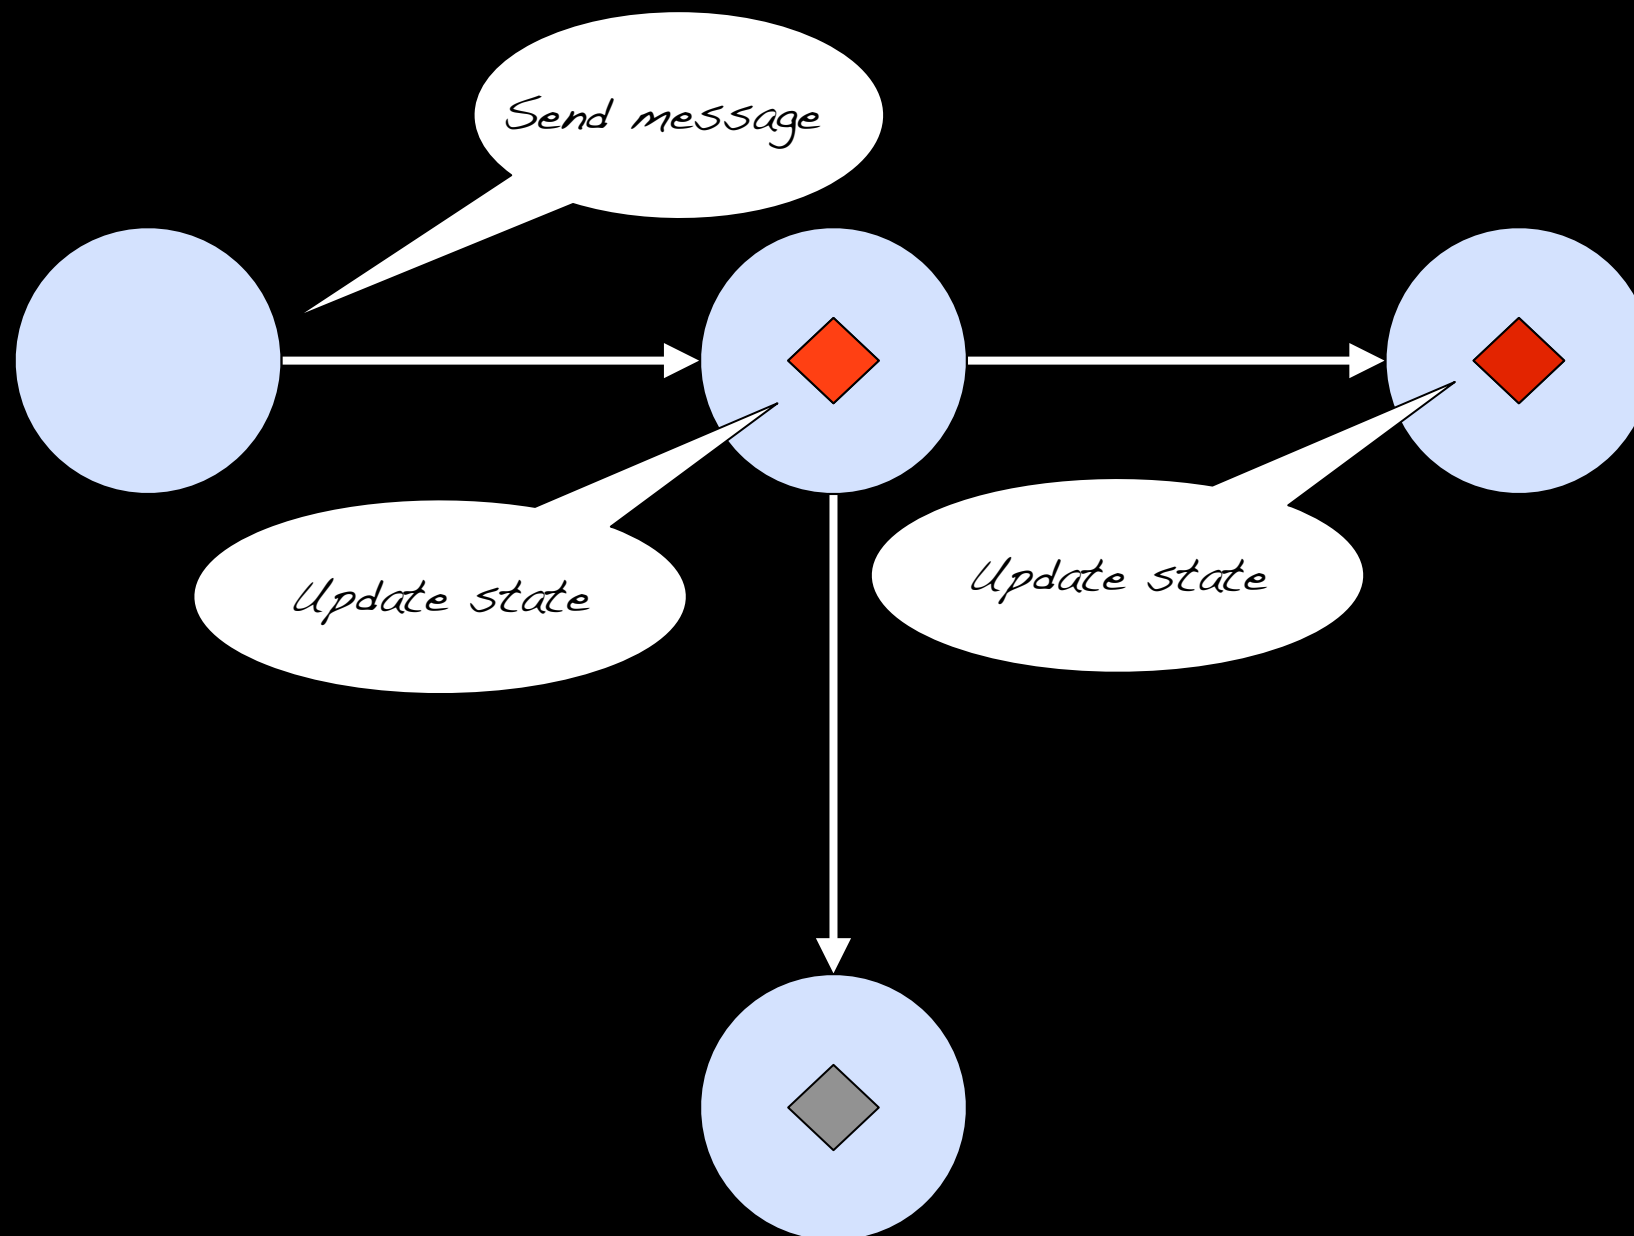
Actors



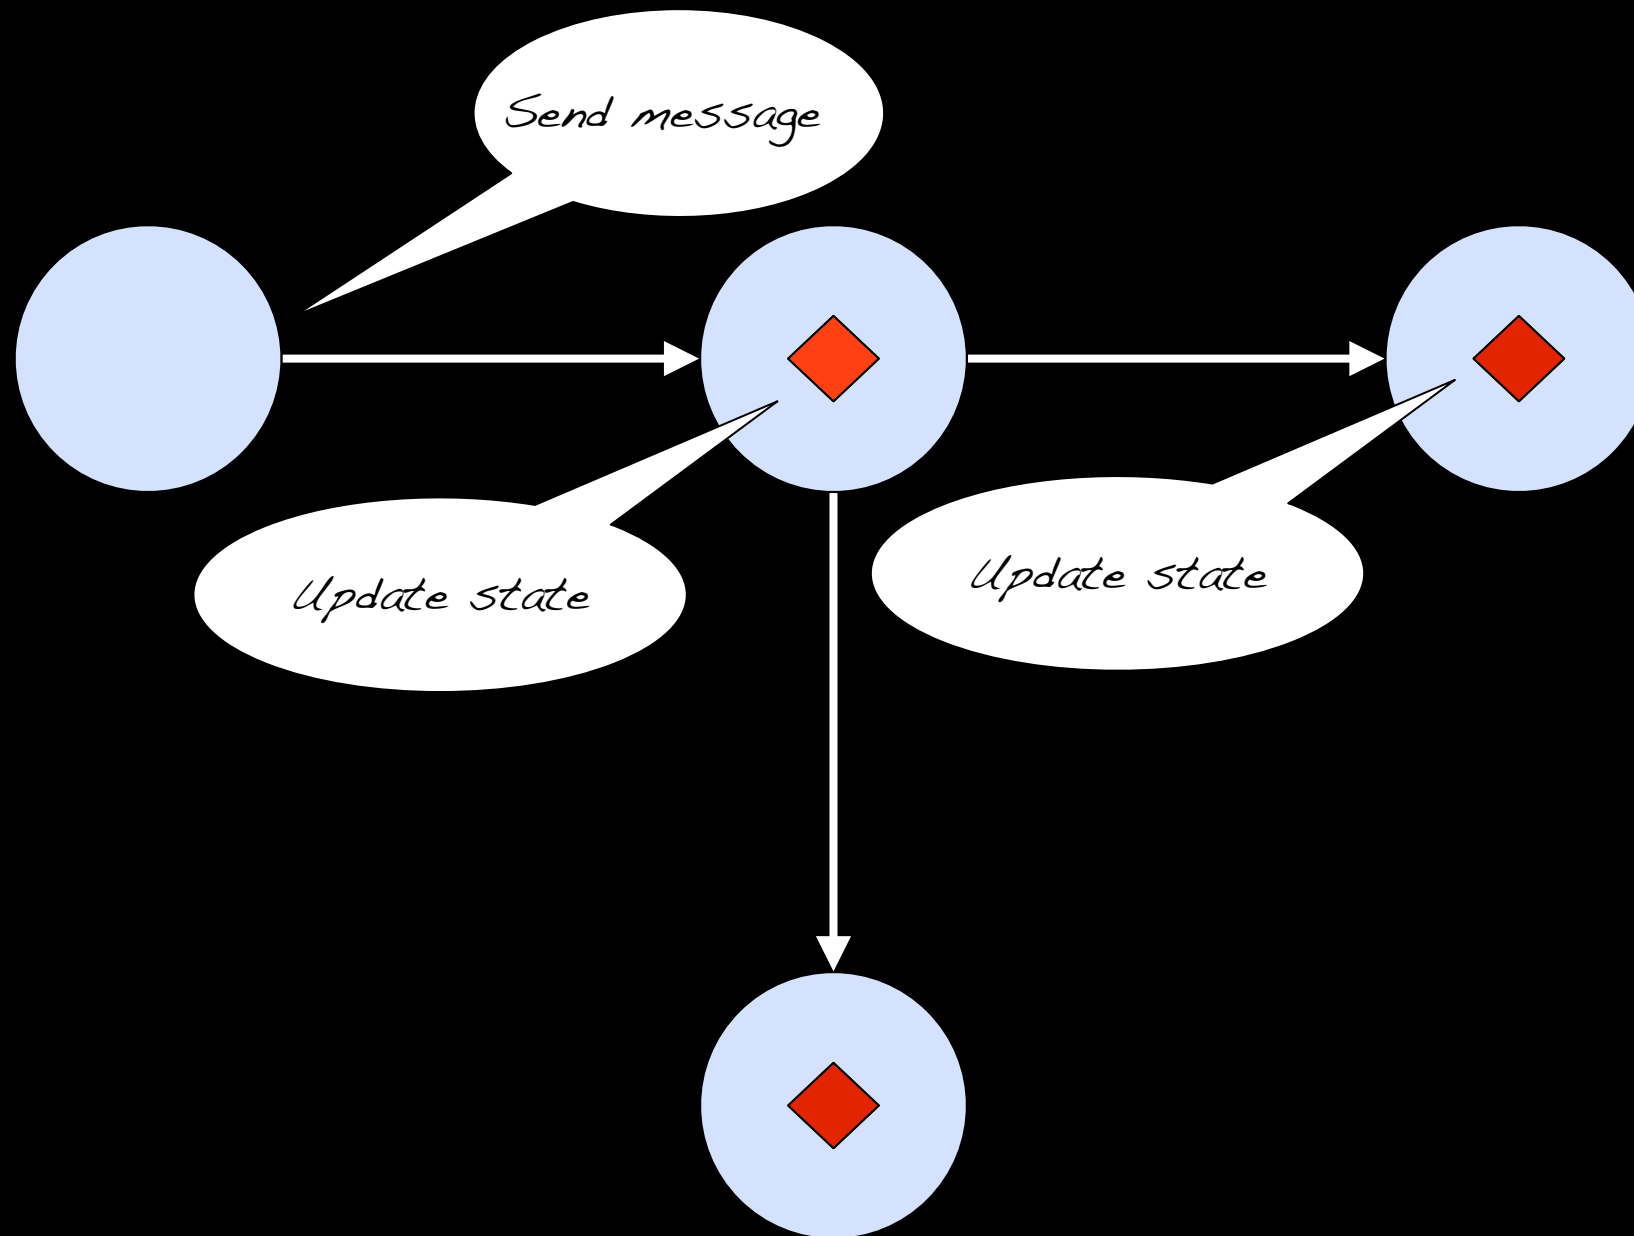
Actors



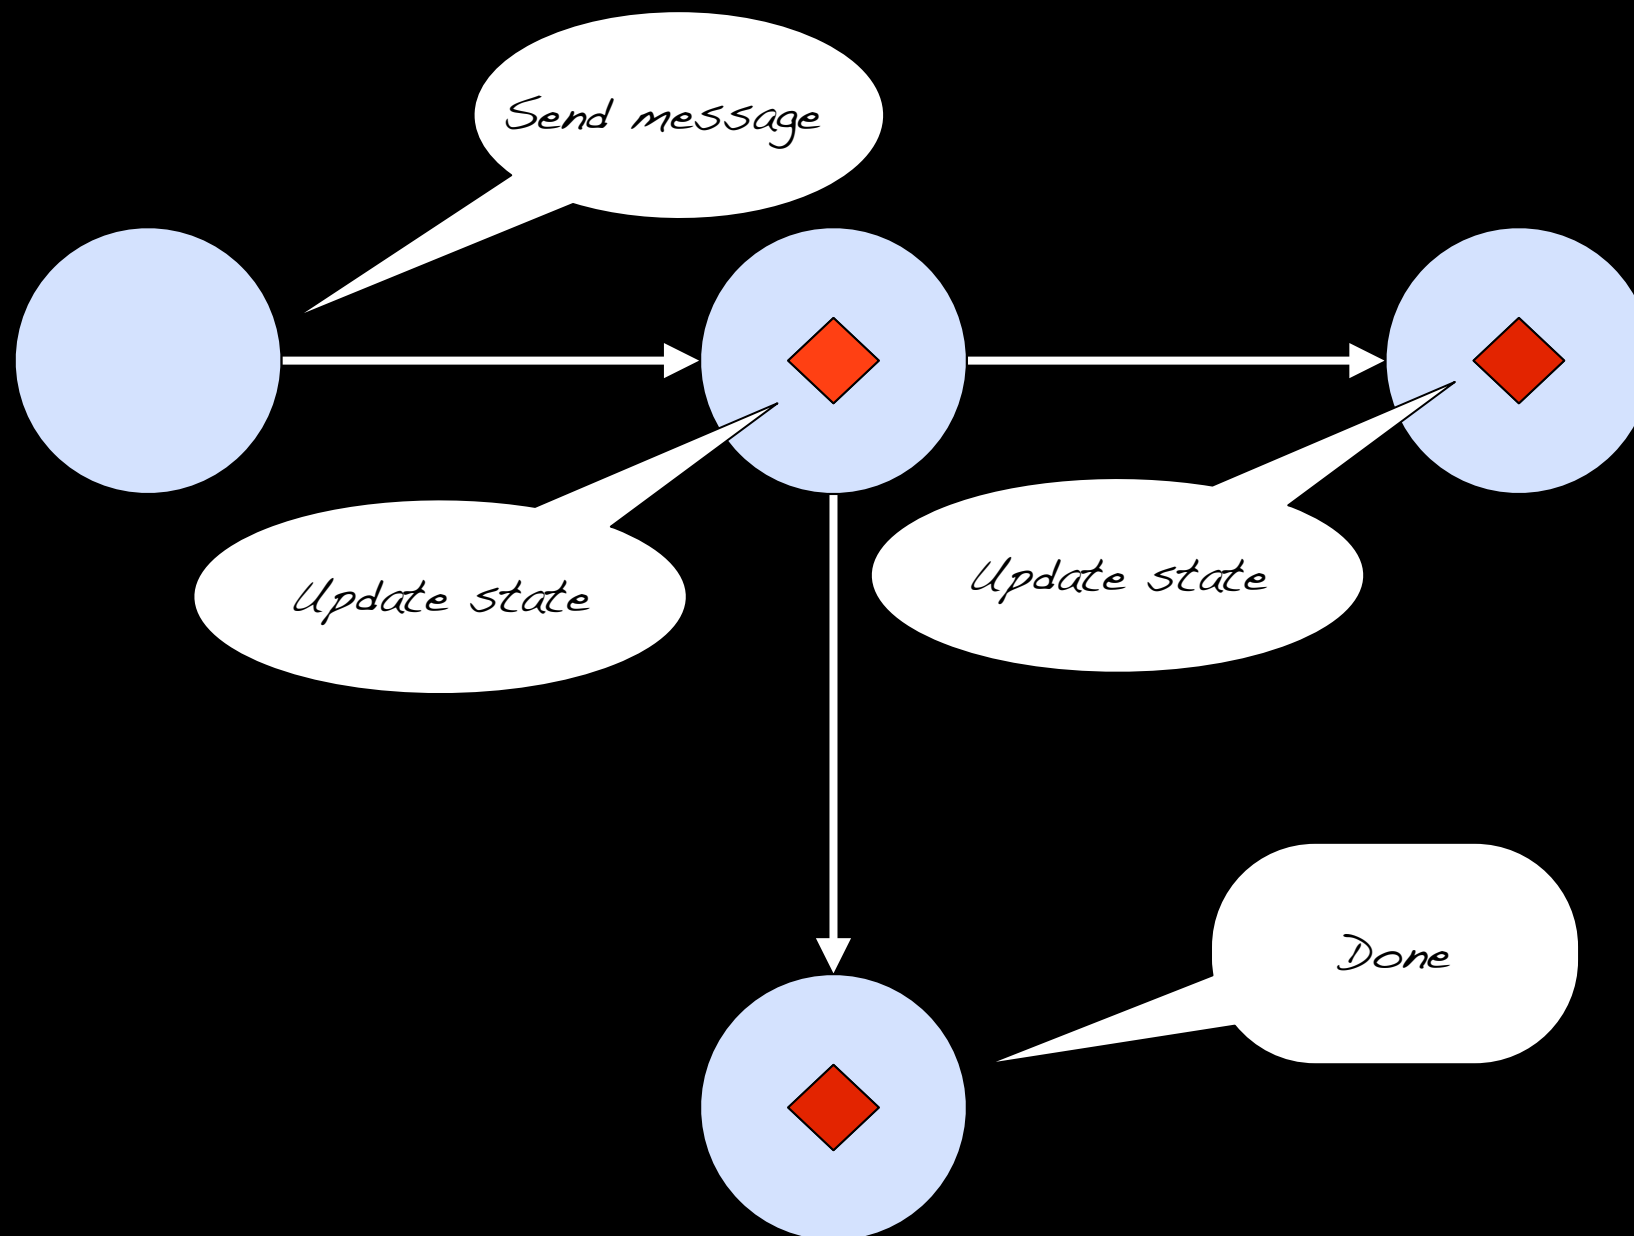
Actors



Actors



Actors



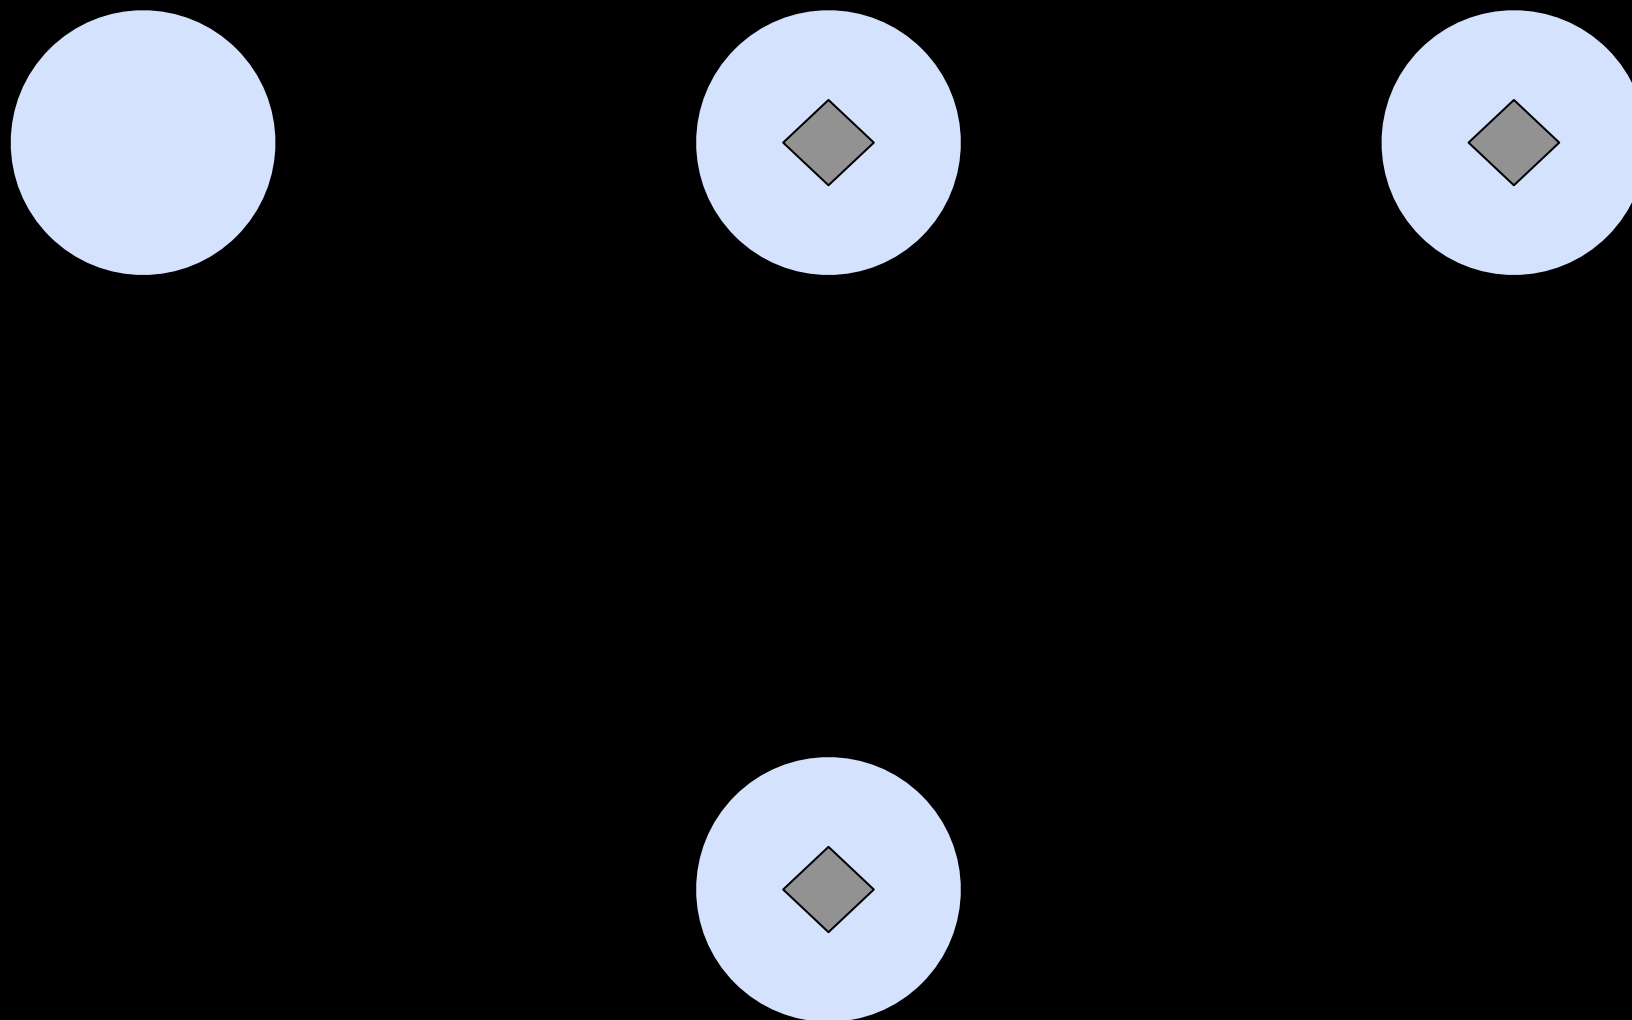
REPL time!

Actor gotchas

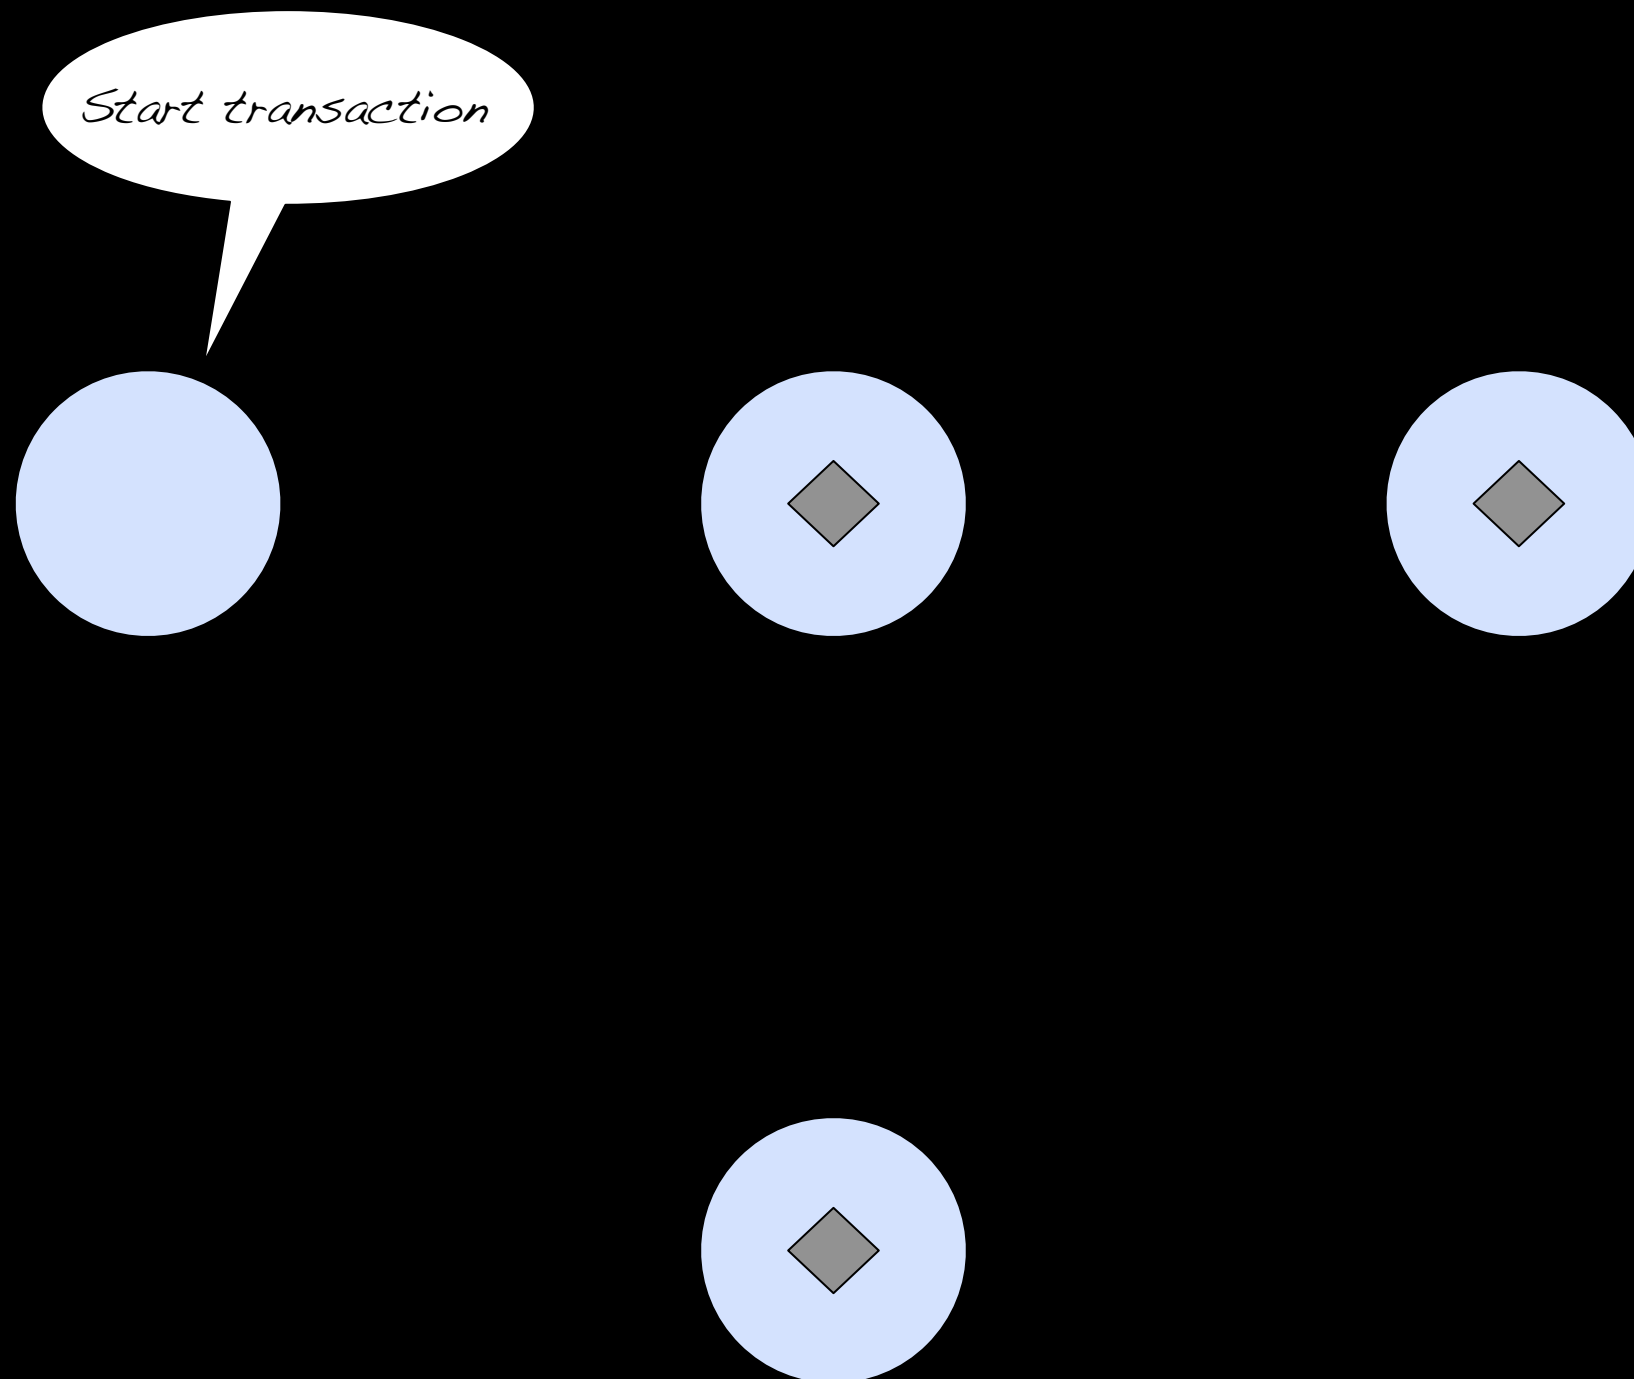
- Messages should be **immutable**
- References to mutable internal state must not escape the Actor
- Coordinating state changes between multiple actors is **hard**
- **Actors are reactive**, not proactive

Actors + STM =
Transactors

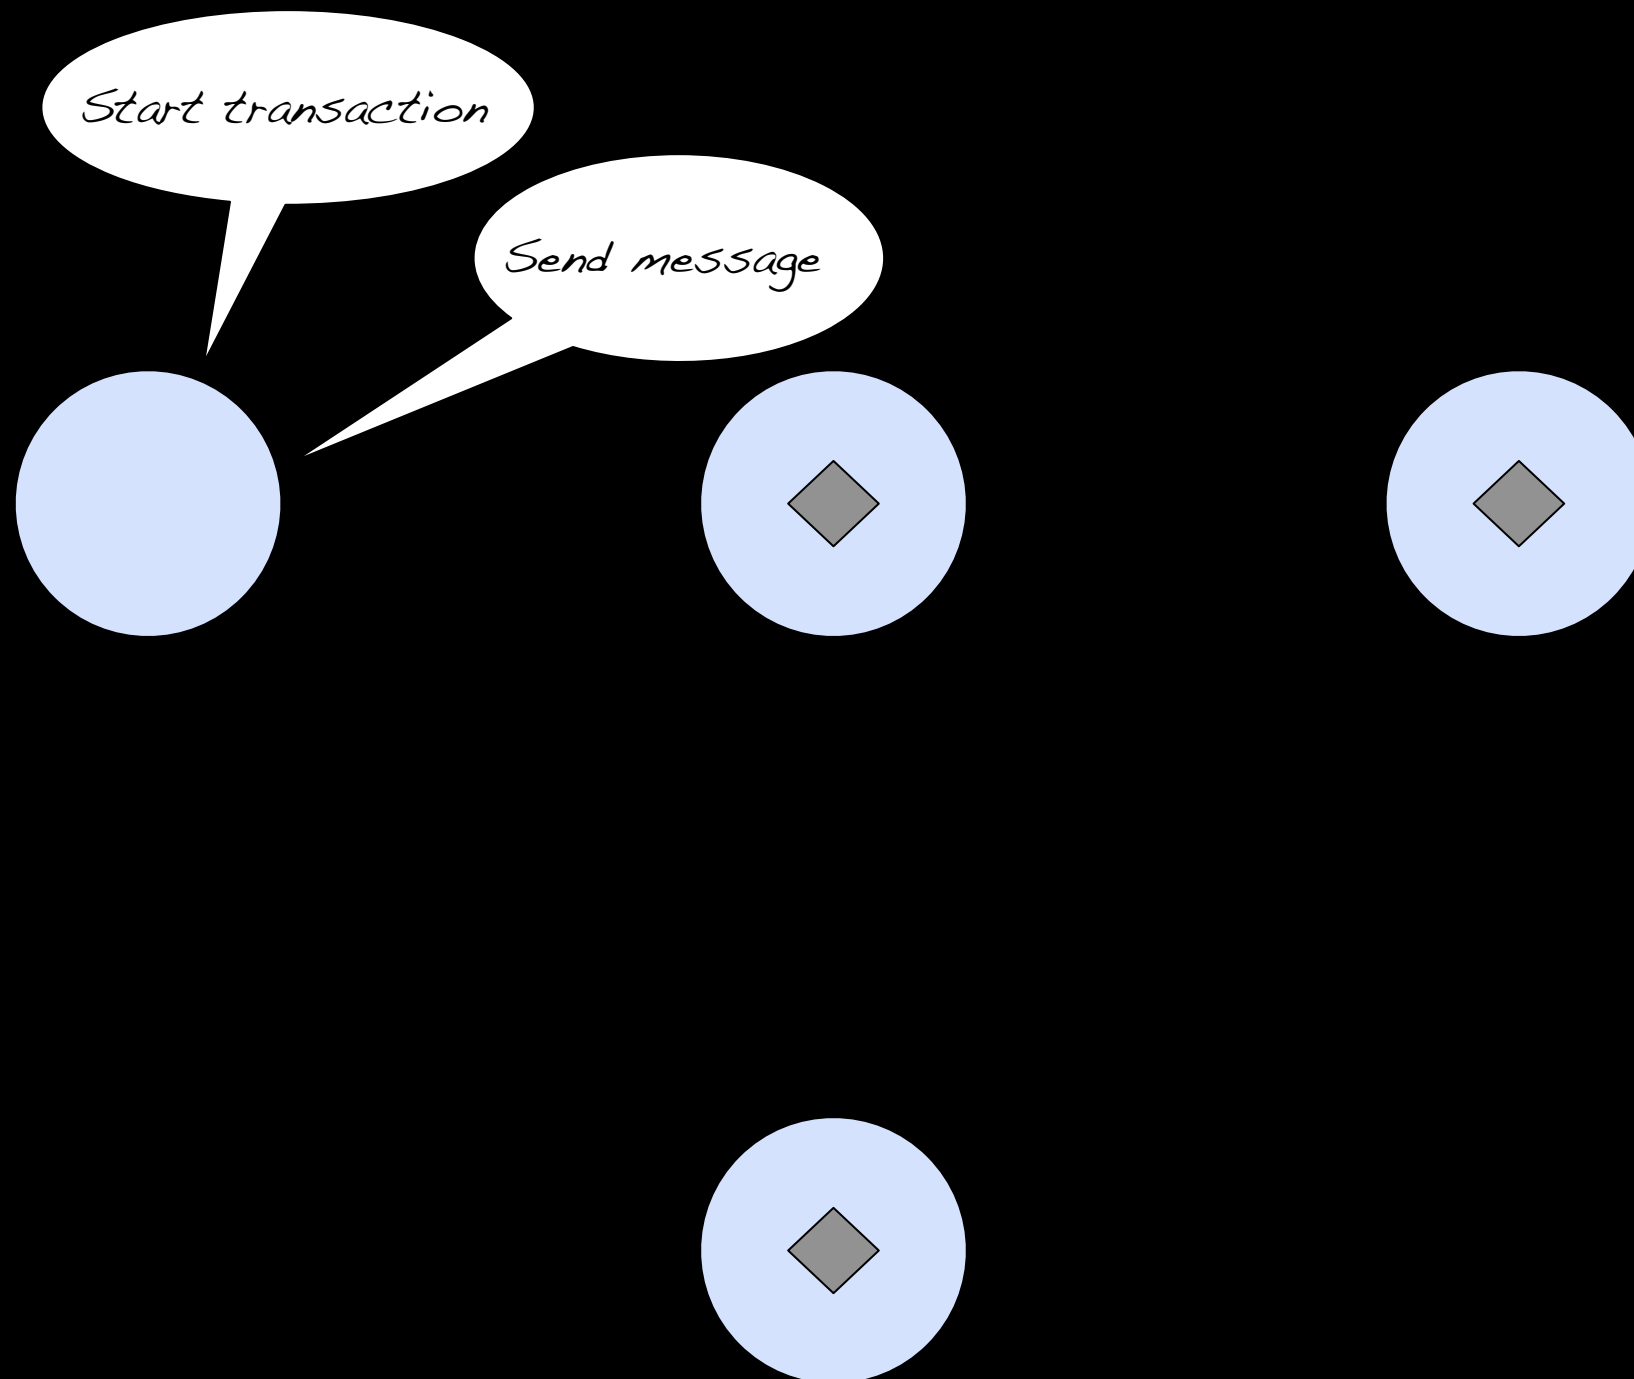
Transactors



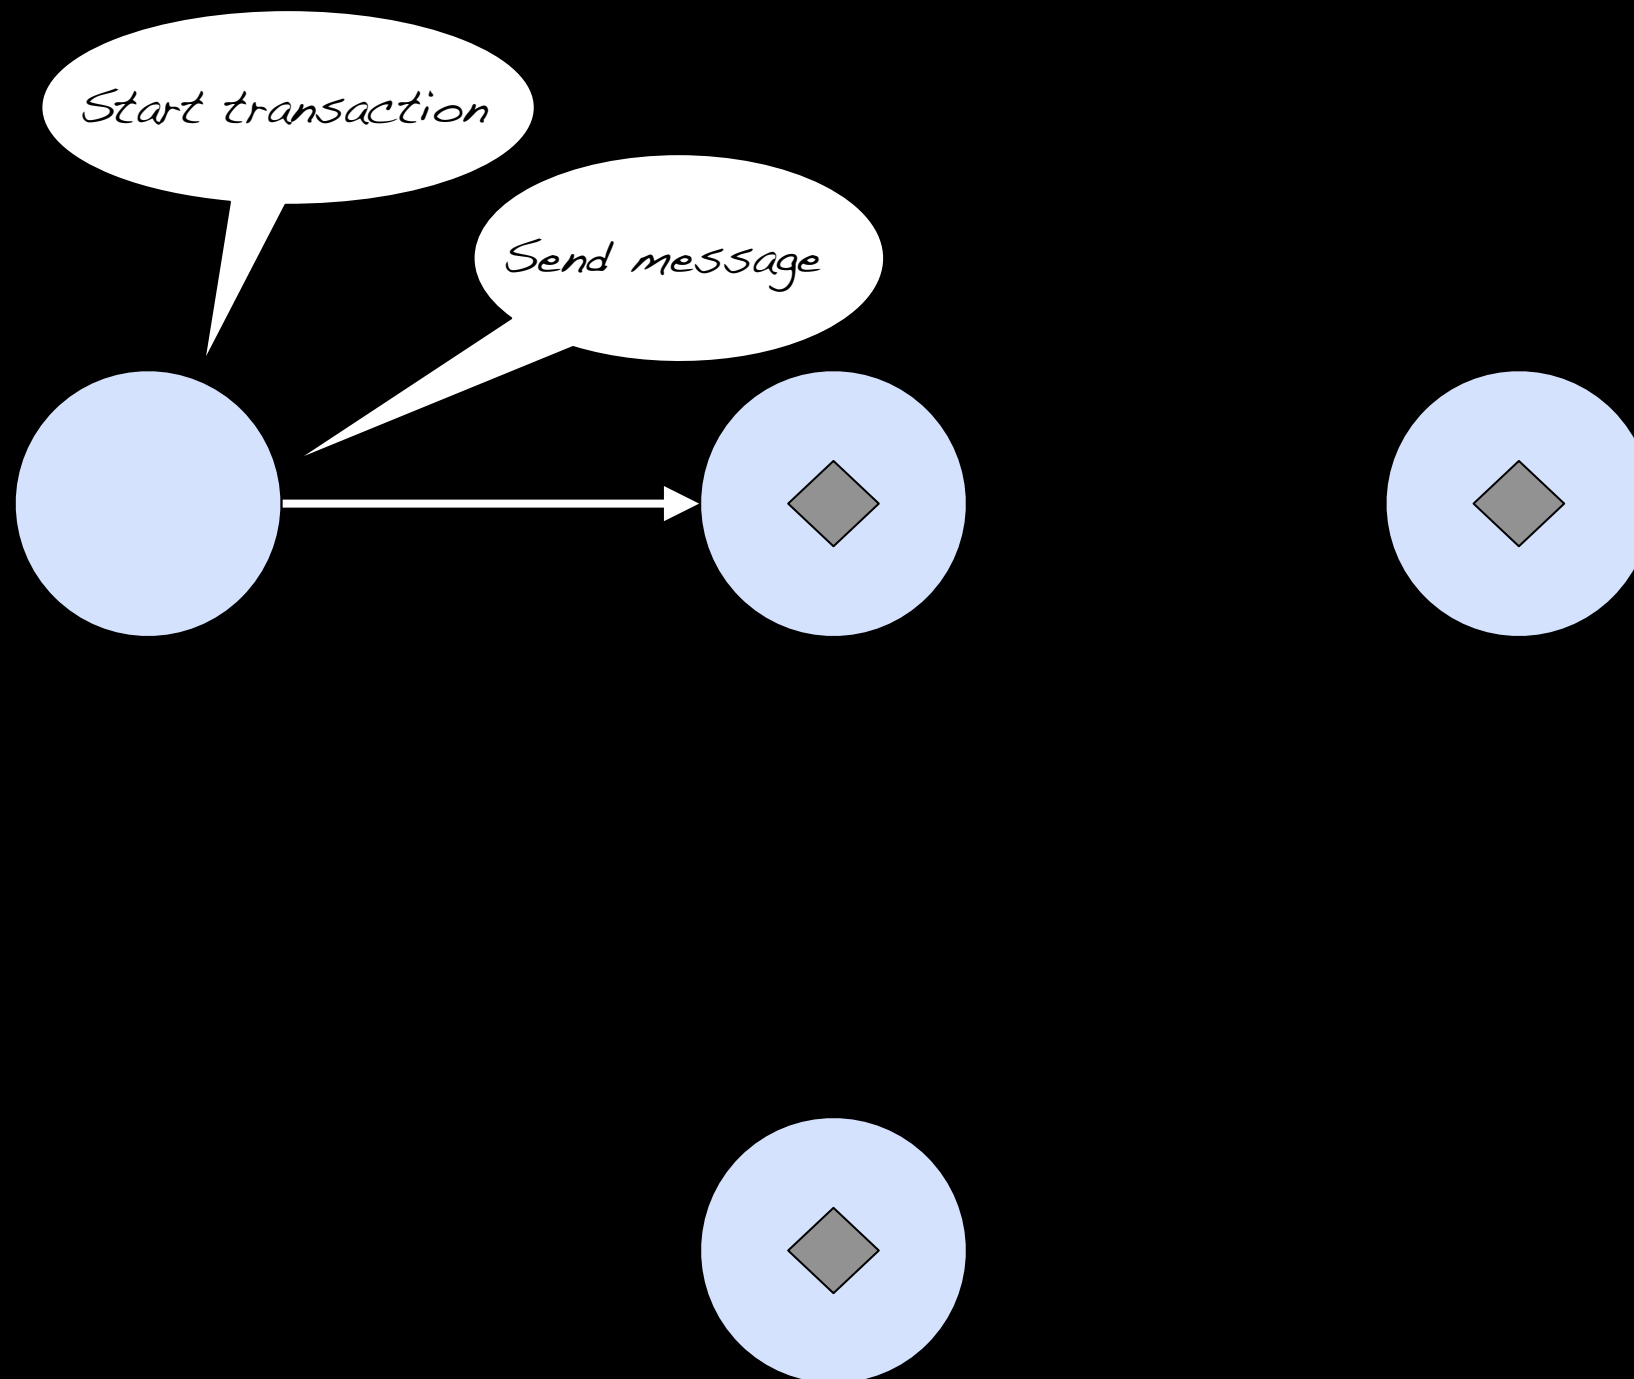
Transactors



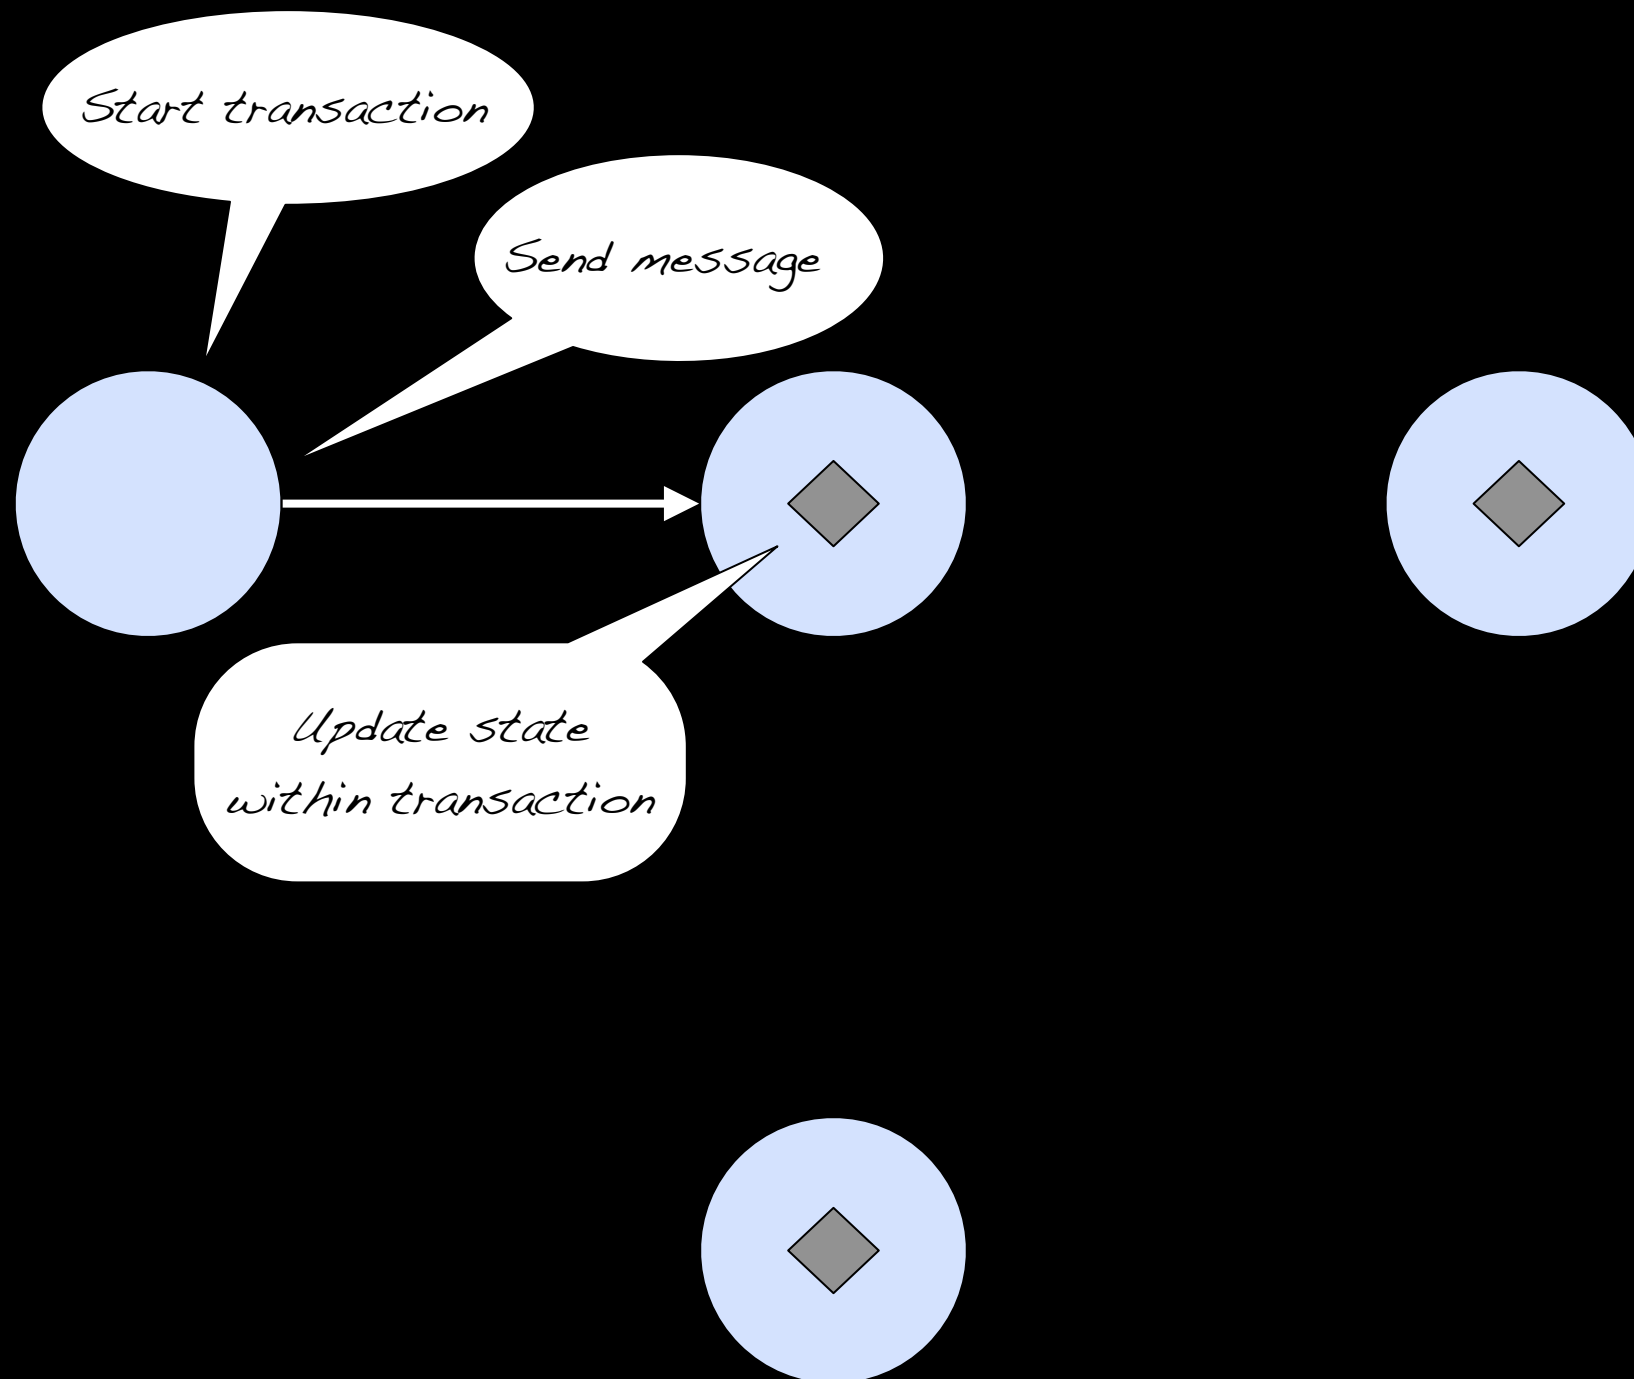
Transactors



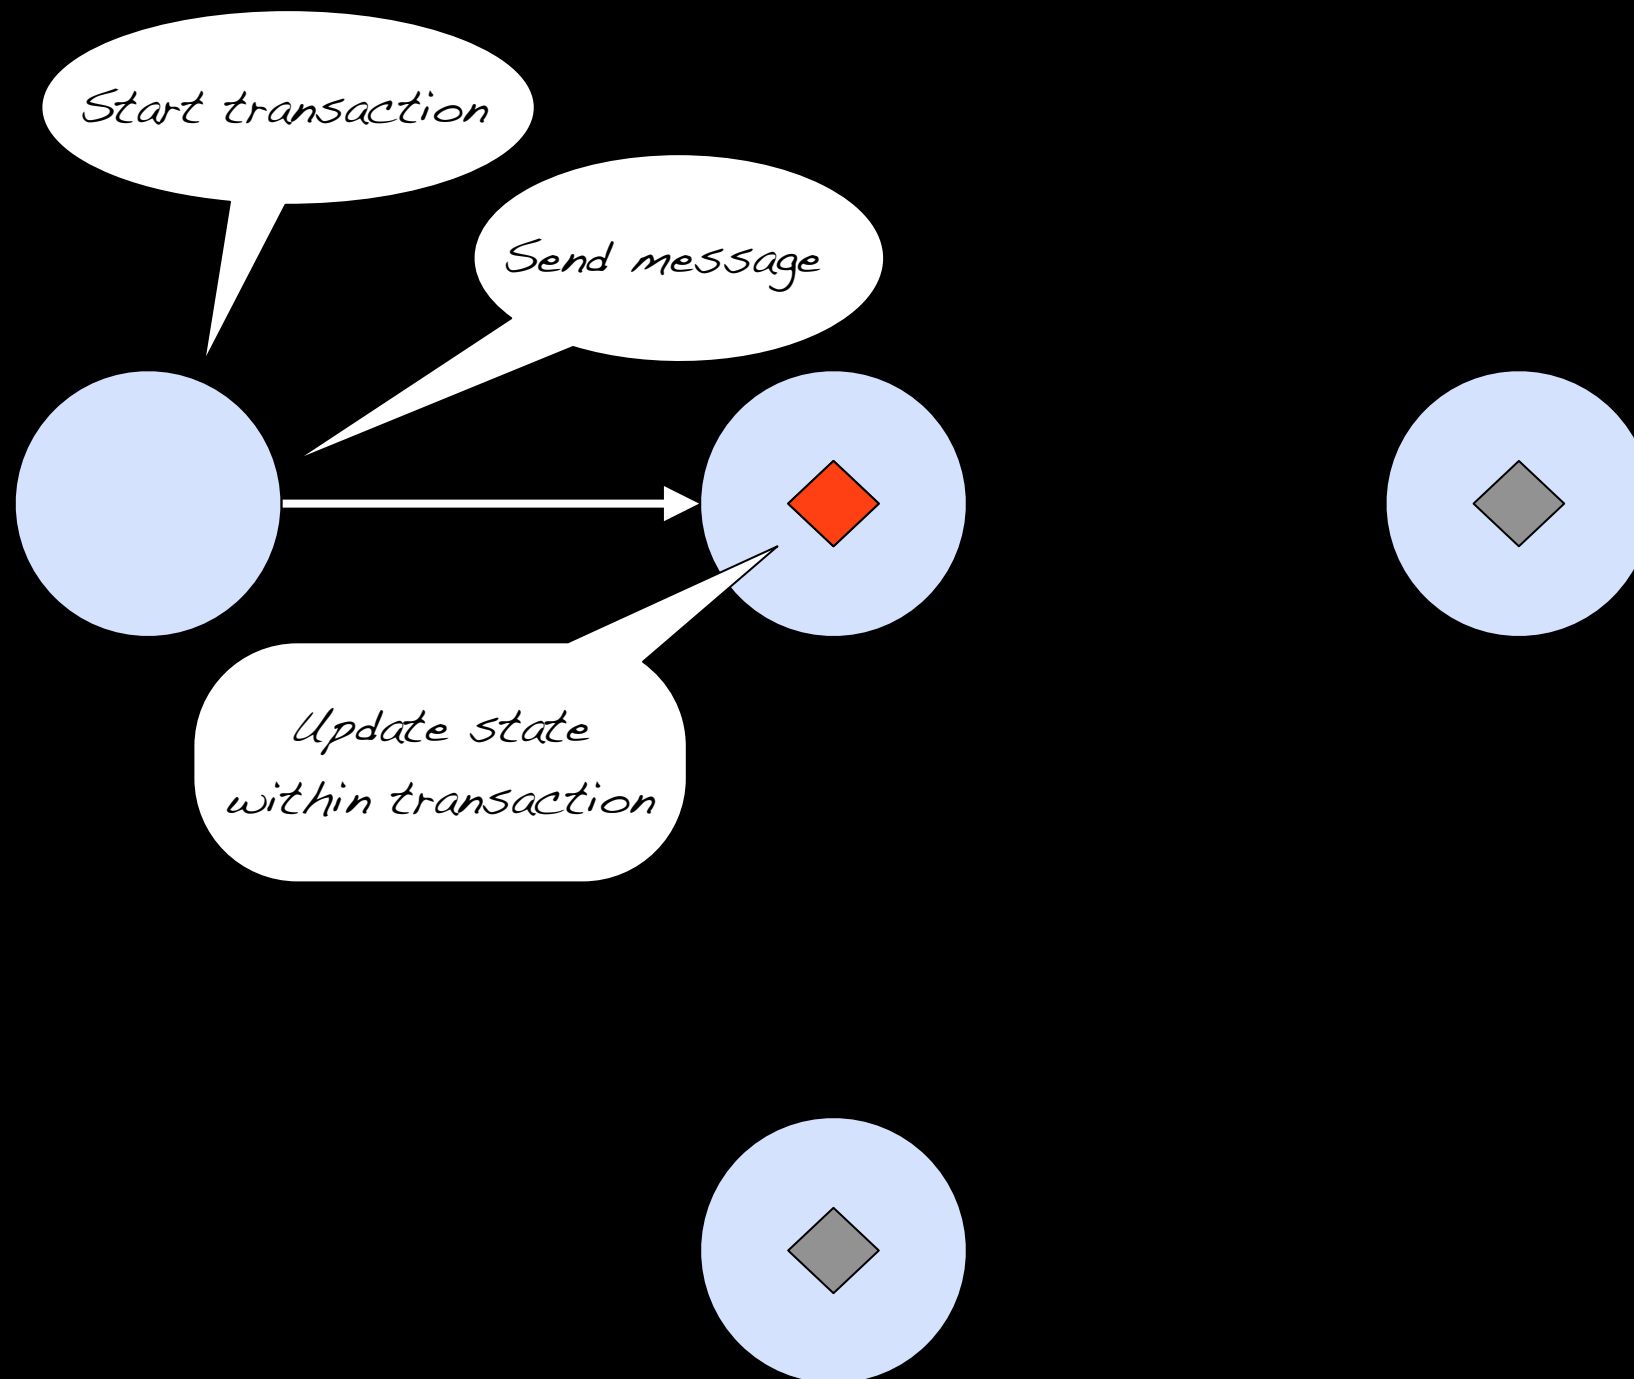
Transactors



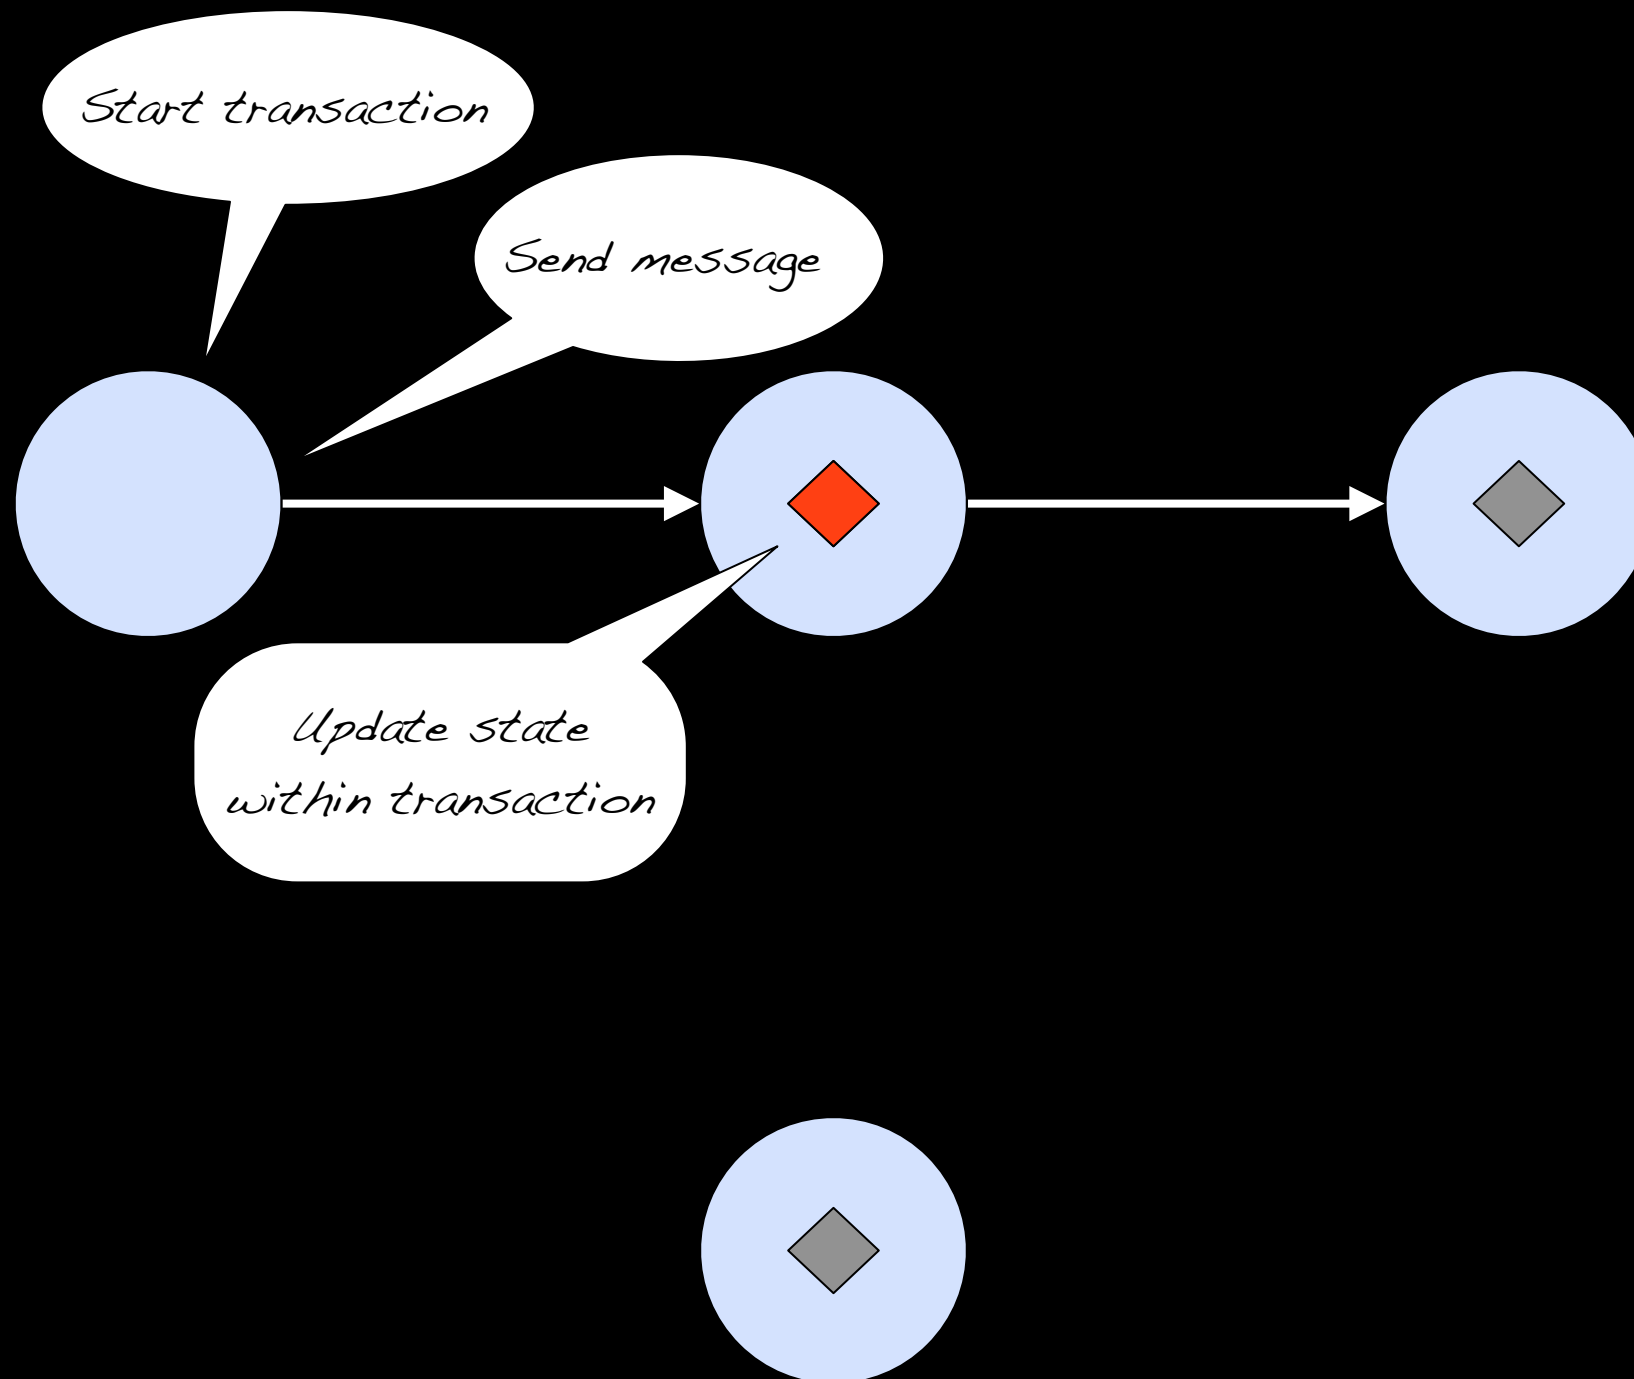
Transactors



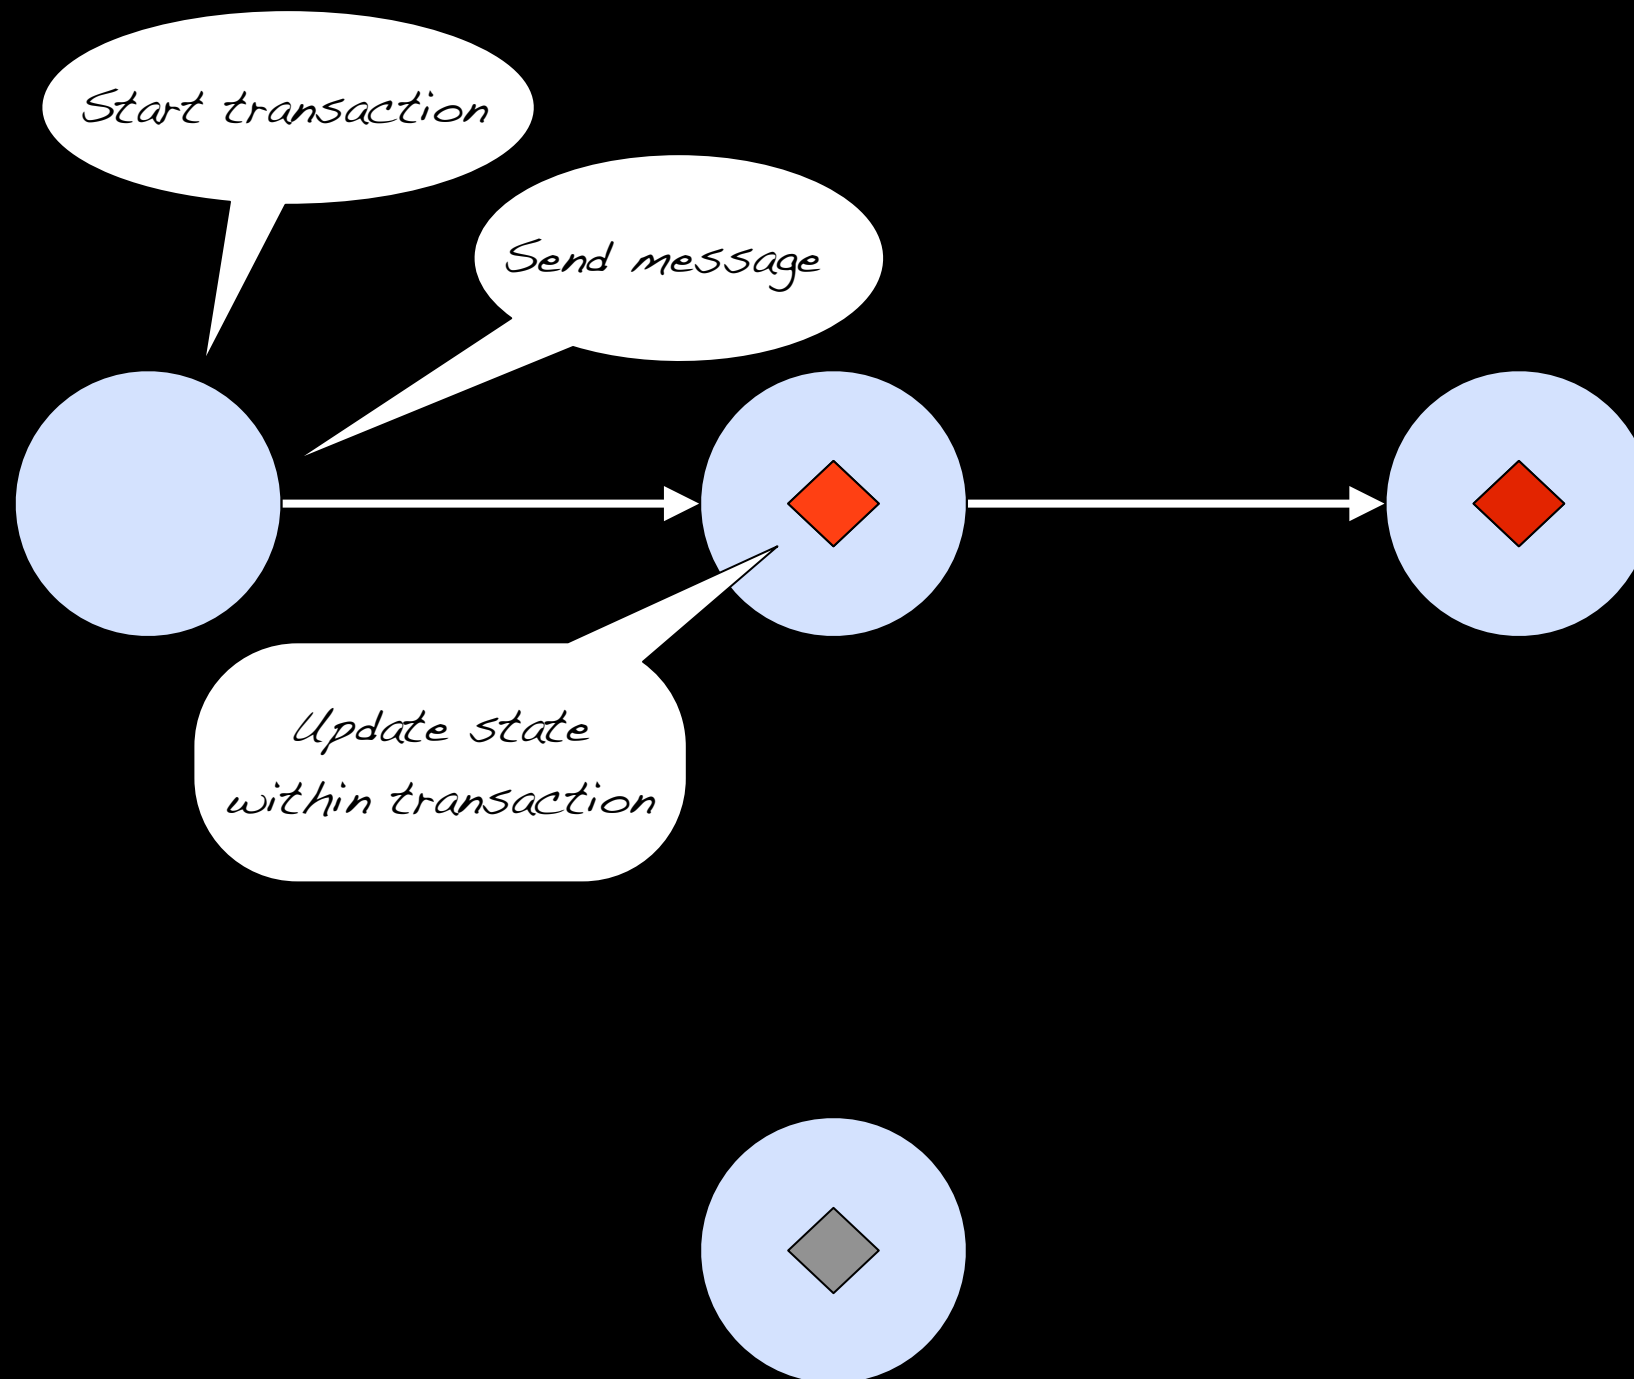
Transactors



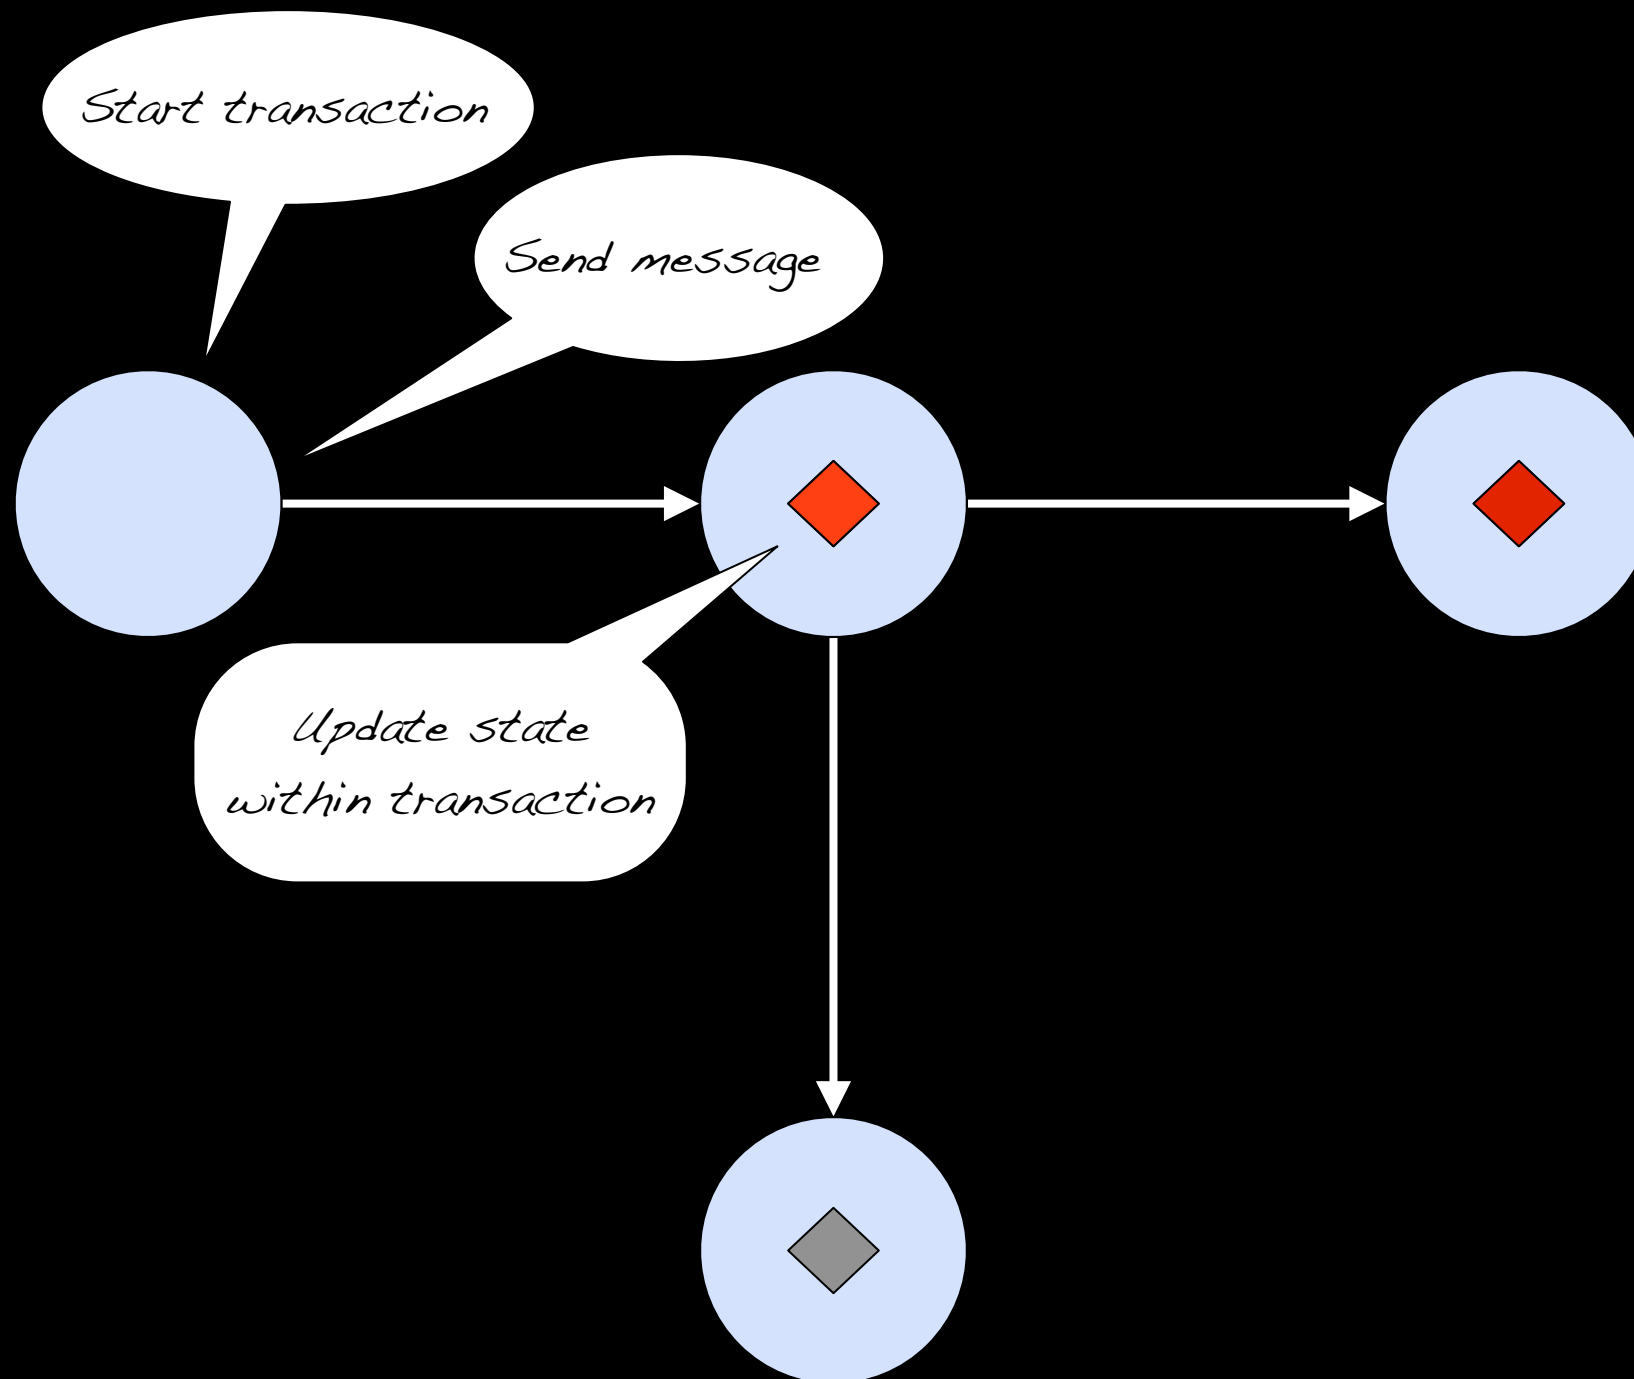
Transactors



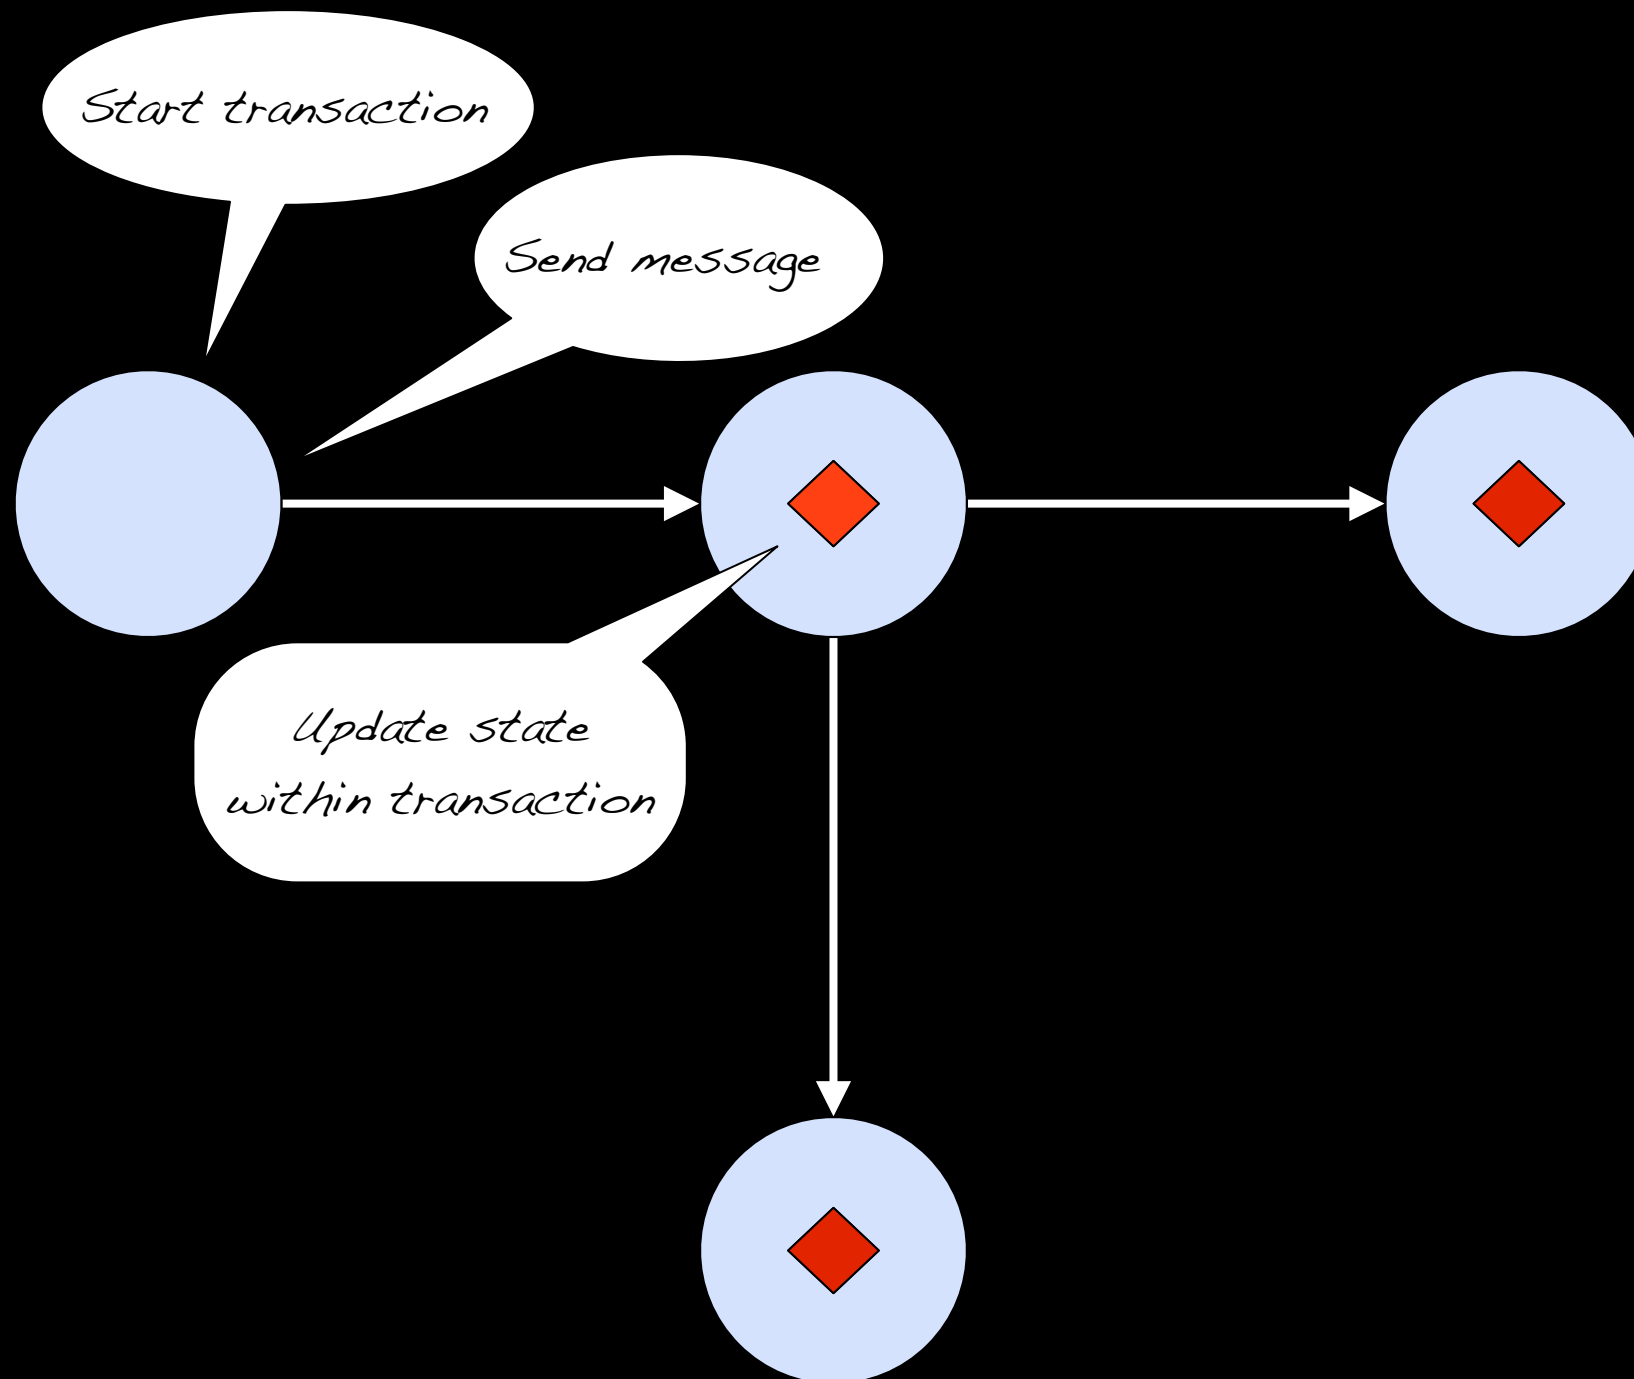
Transactors



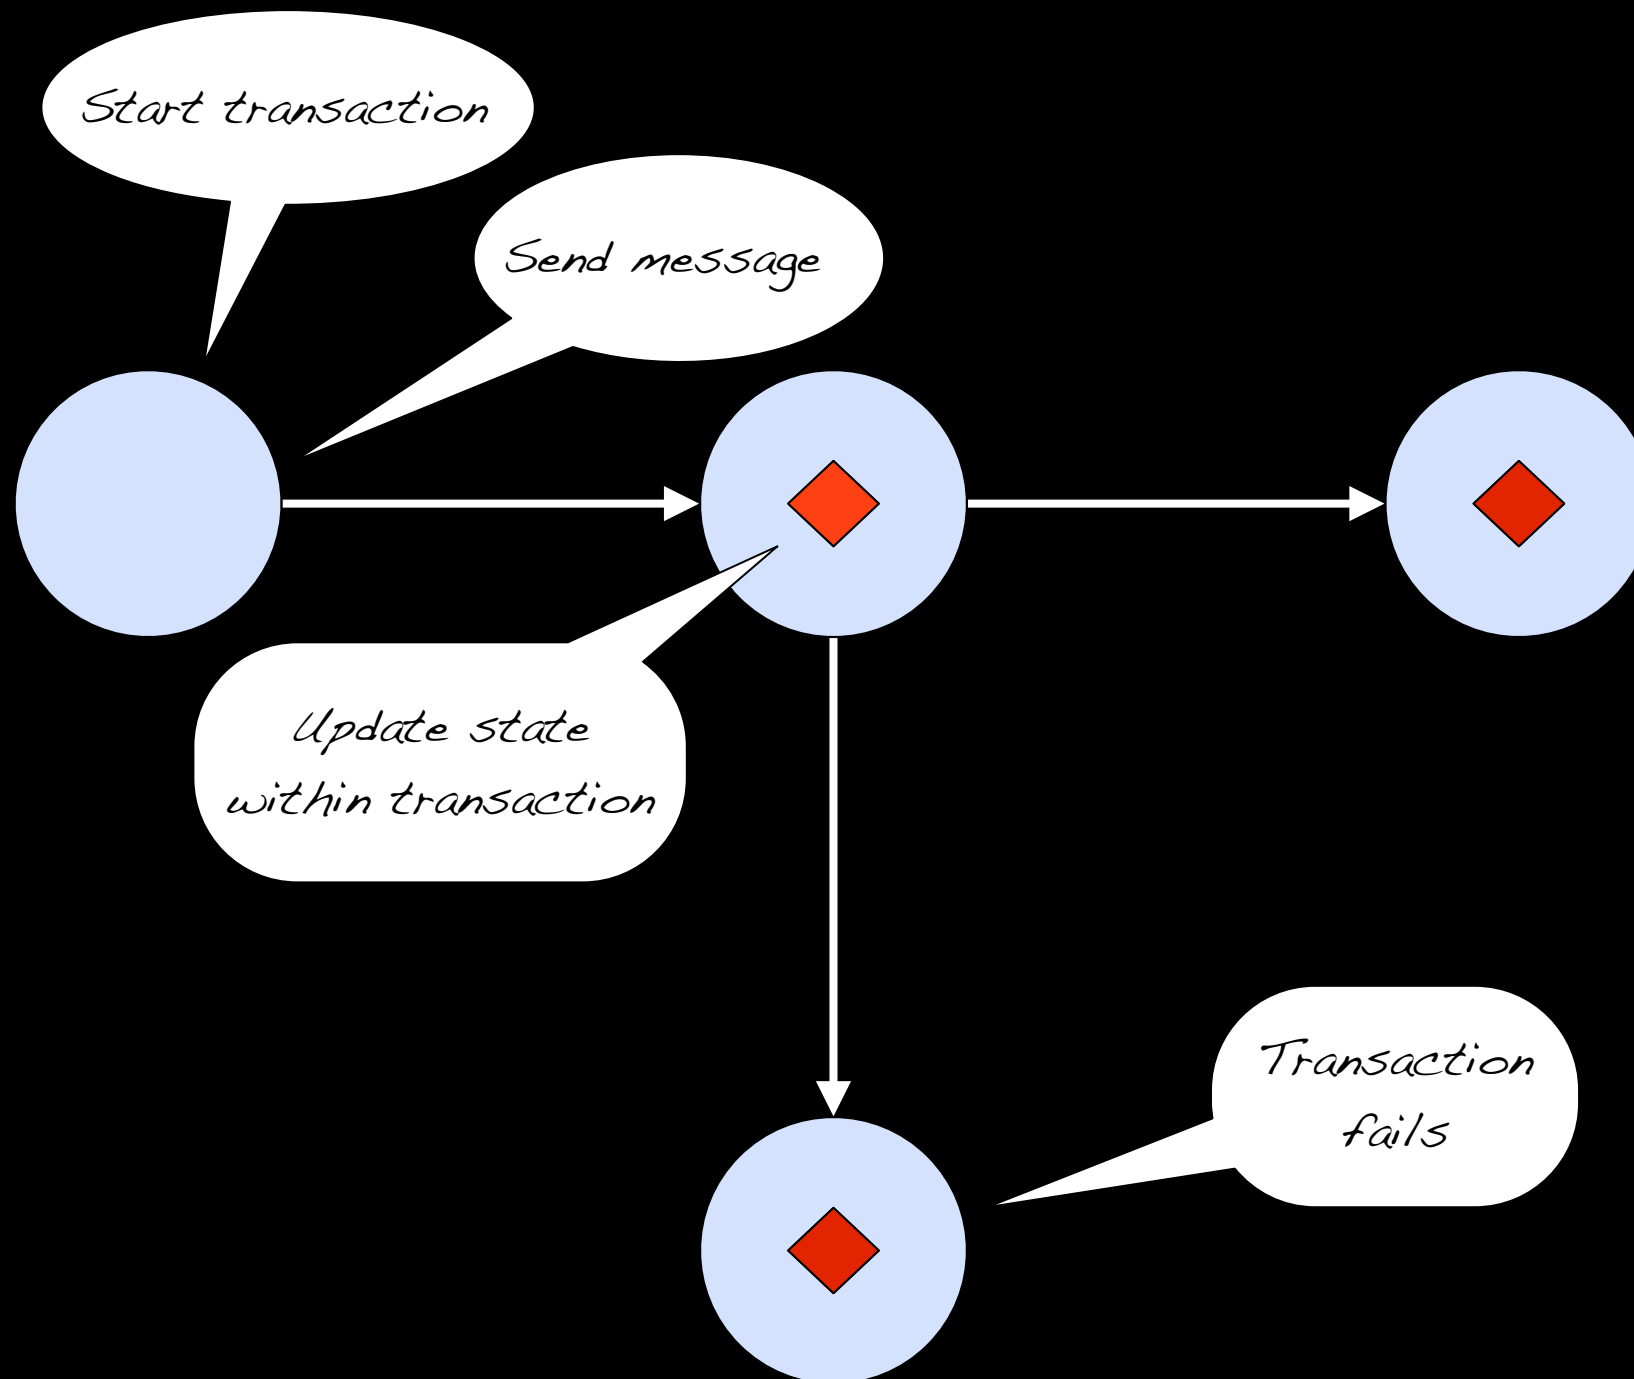
Transactors



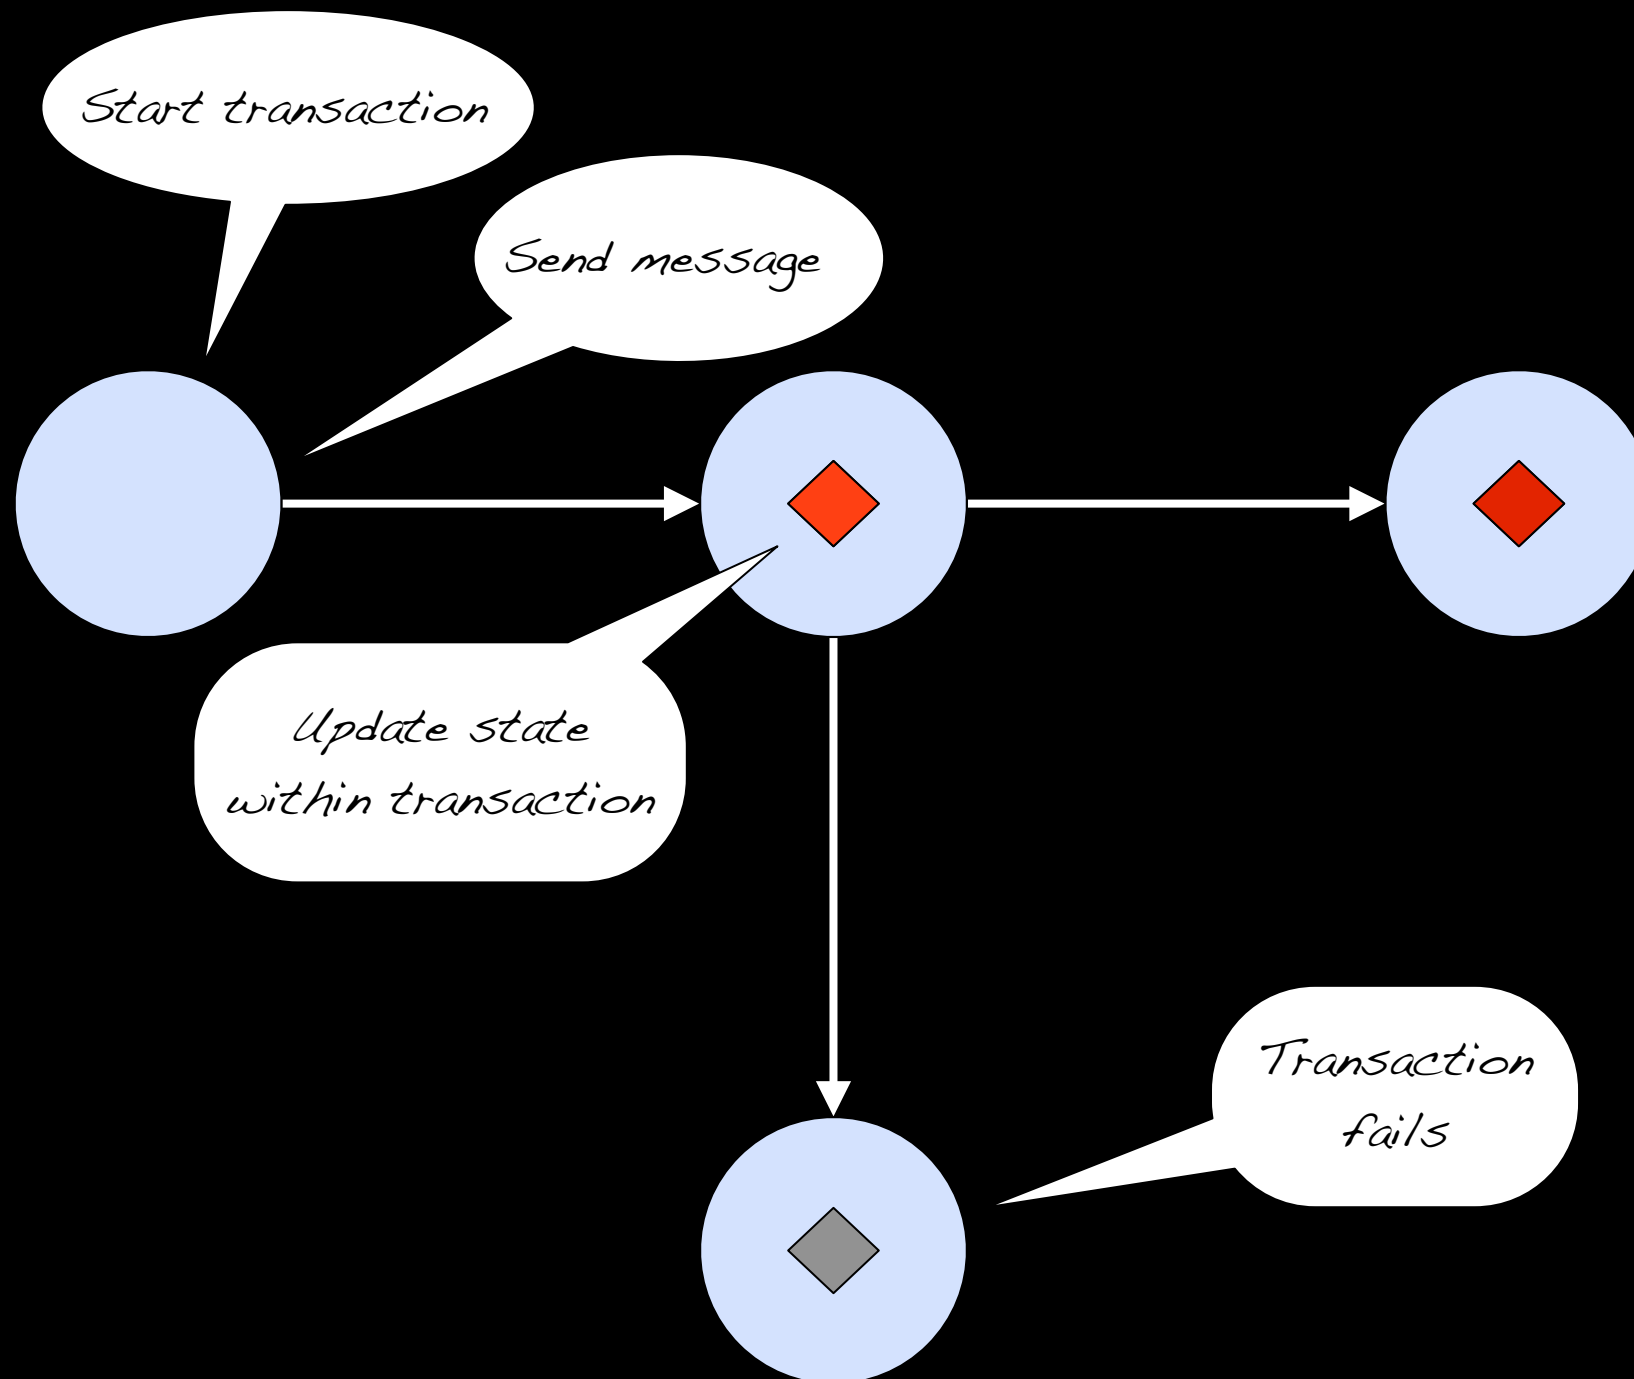
Transactors



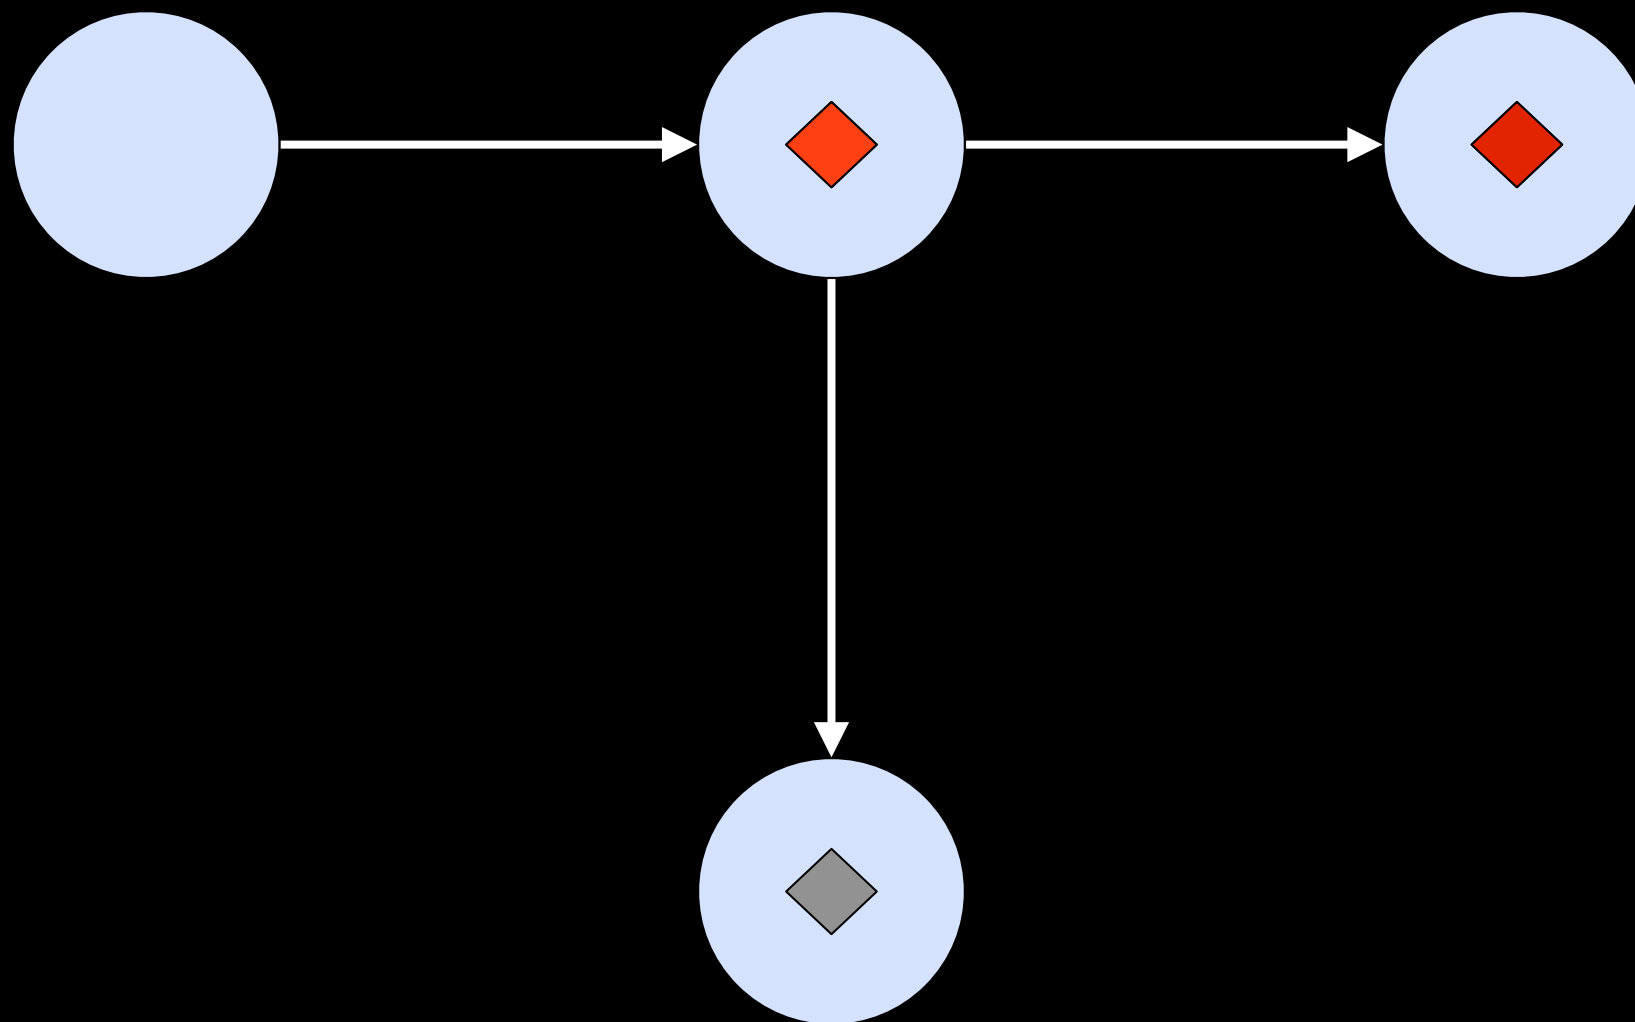
Transactors



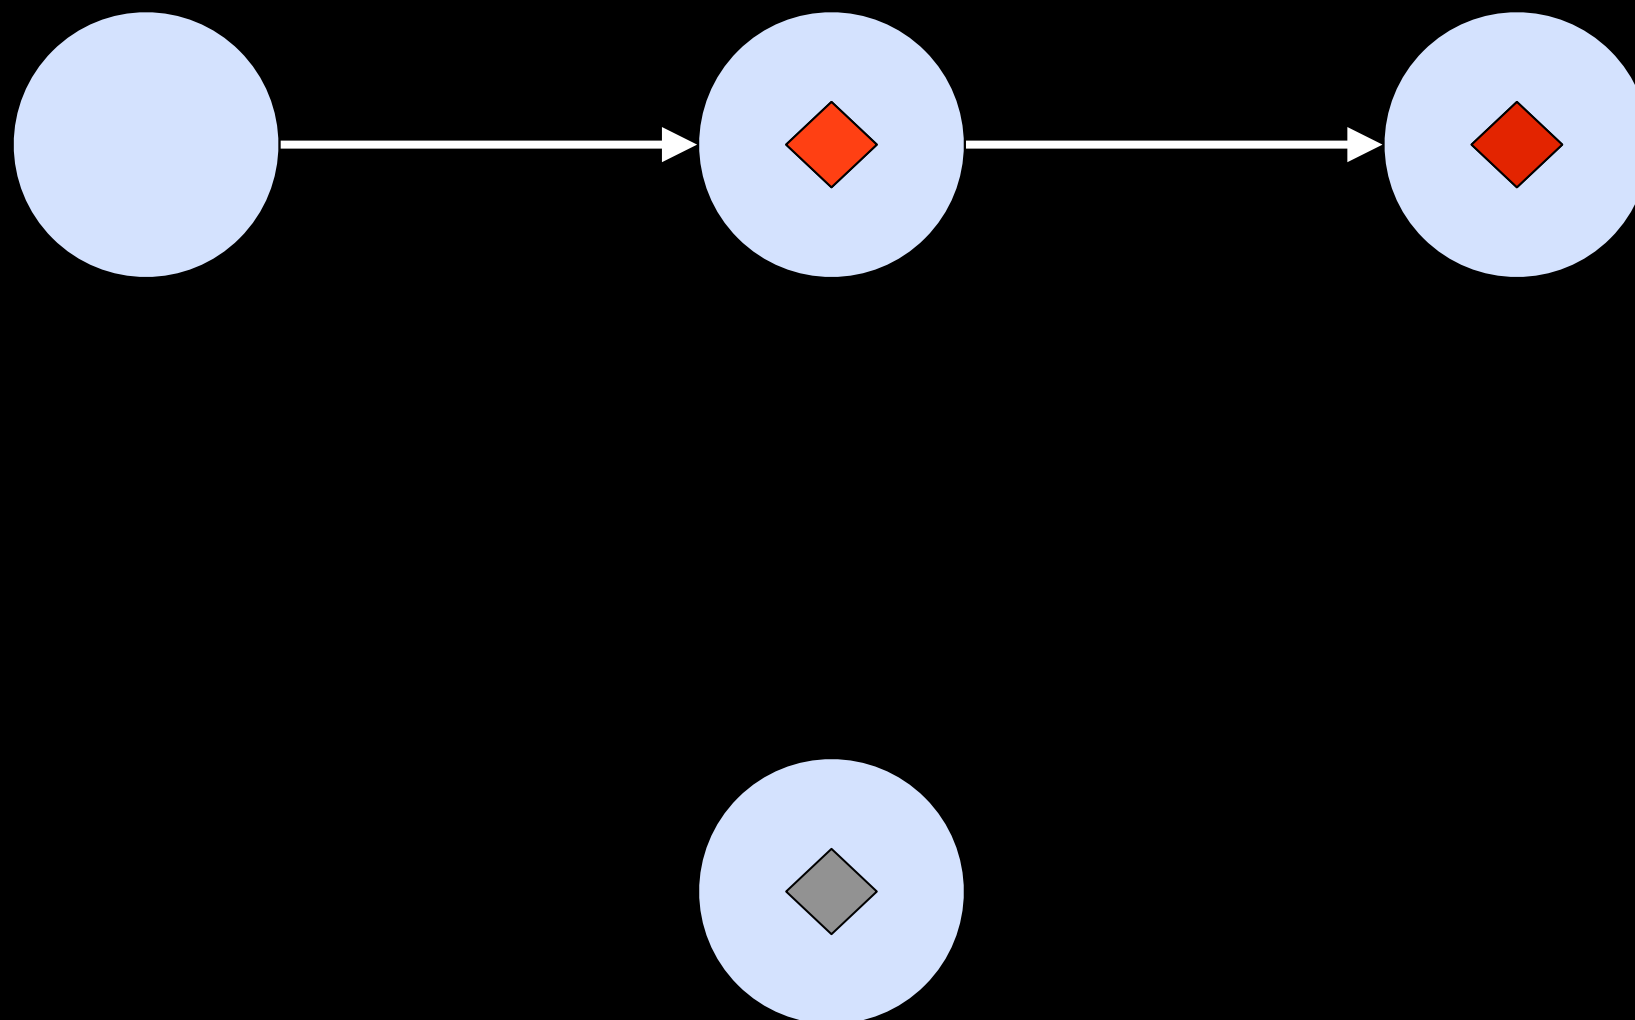
Transactors



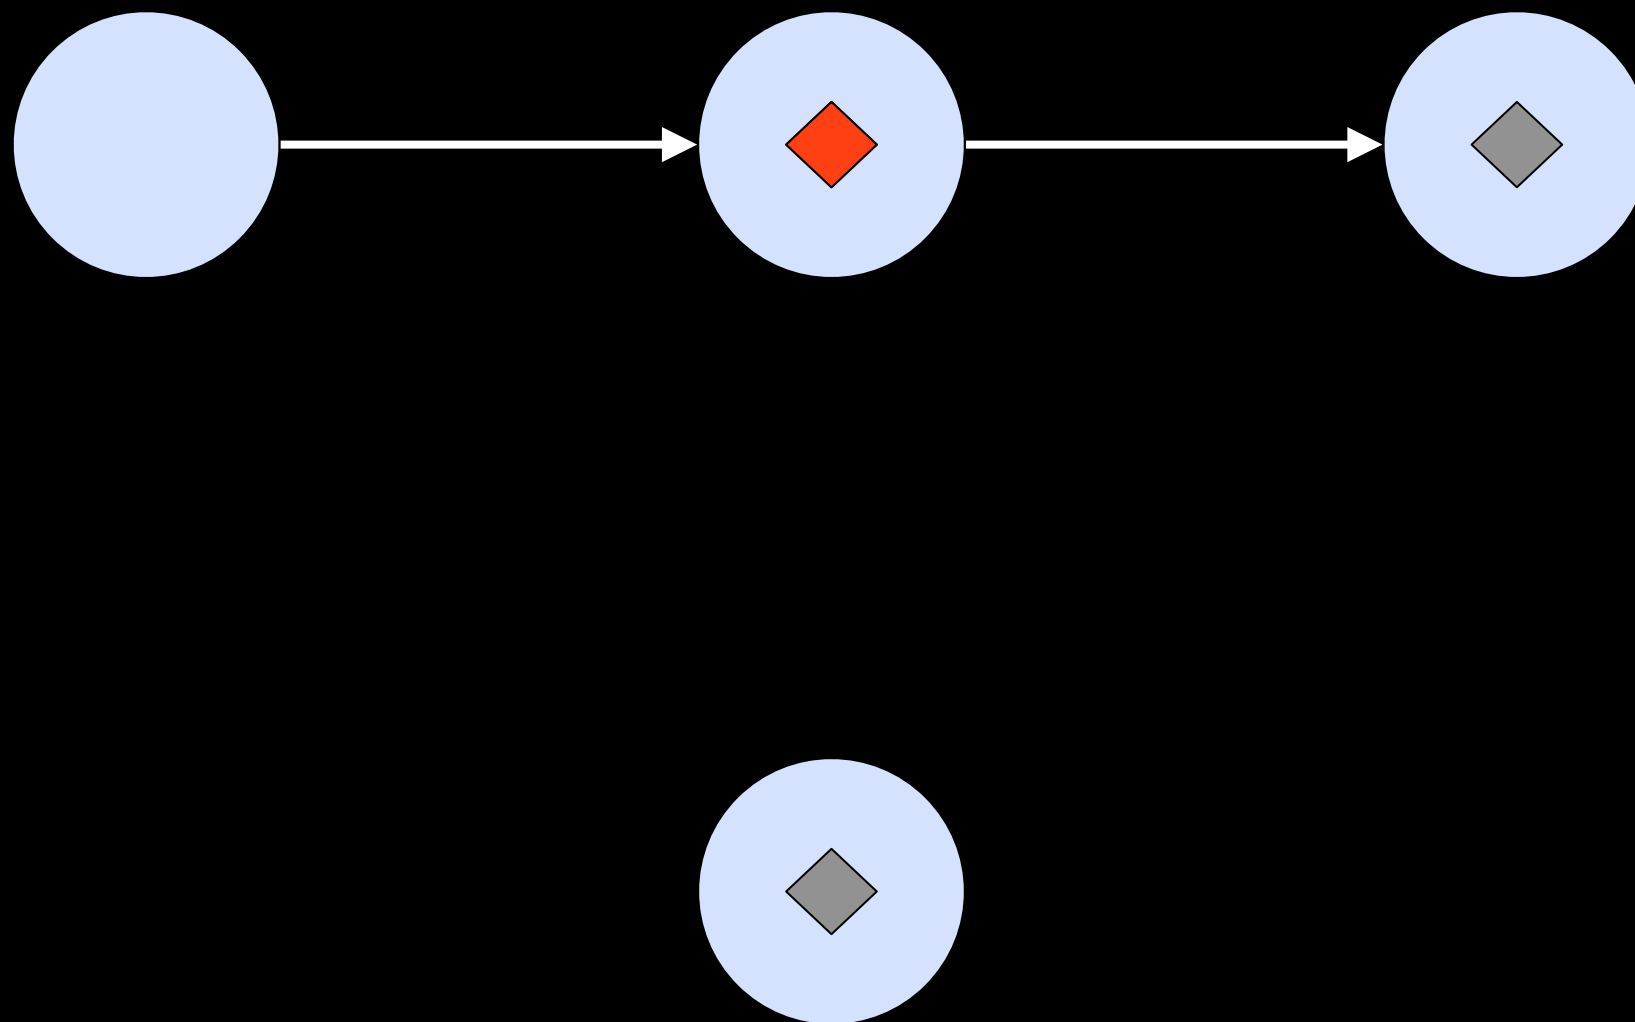
Transactors



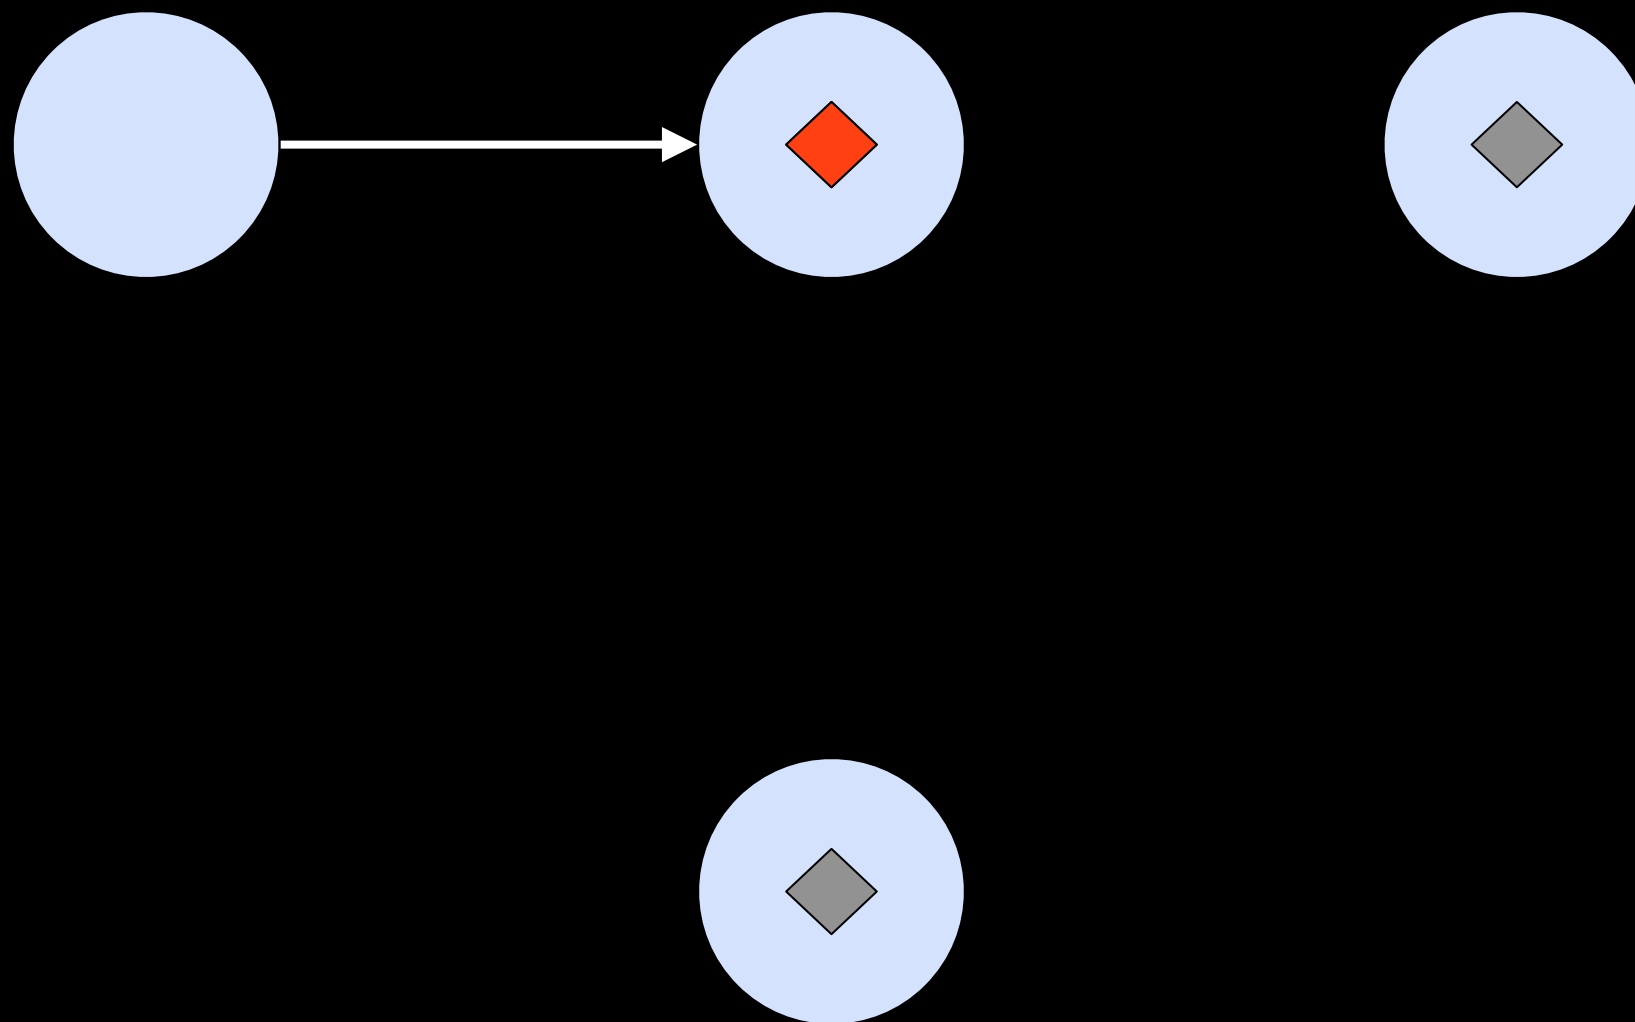
Transactors



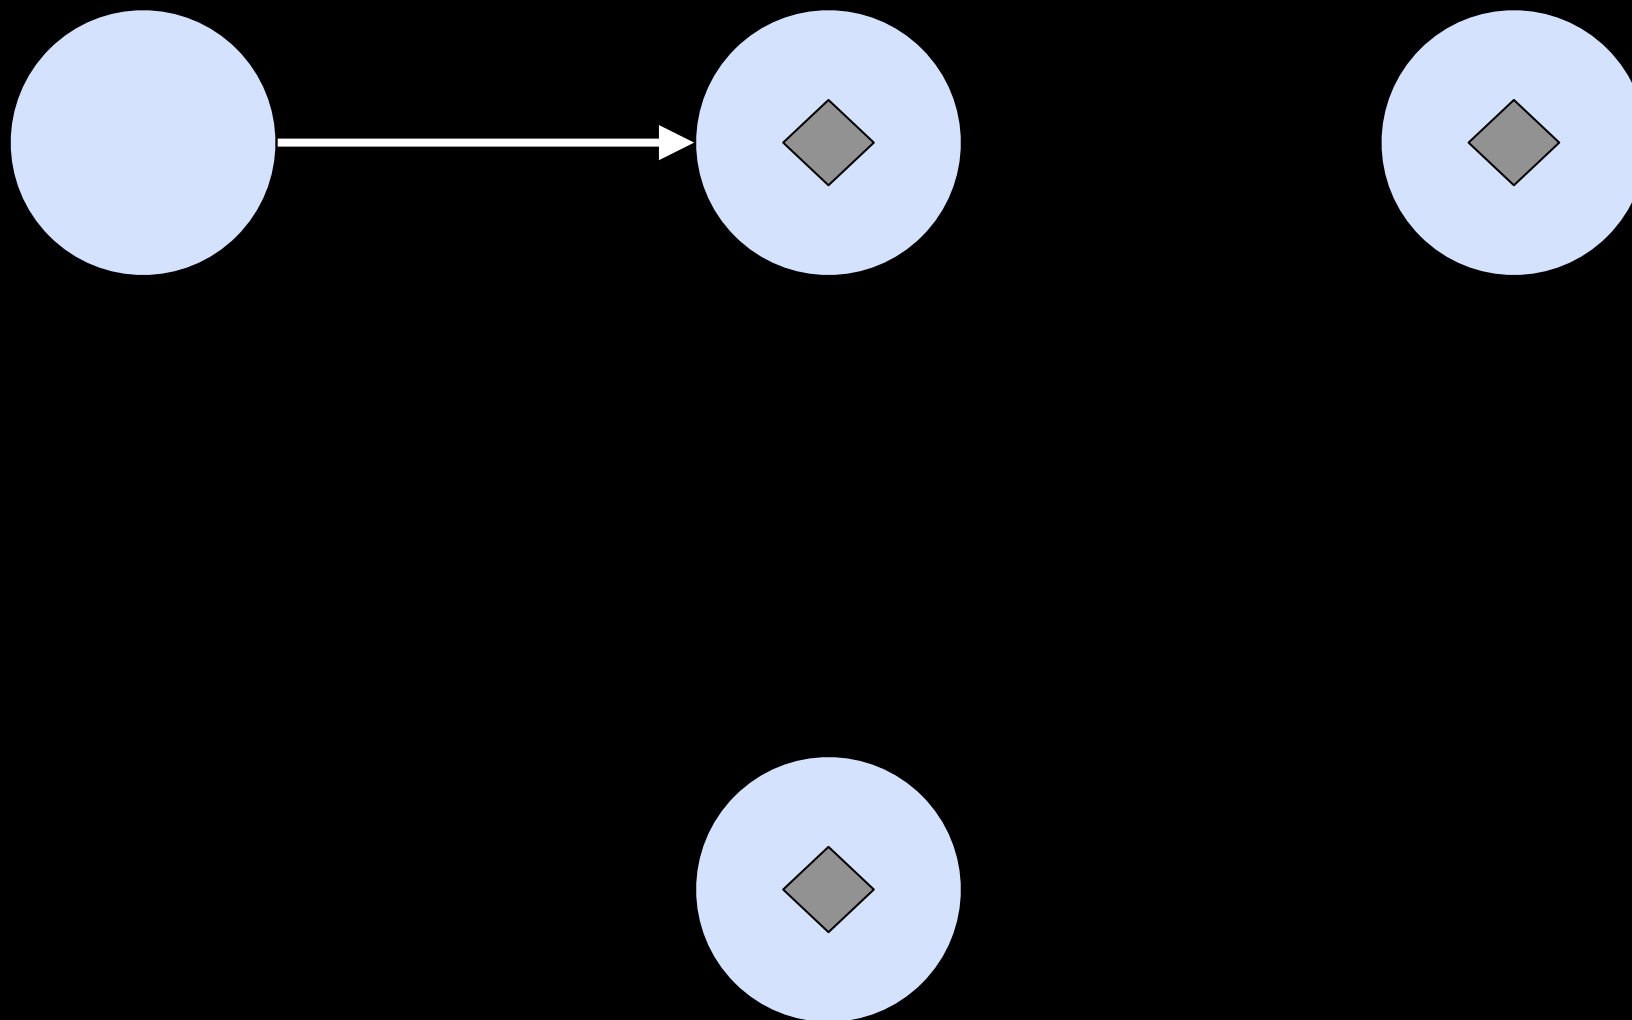
Transactors



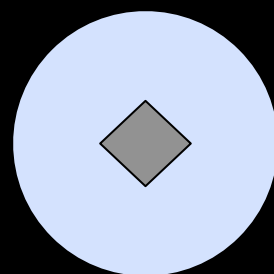
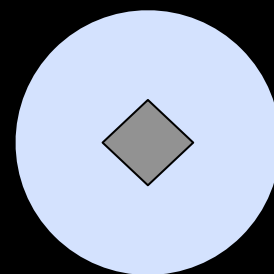
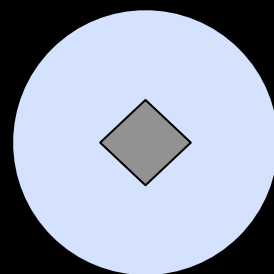
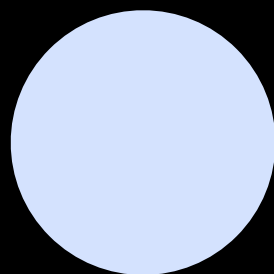
Transactors



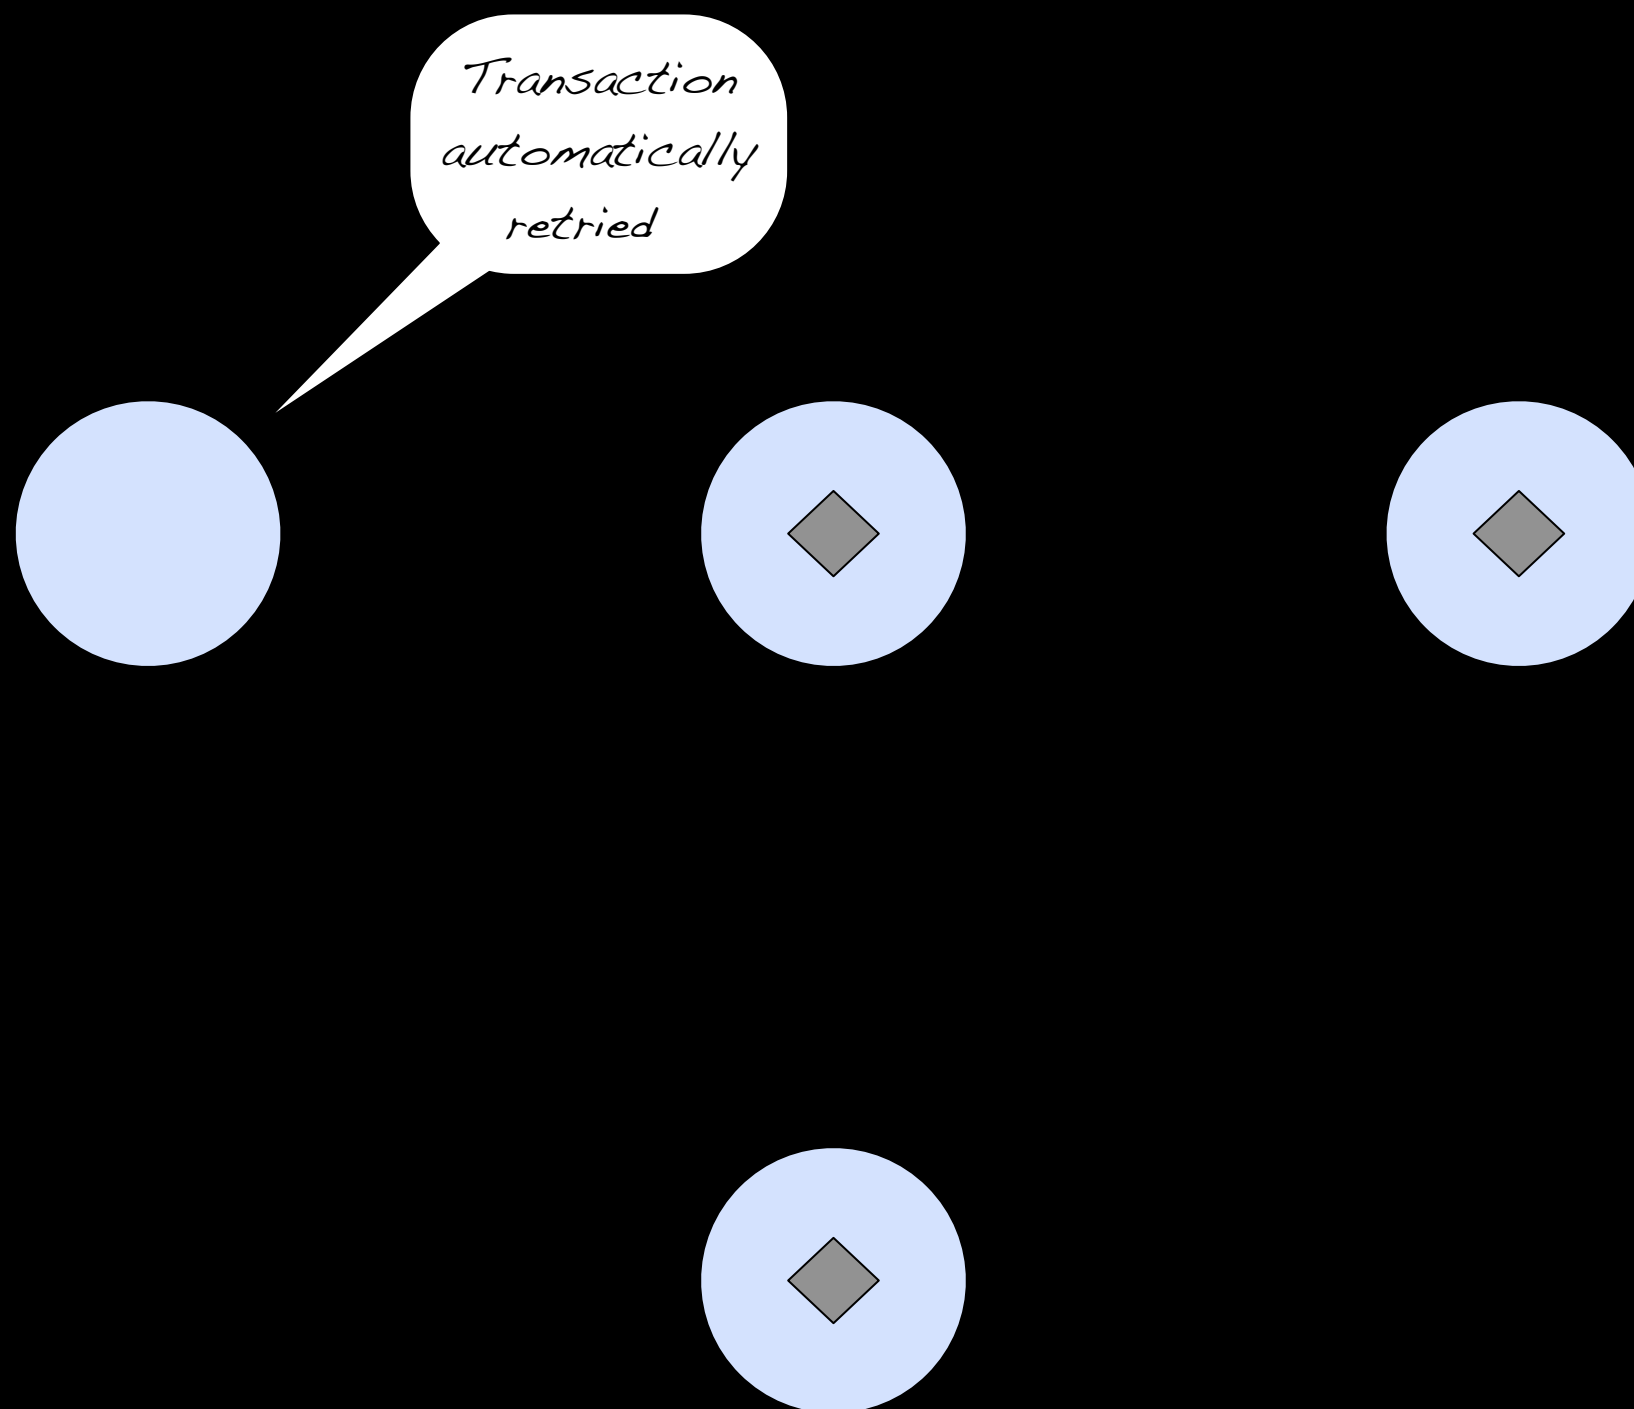
Transactors



Transactors



Transactors



REPL time!

Transactor Overview

- All the good from **Actors**
- All the good from **STM**
- **Removes** state update **coordination issues** from Actors
- **Replaces Threads** as the concurrency mechanism for STM

Transactor gotchas

- Non-transactional changes may occur if such code is called within a transaction
-

Summary

Threads and locks
are...

...sometimes

plain evil

It doesn't have to
be like this

We need to raise the
abstraction level

...but **always** the

wrong default

Introducing



STM

Actors

Agents

Dataflow

Distributed

Open Source

RESTful

Secure

Persistent



www.akkasource.com

<http://github.com/viktorklang/transactors>